

000 Official training rows. Grouped 80/20 split. Official test is not an input.

001 Residual generator. At most 12 inputs, width 16, tanh. AdamW, 2,048 steps, learning rate 0.002.

002 Refit readout a , β , and v on the frozen tanh features.

003 Selection on each method's native objective. Not the shared retention.

004 Common measurement on validation rows. Auditors: CatBoost, linear, and MLP.

005 Retention and charged bytes.