

SUPPLEMENT: DOPE: DOWNSTREAM OBJECTIVE-PRESERVING ENCODING OF REGRESSION TABLES INTO KILOBYTE GENERATORS

Anonymous

Affiliation withheld for review.

ABSTRACT

This supplement holds the method details, the release-gate specification and its conformance test, the data and selection protocol, the statistics, every five-seed contrast and level behind the main paper, the controls, the fidelity and empirical-privacy panels, the ablations, the baseline coverage, the per-lineage tables, and the extended related work. A blank retention cell is an unavailable fit or an auditor that is uninformative on one of the three sample seeds; it is never a zero. BeyondArena is a separate fit check and is never pooled with the PMLB panel.

A METHOD DETAILS

A.1 ARTIFACT CONTENTS AND SAMPLING

Equation (1) of the main paper is the conditional mean; the joint draw is separate. The profile reported in the main paper samples with the symbolic kernel stored in the artifact; the separate neural-joint sampler is unused. The stored dependence is one choice from the compile search: independence, a Chow–Liu tree, a triangular autoregression, a sparse Gaussian copula, a factor model, a truncated vine, or a two-component mixture, and the reported artifacts span several of these families.

Training screens the inputs by absolute rank correlation with the target and keeps at most twelve. The artifact stores a marginal and a transform for each feature, the selected dependence, the residual, a noise model, and a per-feature missing rate. A draw samples a latent vector from the stored dependence, maps each coordinate through the standard normal distribution function, inverts the marginal, and inverts the feature transform; the residual then supplies $\hat{y}$. When the training target is not binary, the stored noise is heteroscedastic: a linear fit of the log absolute residual on up to four screened features, with the scale clamped between the lower and upper deciles of those absolute residuals. The sampled target adds Gaussian noise of that scale and is clamped to the unit interval of the training transform. Each feature is then replaced, independently, by a missing mark with its stored probability. Before encoding, the compiler quantizes the marginals, the dependence, the residual, and the noise model. The charged byte count of a DOPE artifact is the encoded model file plus `projection.json`; a baseline is charged the files its sampler reads, as listed in Appendix C.3.

A.2 TRAINING FREEZE, COMPILE SEARCH, AND THE REPORTED PROFILE

Training maps the screened inputs through the width-16 tanh layer and runs AdamW on the regression mean-squared error, with gradient clipping at 5. The hidden weights and biases are then frozen, and the intercept, the linear skip, and the hidden readout are refit on the frozen features. Coefficients whose absolute value is below 10^{-5} are dropped from the skip.

The compile search records the proxy $0.45u + 0.25d + 0.30j$ and keeps the candidate with the smallest normalized gate shortfall, then the smallest artifact, then the lowest candidate id. It reads




Figure 1: **Pipeline of the reported profile.** The official test is outside the inputs. No lineage denominator and no Holm family.

the fit rows only. A candidate record carries no release certificate, and no global DOPE family has been selected.

The reported research profile, `features12_steps2048`, was chosen among four bounded research profiles (twelve or twenty-four screened inputs, 512 or 2,048 steps) on a six-lineage discovery panel and checked on six disjoint confirmation lineages, by validation retention at fit seed 11. Those twelve lineages belong to the hundred-lineage panel, so the reported profile carries a mild selection on validation rows. The untuned trainer default, used for any other profile, sets the hidden width between 4 and 12 and trains for 96 steps; the main paper reports it as a separate arm. Other neural shapes (micro-TVAEs, TinyMAT, joint TVAE, a masked transformer, and a TabSyn candidate) are outside this measurement, and the direct-rank TabDDPM candidate is a research comparator with no claim on the 10,240-byte tier.

A.3 DATASET SKETCH AND ROUTER VECTOR

An index can store a fixed sketch of a training table, which a separate dataset-action router consumes; neither enters the generator reported in the main paper. `embed-dataset min-max` normalizes each training column onto the unit interval, leaves a missing entry missing, and never reads the official test. Each column becomes twelve content summaries: a missing rate, moments, quantiles, a unique-value share, a tail ratio, and the absolute association with the target. Columns are sorted by those summaries, variance first. A ten-number header records the task, the log row count, the log column count, and the shape of the target. Each of the twelve summaries is then reduced across columns to a six-number distribution (minimum, lower quartile, median, upper quartile, maximum, and mean), which is the 12×6 block. The sketch then keeps the first 64 columns and stores each column’s twelve summaries, padding a shorter table with zeros, and absolute correlations among the kept columns collapse to one further six-number distribution. The width is therefore $10 + 12 \times 6 + 64 \times 12 + 6 = 856$. The router vector appends 24 candidate states, each a hidden state of width 128 plus 10 action predictions, to that sketch, for $856 + 24 \times (128 + 10) = 4,168$ dimensions. A neighbour found with the sketch is still fit with the reported profile and still charged in bytes.

B RELEASE GATE

B.1 MFS-v2, THE HISTORICAL FORMULA

MFS-v2 is the release score that historical receipts pin. With component $c_i \in [0, 1]$, weight $w_i > 0$, $\epsilon = 10^{-6}$, and $W = \sum_i w_i$,

$$\text{MFS-v2} = 100 \exp \left(\frac{1}{W} \sum_i w_i \ln(\epsilon + c_i) \right)$$

when the hard-gate map is non-empty, every hard gate passed, and every positive-weight component is finite and inside $[0, 1]$; otherwise the scalar is null. Its weights are utility transfer 0.30, driver fidelity 0.20, distribution fidelity 0.20, structure fidelity 0.15, coverage realism 0.10, and compactness 0.05. Privacy has soft weight zero, and a membership failure, an attribute failure, or a copied row is a gate that utility cannot offset. Compactness is $\text{clamp}(1 - \text{bytes}/10240, 0, 1)$, and an artifact past 10,240 bytes fails the byte gate.

B.2 WHAT THE COMPONENTS MEASURE

Utility transfer is a one-sided lower bound on retention, clamped into $[0, 1]$; it is distinct from the median retention in the main paper. Driver fidelity is an equal mixture of feature-importance rank agreement and top-feature-set overlap between the real and the synthetic table. Distribution fidelity is half the minimum marginal fidelity, plus a quarter of the sliced-Wasserstein fidelity and a quarter of the maximum-mean-discrepancy fidelity. Structure fidelity is the minimum dependence fidelity. The driver term stored beside the panel is the same equal mixture, applied to absolute correlations with the target; it differs from the permutation-importance auditor in the v2 definition, and this paper reports no component value from it.

B.3 MFS-V3 AND REPRESENTATION CLOSENESS

MFS-v3 adds representation closeness to the geometric mean, as in equation (2) of the main paper. v3 utility transfer is the same one-sided bound on the Kumo large retention: the model is conditioned on the synthetic rows and scored on real holdout rows. Mitra and TabICL must be measured on the same mapped tables for the gate. Their retentions stay out of that bound, and a CatBoost retention cannot fill this component. Real fit rows, real holdout rows, and synthetic rows share one map, fit on the real fit rows only: numeric columns go to the open unit interval by an empirical mid-rank quantile, categorical levels are pre-encoded to fixed codes in that interval, and column names become HMAC digests. The map withholds the minimum and the maximum, and a second fit on the synthetic rows is a protocol failure.

Representation closeness asks whether Kumo large places the mapped synthetic table near the mapped real table. The distance is half the Euclidean length between the two unit-length vectors, so it lies in $[0, 1]$, and closeness is one minus that distance divided by the distance to a matched null, clamped to $[0, 1]$. The synthetic table must also sit closer than an independent-marginal table of the same size under the same map. The matched null is chosen among three independent-marginal tables drawn on the shared map, seeded by the first eight bytes, little-endian, of the SHA digest of the dataset, the method, the configuration, the sample seed, and the candidate index 0, 1, or 2; candidates are ordered by ordinary linear mean squared error on the mapped holdout, then by index, and the middle candidate is the matched null. Those losses are fixed before any encoder runs. Mitra and TabICL must show the same sign as Kumo large on the gap against the independent-marginal table; a sign flip, a zero match distance, or a missing vector makes the component null. Neither is averaged into representation closeness. The licensed encoder set is Kumo large (OpenMDW and Apache), Mitra (Apache), and TabICL (BSD); Foundation, TabPFN, TabDPT, the project's Hyperion, and target-aware blocks are refused. TabDPT is refused because its pretraining tables overlap this panel; whether the pretraining corpora of the three admitted encoders overlap PMLB is unknown, so their retentions are reported as transfer measurements rather than as independent evidence.

A recipient keeps the generator and a hash manifest: a salt identifier, the normalizer name, the encoder pin, one digest per column, a dtype code, and the byte length. The manifest carries no row, header, category string, extremum, or encoder weight. A public sidecar that echoes a source string or an extremum fails the cleartext gate, and an exact copy of a real row fails its own gate. The near-copy floor is the 0.01 quantile, with linear interpolation, of squared Euclidean distances from each distinct real fit row to the nearest other distinct real fit row on the shared map; a synthetic row below that floor, an undefined floor, or a charge above 10,240 bytes makes the scalar null.

B.4 EXECUTABLE CONFORMANCE

The release score uses the contract's weighted geometric mean with a component floor, clipped to $[0, 100]$; the components lie in $[0, 1]$ and a failed gate yields a null scalar. The matched null is the median linear-MSE candidate among three deterministic independent-marginal tables. This is a median selector; a learnability band is outside this implementation. The near-copy floor is the declared quantile of nearest-other squared distances among distinct real fit rows on the shared map, with linear interpolation, and undefined floors fail the gate. The scorer tests in `research/benchmark/tests/test_mfs_v3.py` verify exactly 100 for a perfect profile, the byte boundary, and each prerequisite; the panel tests verify median null selection and the

Table 1: **MFS-v3 contract fields as implemented.** Each row is read from `production/kpi-contract-v3.json` by the scorer. No Holm family.

Scalar	$100 \exp(\sum_i w_i \log \max(s_i, \epsilon) / \sum_i w_i)$, clipped to $[0, 100]$
Weights	compactness = 0.05, coverage_realism = 0.05, distribution_fidelity = 0.15, driver_fidelity = 0.15, representation_closeness = 0.15, structure_fidelity = 0.1, utility_transfer = 0.35
Epsilon	1×10^{-6}
Bytes	$b \in \mathbb{Z}$, $0 < b \leq 10240$
Prerequisites	Real-vs-real control; required sizes; artifact sampling; locked metric implementations.
Matched null	Median linear MSE of three independent-marginal candidates; candidate-index tie break.
Near-copy floor	$q = 0.01$ of nearest-other distinct real-row squared distances; linear interpolation.

Table 2: **Invented applications of the bounded MFS-v3 formula.** Arithmetic illustrations only. No Holm family.

Invented input	Scalar
All seven components equal to one	100
Compactness component set to zero	50.12
Byte gate failed	null

quantile floor. Table 1 lists the implemented fields, and Table 2 applies the bounded formula to invented components.

B.5 CLEAN-LICENSE TRANSFER ON THE DENSITY BLOCK

The density block was also scored by Kumo large, Mitra, and TabICL on the shared map, with the three sample seeds of the main panel. These auditors are separate from the CatBoost, linear, and MLP rows, and their Holm adjustment is one family per auditor of size three (that auditor against the three density comparators). Kumo large averages 16 inverse-transformed quantile forecasts. The receipt scores 0 of 388 method-lineage cells ($97 \text{ lineages} \times 4 \text{ methods}$), so the scalar is null. Cleartext is unscanned on 388 method-lineage cells.

C DATA, SELECTION, COST, AND AVAILABILITY

C.1 DATA LOCK AND PARTITIONS

The data lock admits 104 MIT-licensed entries and prepares 100. 4 lineages have no split object because the same source-row digest appears in two partitions (`check_nonoverlap` raises on that overlap), and another 3,755 catalog entries lack a redistribution grant and are excluded. The prepared cut is a grouped 80/20 split of the official training rows, seed 1729. Worker directories hold `train.csv` and `validation.csv` only; the official test is held elsewhere and its contents are unread. The fit view has 72 to 800 rows; sources above the quota use a hybrid coreset that keeps the observed feature and target range. Of these prepared lineages, 20 names begin with feynman, which the PMLB catalog describes as synthetic physics, 14 begin with strogatz, which that catalog describes as simulated dynamics, 25 contain the Friedman-style fri marker, 6 contain bng and are Bayesian-network generators, and 35 are the remainder. This paper does not re-verify which of the remainder are observational measurements. Distinct lineage identifiers are not a claim of independent source families. Related names, including the `cpu_small` lineages, stay inside the one-row-per-lineage bootstrap. The resampling unit is the lineage, so related variants move together only by chance. The printed family-cluster intervals use a two-stage bootstrap. Lineages in the four simulated name families (feynman, strogatz, fri, and bng) form one cluster per family, and every other lineage is a singleton cluster. Each draw resamples clusters with replacement, then

Table 3: **Kumo large, Mitra, and TabICL on the density block at $4n$.** Each cell is a retention; the denominator is the paired informative lineages in that row. Holm family size is three per auditor. W/T/L counts lineages; a tie is a difference of zero.

Auditor	Comparator	n	DOPE	Other	Difference	W/T/L	Holm p
Kumo large	Gaussian copula	97	0.863 [0.809, 0.880]	0.747 [0.668, 0.778]	0.059 [0.0098, 0.116]	62/0/35	5.70×10^{-4}
Kumo large	Chow-Liu	97	0.863 [0.809, 0.880]	0.562 [0.501, 0.629]	0.247 [0.195, 0.306]	85/0/12	1.34×10^{-11}
Kumo large	Indep. marginals	97	0.863 [0.809, 0.880]	-0.000819 [-0.00189, 7.58e-05]	0.874 [0.834, 0.888]	94/0/3	7.46×10^{-17}
Mitra	Gaussian copula	97	0.858 [0.803, 0.887]	0.713 [0.585, 0.771]	0.066 [0.023, 0.135]	65/0/32	1.34×10^{-4}
Mitra	Chow-Liu	97	0.858 [0.803, 0.887]	0.583 [0.527, 0.643]	0.242 [0.182, 0.293]	85/0/12	2.71×10^{-11}
Mitra	Indep. marginals	97	0.858 [0.803, 0.887]	-0.00907 [-0.014, -0.00377]	0.890 [0.840, 0.909]	94/0/3	6.01×10^{-17}
TabICL	Gaussian copula	97	0.871 [0.835, 0.893]	0.754 [0.672, 0.787]	0.079 [0.00947, 0.116]	62/0/35	1.61×10^{-4}
TabICL	Chow-Liu	97	0.871 [0.835, 0.893]	0.587 [0.484, 0.650]	0.267 [0.200, 0.302]	87/0/10	7.12×10^{-12}
TabICL	Indep. marginals	97	0.871 [0.835, 0.893]	-0.00446 [-0.012, 0.0012]	0.891 [0.846, 0.919]	92/0/5	5.83×10^{-17}

Holm family: 3 tests per auditor; DOPE versus Gaussian copula, Chow-Liu, and independent marginals.

Table 4: **Pinned representation vectors.** Each row names the activation that the density-panel receipt reads as that auditor’s table vector. No Holm family.

Auditor	Pinned vector
Kumo large	<code>row_project_output_mean_labeled_rows</code>
Mitra	<code>final_layer_norm_input_mean_self_query</code>
TabICL	<code>tf_icl_output_mean_self_query</code>

resamples lineages with replacement inside each drawn cluster. A verified parent map is absent from the ledger, so those intervals put remainder lineages that share a parent in separate clusters.

The fit view is not the full PMLB table. Frozen catalog metadata draws each official training and test file by quota stratification on the target distribution and high-variance feature quantile bins, aiming at 1000 training rows and 1000 test rows, with at most 2000 rows in total and at least 101 rows. A source smaller than that quota keeps every row and still reserves a test share of at least 0.2. Of the 100 prepared lineages, 60 were already inside the quota. The other lineages use a hybrid coreset that keeps the observed feature and target range; the bin counts are 1 lineage at 11 bins, and 39 lineages at 50 bins. The generator then groups identical official-training rows and assigns those groups with seed 1729 so that about 80% of the rows become the fit view and the rest become the validation holdout. Duplicate rows are not split across that cut. On these lineages the fit view has 72 to 800 rows. The bound keeps every lineage on one worker view and leaves the official test sealed. It is not a second draw at scoring time.

C.2 SELECTION RULE FOR EACH METHOD

Table 6 states what chose each configuration and which rows it read. Only DOPE’s compile search and the density baselines’ fixed or density-selected configurations avoid the scoring auditors; CTGAN and TVAE select on the same validation rows and the same linear and MLP auditors that score retention, and DOPE’s research profile was chosen on a discovery panel inside the evaluated lineages. The sealed test run removes all of these overlaps for its estimate.

C.3 COMPUTE LEDGER AND BUDGET PROTOCOL

Table 7 lists fit time, peak memory, GPU models, GPU hours, and charged bytes for receipts that store a fit timer. Peak RAM is resident bytes in mebibytes, peak VRAM is `peak_gpu_used_mib`, and GPU hours sum the walls on which the receipt recorded a GPU process. No receipt stores a separate CPU timer. Sample-timer medians (maxima) are 186.7 (197.0) for CTGAN, 186.1 (196.6) for TVAE, and 342.4 (511.7) for Forest-Flow. Gaussian copula, Chow-Liu, and independent-marginal attempts store one unsplit `wall_seconds`, with medians 1.71, 0.168, and 0.162 and maxima 39.9, 1.29, and 0.429. The logged replay medians 4.56 and 14.4 are replay walls, distinct from the scored runs. Forest-Flow medians use the 12 native-grid operation walls; 953 sums `fit_seconds` and 8,801 sums receipt operations. Byte charges count the files a sampler reads: the DOPE model plus `projection.json`; every TabSyn file except intermediate diffusion-epoch checkpoints, which the sampler never reads; and the stored per-lineage charge for every other baseline.




Figure 2: **Kumo large, Mitra, and TabICL retention at $4n$.** Each bar is the lineage median and its lineage-bootstrap interval; the denominator is the paired informative lineages in that panel. Holm family size is three per auditor.

Table 5: **Test-split provenance.** File-name counts only. No retention denominator and no Holm family.

Item	Record
Catalog SHA-256	ab9fda8d2dea46067b70a42812e9d3c1d9dc6ba025780100df34377e81aa1120
Eligible / prepared / overlap exclusions	104 / 100 / 4
Unknown-rights catalog entries held out	3755
Split	official_test_grouped_training_80_20_seed_1729
Overlap lineage ids	590210721973399a, 9fe58a656fdb458a, d06a584f07d0acd5, d0ba629ed07b5d63
Final evaluation authorized	false
S3 worker directories / train.csv / validation.csv	100 / 100 / 100
test.csv under S3 workers	0
BeyondArena workers / worker test.csv	5 / 0
Evaluator test.csv files, contents not read	5

The frozen Forest-Flow budget keeps the 600-second fit cap and every timed-out attempt receipt; a timeout cell has missing quality metrics and counts as no DOPE win. The October 8 disposition addendum was written after 75 observed timeouts and is therefore no retroactive pre-registration. predeclare.v2 instead fixes a common budget (B1: 1,800 seconds per fit on one 16-core slot and one GPU) for every method and every lineage, applied non-selectively; docs/whitepaper/BASELINE_PROTOCOL.md links the frozen checkpoint and every disposition.

C.4 PUBLIC DATA AND ARTIFACT AVAILABILITY

A credential-free path to the source data is

```
python3 -m research.benchmark.fetch_pmlb --list
python3 -m research.benchmark.fetch_pmlb DATASET_ID OUTPUT_DIRECTORY
```

It authenticates PMLB data against an upstream Git LFS SHA pin at a fixed public revision. The public manifest lists every prepared lineage’s PMLB identifier and the historical train, validation, and test split hashes from the committed data lock. The historical staging used a hybrid coreset and target-stratified selection whose exact row indices were never captured, so the public fetch reproduces the source datasets and the split identities but cannot rebuild the exact historical partitions; publishing the fit and validation partitions themselves (PMLB is MIT-licensed) closes that gap and is scheduled with the artifact release. The originally scored model files and synthetic rows have no published bundle yet; the committed validation ledgers, the original metadata, and the compressed numeric logs, authenticated by the input pins and rebuilt by just paper, reproduce every reduction and figure.




Figure 3: **Paired retention differences for Kumo large, Mitra, and TabICL, sorted by lineage.** The denominator is the paired informative lineages in that panel. Holm family size is three per auditor.

C.5 REPRODUCIBILITY AND HASHES

Regenerate every generated/ file and the figure PDFs with just paper from the commit that contains these files. The same recipe writes the main PDF, the anonymous main PDF, the supplement PDF, and the anonymous supplement PDF. That recipe reads the committed validation ledgers, original fit metadata, and compressed numeric replay logs. The replay cost, loss curves, and BeyondArena extracts have file and field mappings in generated/original-field-map.json. Those original metadata snapshots were captured after the historical runs; they leave historical custody and the scored artifacts in place. Pinned LICENSE extracts are reread when their local store is available. Official test files are sealed until the single pre-registered run. The catalog hash is ab9fda8d2dea46067b70a42812e9d3c1d9dc6ba025780100df34377e81aa1120. The grouped training split uses seed 1729. The historical DOPE and legacy comparison fits use fit seed 11 and sample seeds 101, 211, and 307; the five-seed DOPE refit and its pins are in research/benchmark/review_fixes/predeclare_v2.json. Retained panels document the additional ARF and Forest-Flow fit seeds and the measured TabSyn and TabDDPM subsets separately. python3 research/benchmark/verify_paper_numbers.py exits nonzero when a displayed macro disagrees with those ledgers. The remaining baselines, the shadow-model and artifact attacks, and the full ablation grid enter the tables only after their runs close. The fidelity table is a separate empirical ledger on the grouped validation split. Empirical resemblance is no differential-privacy guarantee and no HIPAA de-identification.

Table 6: **Selection rule for each method.** The scored column is retention on the validation cut. No Holm family.

Method	Criterion	Data used	Budget
DOPE	Profile chosen on a six-lineage discovery panel and checked on six confirmation lineages; compile search keeps the smallest gate shortfall, then the smallest artifact.	Discovery and confirmation lineages' validation cut; compile search reads fit rows only.	At most twelve inputs, width 16, 2,048 steps, five fit seeds.
DOPE untuned default	No selection.	Fit rows only.	Trainer default, five fit seeds.
Gaussian copula	Held-out mean log density. Copulas 0.14.1 GaussianMultivariate.	Validation cut, distinct from the retention auditors.	One library configuration.
CTGAN	Mean of linear and MLP training-mean skill.	The validation rows and the two auditors used at scoring.	Author implementation; stopped runs omitted.
TVAE	Same skill as CTGAN.	Same rows and auditors as CTGAN.	Author implementation; stopped runs omitted.
Chow-Liu	Native: held-out mean log density over a fixed trial set. Default: no search.	Native: validation cut. Default: fit view, eight bins, Laplace smoothing 1.0. Written for this comparison.	Fixed trial set.
Independent marginals	Native: held-out mean log density over a fixed trial set. Default: no search.	Native: validation cut. Default: fit-view marginals. Written for this comparison.	Fixed trial set.
ARF	Author default; separately, the native objective.	Author default reads no validation rows; the native arm reads the validation cut.	Five fit seeds.
TabSyn	No selection.	Author architecture with a scaled schedule.	200 VAE and 1,000 diffusion epochs, fit seed 11.
Forest-Flow	Author regression objective on a frozen grid.	6 confirmation lineages.	Fit cap 600 seconds.

Table 7: **Fit time, memory, GPU, and charged bytes where a fit timer was stored.** The denominator is the receipts that stored a timer. No Holm family.

Method	Fit s	RAM / VRAM	Hardware	GPU h	Bytes
DOPE	7.25(13.5)	920.7/ 1067	one NVIDIA GeForce RTX 4090 and two NVIDIA GeForce RTX 3090 GPUs	0.193	1,895.5(758-22,233)
CTGAN	213.5(268.1)	1557.8/ 1101	NVIDIA GeForce RTX 4090	2.52	741,834(528,075-1,428,695)
TVAE	214.2(253.2)	1252.9/ 1069	NVIDIA GeForce RTX 4090	2.45	317,141.5(157,578-741,659)
Forest-Flow	185.9(380.6)	6112.0/ 3181	NVIDIA GeForce RTX 4090	2.44	325,848,913.5(115,202,253-1,085,155,062)

PMLB is MIT. Every dataset entry in the data lock records SPDX MIT and the note that the source license is MIT (PMLB). CTGAN and TVAE share one pinned LICENSE file: Business Source License 1.1, licensor DataCebo, Inc., licensed work CTGAN. The Gaussian copula uses the pinned Copulas LICENSE: Business Source License 1.1, licensor DataCebo, Inc., licensed work Copulas. TabSyn is Apache License 2.0. TabDDPM is MIT License, Copyright (c) 2022 Akim Kotelnikov. Forest-Diffusion is MIT License, Copyright (c) 2023 Samsung Electronics Co., Ltd. All Rights Reserved. ARF is MIT License, Copyright (c) 2023 Kristin Blesch, Marvin Wright. Chow-Liu and the independent marginals are implementations written for this comparison, under this repository's license: MIT License, Copyright (c) 2026 Dope contributors. The SDV repository LICENSE is absent from the pinned snapshot, so this paragraph leaves that license unrecorded.

The rows below are the SHA-256 digests of each ledger and lock the paper reads, plus the commit and the file hash of the harness that measured the grouped-validation fidelity ledger. Regenerating

the macros never rewrites a ledger hash.

Table 8: Ledger and lock digests. No Holm family.

Ledger or lock	SHA-256
Code commit of the measured harness Harness SHA-256	17f3294d139e19f0ec831605f8750368eebc30bb a952062c0f83805f6442a440a5ae15293a843349 d423eaa106a9c56a20d14c0f
research/benchmark/results/paper-original-m etadata-v1/index.json	6d73d5f7b37849526a98a451b41a49e045ef21e9 11058d920940c5ea989f7fd9
docs/whitepaper/generated/original-field-m ap.json	5b164d434414d3389f639d3645f40b11a999b32c 56daa2cfe36d25b2205082b6
docs/whitepaper/scripts/paper-receipts.py	a2d6d45e3ff06a38ed61ab029fadfd14529978bc fc89618aa9ceb3f5e57238ee
research/benchmark/expanded.valida tion.metrics.py	a952062c0f83805f6442a440a5ae15293a843349 d423eaa106a9c56a20d14c0f
research/benchmark/methods.lock.json	899c355cb55769710ed369fea6b082589b318588 60fb74929bf4190de28aca33
research/benchmark/results/expanded-validat ion-diagnostics.json	cd20ac0d07a1dcab7c5e46db195adb463172b517 4387bb71a054b2e5f340beeb
research/benchmark/results/s3-coreset-proce dure.json	73251e43cf7838687427569931bf6127f42e9323 74af33137c57fd8fe8544704
research/benchmark/results/s3-lineage-recor d.json	bbd0852c49f595315104261efa4eaa6f50db257e 323b8c002aac6737a2dbe7d7
research/benchmark/results/density-matched- population-validation.json	5264b88a40ad5efb21d9b119789caeb13e71e911 a17f1d8b57f1d8e8ac31732a
research/benchmark/results/sdv-matched-popu lation-validation.json	4dc367eb78159a0386882d23ae1f99b0aa2cc9f4 25d1aae499b6f3453222a44e
research/benchmark/results/s3-matched-fores t-confirmation-validation.json	97a25be902954cad16c5b5802d0e0ba468dea104 221efa2465434c94e6774a1f
research/benchmark/results/s3-data.lock.json	857b61324d0a5b7dbbb98e0c147fcc79c5bb8616 e7ca61afd6785469c17f038f
research/benchmark/results/arf-native-closu re-watch-v1.receipt.json	7e8eeb1ba94ca747367e98d5a016d2ed3ed4c40c 5a98ca76d7a5fe27ce9c173c
research/benchmark/results/beyondarena-s3-i nventory.json	504166764631dbdf5e02d116334af5cc587731fc f3efe130d62240e6dd512d64
docs/whitepaper/generated/fit-trace.json	d8c31b9e4eb59cb1a087cf5dcdd4e4a91a73eca00 f3704fbff647a7da6ce65522
docs/whitepaper/generated/loss-curves.json	0124915497c98e560ef99c1e5ecf7ad876f91744 39ff9c763cflcf906b4aeeb1
docs/whitepaper/generated/replay-cost.json	faf77e7ce1ed1d22127701bf68bbfe2f8d183f89 cbd3475dd80094888c36f086
docs/whitepaper/generated/provenance.json	dea64ea61eaacac0e287298bba3ba76aa7900819 aa00feb3a5d66cdcb4fb8a17
docs/whitepaper/generated/beyond-fit.json	a1447f908bc9c51876d345597039607c734b507b f85aa8e08db0639e80a1f036

D STATISTICS

D.1 RETENTION ALGEBRA AND VARIANCE

With r the retention, the ratio rearranges to

$$L_{\text{synthetic}} = L_{\text{real}} + (1 - r)(L_{\text{null}} - L_{\text{real}}).$$

Invented losses 1, 0.01, and retention 0.94 give a synthetic-trained loss of 0.0694, which is 6.94 times the real-trained mean squared error. The CatBoost median in this paper is not this example. Let $D = L_{\text{null}} - L_{\text{real}}$, so $R = (L_{\text{null}} - L_{\text{synthetic}})/D$ when $D \neq 0$. The partial derivatives are

$$\frac{\partial R}{\partial L_{\text{null}}} = \frac{L_{\text{synthetic}} - L_{\text{real}}}{D^2}, \quad \frac{\partial R}{\partial L_{\text{synthetic}}} = -\frac{1}{D}, \quad \frac{\partial R}{\partial L_{\text{real}}} = \frac{R}{D}.$$

If the three losses have covariance Σ , the delta-method variance of one lineage's retention is $\nabla R^\top \Sigma \nabla R$, which diverges as D approaches zero. The informativeness rule drops a lineage whose real improvement is below 0.01 times the absolute null loss as a stability cut; the tables report bootstrap intervals rather than this standard error.

D.2 INFORMATIVENESS THRESHOLD

The primary threshold is 0.01 times the absolute null loss. Table 9 lists the CatBoost median by threshold and the informative counts; each column uses its own informative count, so the Chow–Liu column differs from the density-block paired median of 0.586 in Table 20. Thresholds 0, 0.001, and 0.05 keep that median and that count. At 0.1 the count is 96 and the median is 0.942.

Table 9: **Informativeness threshold and auditor counts.** The denominator for each median is the informative count beside it. No Holm family.

Threshold	DOPE	Gaussian	Chow-Liu	Indep.
0	0.940 (97)	0.656 (99)	0.582 (99)	-0.012 (99)
0.001	0.940 (97)	0.656 (99)	0.582 (99)	-0.012 (99)
0.01	0.940 (97)	0.656 (99)	0.582 (99)	-0.012 (99)
0.05	0.940 (97)	0.656 (99)	0.582 (99)	-0.012 (99)
0.1	0.942 (96)	0.668 (98)	0.584 (98)	-0.012 (98)

Threshold	CatBoost	Linear	MLP
0.001	97	93	75
0.01	97	92	75
0.05	97	90	73

Reading the threshold counts from top to bottom, CatBoost is 97, 97, 97; linear is 93, 92, 90; and the MLP is 75, 75, 73.

D.3 BOOTSTRAP, FAMILIES, AND EQUIVALENCE

The v2 analyses follow `research/benchmark/review_fixes/predeclare_v2.json`, committed before any five-seed result existed. A fit value is the median of its three sample-seed retentions; a lineage value is the median of its complete fits (at least three for a five-seed arm); the panel level is the median of lineage values. Paired contrasts use the Hodges-Lehmann pseudo-median of lineage differences, which is the location the two-sided Wilcoxon signed-rank test addresses, and report the median difference beside it. Intervals come from a nested percentile bootstrap with 10,000 draws: the four simulated name prefixes are clusters and every other lineage is its own cluster; a draw resamples clusters, then lineages inside each drawn cluster, then fit seeds inside each drawn lineage, separately for each arm. Holm families are declared per size: the primary family of 12 tests, the density family of 9, a deep-generator family of 12, an other-generator family of 9, and a control family of 6; a member without data keeps its slot. Equivalence requires the 90% nested interval of the Hodges-Lehmann difference to lie inside ± 0.02 ; a 20% trimmed-mean (Yuen) test is a sensitivity. The v1 contract tested the mean paired difference with two one-sided t tests at ± 0.02 and unadjusted $\alpha = 0.05$; paired retention differences have heavy tails when the denominator is small, so that estimand is kept only as a continuity column below.

A Friedman test and a Nemenyi critical difference compare the four density methods on their common complete-informative set, with rank 1 for the higher retention; on CatBoost the Holm sign-test values are 2.15×10^{-15} , 6.72×10^{-24} , and 4.95×10^{-27} . A Bayesian signed-rank procedure would be an alternative to these frequentist tests and is not computed here.

E FIVE-SEED PANEL

Table 10 gives the retention level of every arm, auditor, and size, and Table 11 every DOPE-minus-comparator contrast with its family-adjusted p -value and both equivalence rules. Table 12 gives the seed-exchangeability diagnostics: how often seed 11 is the best of five, the estimate from seeds 23 to 71 alone, and the binary-by-seed diagnostic on the lineages common to every cell. The historical-binary refit reproduces the samples of the historical seed-11 population byte for byte on 86 of 98 lineages. The later binary is the one an earlier four-seed run used; it writes the same seed-11 artifacts as the historical binary on 100 of 100 lineages, and its samples match the refit at seeds 23 to 71 on 400 of 400 fits. Every retention cell in these tables, for DOPE and for every baseline including Forest-Flow, was scored again from its stored synthetic sample file by one pinned evaluator with auditor seed 1729; a cell whose sample file was missing stays out. Each statistic draws its own bootstrap stream, so two rows built from identical samples, such as the seeds-23-to-71 sensitivity and the later-binary four-seed arm, can differ in the third decimal of an interval. The fit-seed-11 privacy and detector panels in Appendix G come from separate earlier runs, which record the Forest-Flow cells as unavailable. The historical density, SDV, Forest-Flow, and fit-seed-11 DOPE receipts had used the sample seed as the auditor seed, while the ARF, TabSyn, and control receipts had used 1729; the uniform re-score removes that difference, so the

five-seed tables and the historical block tables below differ slightly for this reason as well as for the extra fit seeds.

Table 10: **Retention level of every arm.** Median of lineage values; a five-seed arm needs three complete fits per lineage. No Holm family.

Arm	Size	Auditor	Lineages	Fits	Median	95% nested interval
DOPE (ours)	n	CatBoost	99	495	0.824	[0.725, 0.942]
DOPE (ours)	n	Linear	94	470	0.986	[0.964, 0.993]
DOPE (ours)	n	MLP	88	440	0.969	[0.895, 0.992]
DOPE (ours)	$4n$	CatBoost	99	495	0.898	[0.804, 0.972]
DOPE (ours)	$4n$	Linear	94	470	0.991	[0.972, 0.997]
DOPE (ours)	$4n$	MLP	88	440	1.133	[1.029, 1.246]
DOPE, untuned default	n	CatBoost	99	495	0.636	[0.387, 0.825]
DOPE, untuned default	n	Linear	94	470	0.986	[0.974, 0.992]
DOPE, untuned default	n	MLP	88	440	0.887	[0.801, 0.957]
DOPE, untuned default	$4n$	CatBoost	99	495	0.704	[0.424, 0.877]
DOPE, untuned default	$4n$	Linear	94	470	0.993	[0.983, 0.998]
DOPE, untuned default	$4n$	MLP	88	440	0.991	[0.960, 1.058]
TabSyn (scaled schedule)	n	CatBoost	98	98	0.856	[0.696, 0.946]
TabSyn (scaled schedule)	n	Linear	93	93	0.911	[0.763, 0.989]
TabSyn (scaled schedule)	n	MLP	87	87	0.917	[0.741, 0.984]
TabSyn (scaled schedule)	$4n$	CatBoost	98	98	0.897	[0.816, 0.979]
TabSyn (scaled schedule)	$4n$	Linear	93	93	0.958	[0.830, 0.990]
TabSyn (scaled schedule)	$4n$	MLP	87	87	1.031	[0.969, 1.186]
ARF (author default)	n	CatBoost	99	495	0.718	[0.570, 0.844]
ARF (author default)	n	Linear	94	470	0.903	[0.794, 0.965]
ARF (author default)	n	MLP	88	440	0.828	[0.631, 0.886]
ARF (author default)	$4n$	CatBoost	99	495	0.822	[0.689, 0.894]
ARF (author default)	$4n$	Linear	94	470	0.938	[0.846, 0.982]
ARF (author default)	$4n$	MLP	88	440	0.991	[0.886, 1.090]
ARF (native selected)	n	CatBoost	99	495	0.671	[0.551, 0.827]
ARF (native selected)	n	Linear	94	470	0.894	[0.742, 0.966]
ARF (native selected)	n	MLP	88	440	0.805	[0.583, 0.897]
ARF (native selected)	$4n$	CatBoost	99	495	0.804	[0.647, 0.885]
ARF (native selected)	$4n$	Linear	94	470	0.920	[0.828, 0.977]
ARF (native selected)	$4n$	MLP	88	440	0.966	[0.843, 1.027]
Gaussian copula	n	CatBoost	99	99	0.610	[0.310, 0.778]
Gaussian copula	n	Linear	94	94	0.961	[0.915, 0.990]
Gaussian copula	n	MLP	88	88	0.800	[0.618, 0.890]
Gaussian copula	$4n$	CatBoost	99	99	0.669	[0.369, 0.867]
Gaussian copula	$4n$	Linear	94	94	0.987	[0.963, 0.996]
Gaussian copula	$4n$	MLP	88	88	0.978	[0.892, 1.072]
Chow-Liu tree	n	CatBoost	99	99	0.466	[0.246, 0.578]
Chow-Liu tree	n	Linear	94	94	0.636	[0.267, 0.756]
Chow-Liu tree	n	MLP	88	88	0.454	[0.180, 0.631]
Chow-Liu tree	$4n$	CatBoost	99	99	0.576	[0.294, 0.638]
Chow-Liu tree	$4n$	Linear	94	94	0.659	[0.282, 0.763]
Chow-Liu tree	$4n$	MLP	88	88	0.660	[0.347, 0.804]
Independent marginals	n	CatBoost	99	99	-0.046	[-0.083, -0.027]
Independent marginals	n	Linear	94	94	-0.050	[-0.082, -0.015]
Independent marginals	n	MLP	88	88	-0.215	[-0.484, -0.069]
Independent marginals	$4n$	CatBoost	99	99	-0.012	[-0.040, -0.002]
Independent marginals	$4n$	Linear	94	94	0.002	[-0.029, 0.019]
Independent marginals	$4n$	MLP	88	88	-0.049	[-0.161, -0.018]
TVAE (native selected)	n	CatBoost	21	21	0.606	[0.412, 0.739]
TVAE (native selected)	n	Linear	21	21	0.764	[0.552, 0.862]
TVAE (native selected)	n	MLP	20	20	0.721	[0.254, 0.796]
TVAE (native selected)	$4n$	CatBoost	21	21	0.674	[0.452, 0.776]
TVAE (native selected)	$4n$	Linear	21	21	0.796	[0.581, 0.875]
TVAE (native selected)	$4n$	MLP	20	20	0.812	[0.644, 0.877]
CTGAN (native selected)	n	CatBoost	21	21	-0.074	[-0.477, 0.118]
CTGAN (native selected)	n	Linear	21	21	-0.056	[-0.384, 0.118]
CTGAN (native selected)	n	MLP	20	20	-0.247	[-1.072, 0.104]
CTGAN (native selected)	$4n$	CatBoost	21	21	-0.080	[-0.282, 0.223]
CTGAN (native selected)	$4n$	Linear	21	21	-0.033	[-0.454, 0.124]
CTGAN (native selected)	$4n$	MLP	20	20	0.033	[-0.338, 0.328]
Forest-Flow (historical)	n	CatBoost	6	6	0.977	[0.807, 0.999]
Forest-Flow (historical)	n	Linear	5	5	0.981	[0.754, 1.083]
Forest-Flow (historical)	n	MLP	6	6	0.885	[0.578, 1.000]
Forest-Flow (historical)	$4n$	CatBoost	6	6	1.011	[0.905, 1.041]
Forest-Flow (historical)	$4n$	Linear	5	5	0.985	[0.977, 1.001]
Forest-Flow (historical)	$4n$	MLP	6	6	1.054	[0.929, 2.732]
Predictor-only control	n	CatBoost	99	495	0.596	[0.340, 0.726]
Predictor-only control	n	Linear	94	470	0.970	[0.916, 0.994]
Predictor-only control	n	MLP	88	440	0.849	[0.729, 0.913]
Predictor-only control	$4n$	CatBoost	99	495	0.631	[0.375, 0.784]
Predictor-only control	$4n$	Linear	94	470	0.981	[0.955, 0.997]
Predictor-only control	$4n$	MLP	88	440	0.947	[0.883, 1.018]

Arm	Size	Auditor	Lineages	Fits	Median	95% nested interval
Real rows resampled	4n	CatBoost	99	99	1.018	[1.005, 1.028]
Real rows resampled	4n	Linear	94	94	0.998	[0.991, 1.000]
Real rows resampled	4n	MLP	88	88	1.174	[1.051, 1.234]
Chow-Liu tree (default)	n	CatBoost	99	99	0.564	[0.455, 0.677]
Chow-Liu tree (default)	n	Linear	94	94	0.661	[0.608, 0.787]
Chow-Liu tree (default)	n	MLP	88	88	0.523	[0.393, 0.684]
Chow-Liu tree (default)	4n	CatBoost	99	99	0.636	[0.545, 0.754]
Chow-Liu tree (default)	4n	Linear	94	94	0.727	[0.640, 0.810]
Chow-Liu tree (default)	4n	MLP	88	88	0.755	[0.649, 0.949]
CTGAN (author default)	n	CatBoost	17	17	-0.284	[-0.562, -0.054]
CTGAN (author default)	n	Linear	17	17	-0.322	[-1.564, -0.070]
CTGAN (author default)	n	MLP	16	16	-0.399	[-0.817, -0.147]
CTGAN (author default)	4n	CatBoost	17	17	-0.237	[-0.509, -0.022]
CTGAN (author default)	4n	Linear	17	17	-0.294	[-1.606, -0.031]
CTGAN (author default)	4n	MLP	16	16	-0.170	[-0.658, -0.003]
DOPE, later binary, seed 11	n	CatBoost	99	99	0.883	[0.820, 0.970]
DOPE, later binary, seed 11	n	Linear	94	94	0.979	[0.964, 0.994]
DOPE, later binary, seed 11	n	MLP	88	88	0.988	[0.924, 0.995]
DOPE, later binary, seed 11	4n	CatBoost	99	99	0.944	[0.893, 0.985]
DOPE, later binary, seed 11	4n	Linear	94	94	0.994	[0.972, 0.998]
DOPE, later binary, seed 11	4n	MLP	88	88	1.175	[1.044, 1.349]
DOPE, later binary, seeds 23-71	n	CatBoost	99	396	0.791	[0.686, 0.930]
DOPE, later binary, seeds 23-71	n	Linear	94	376	0.984	[0.960, 0.992]
DOPE, later binary, seeds 23-71	n	MLP	88	352	0.958	[0.883, 0.990]
DOPE, later binary, seeds 23-71	4n	CatBoost	99	396	0.862	[0.772, 0.960]
DOPE, later binary, seeds 23-71	4n	Linear	94	376	0.989	[0.970, 0.996]
DOPE, later binary, seeds 23-71	4n	MLP	88	352	1.125	[1.026, 1.229]
DOPE, historical seed 11	n	CatBoost	97	97	0.880	[0.822, 0.971]
DOPE, historical seed 11	n	Linear	92	92	0.984	[0.968, 0.994]
DOPE, historical seed 11	n	MLP	87	87	0.988	[0.923, 0.996]
DOPE, historical seed 11	4n	CatBoost	97	97	0.949	[0.893, 0.987]
DOPE, historical seed 11	4n	Linear	92	92	0.994	[0.972, 0.997]
DOPE, historical seed 11	4n	MLP	87	87	1.174	[1.045, 1.349]
DOPE historical, 12 inputs, 512 steps	n	CatBoost	97	97	0.750	[0.415, 0.923]
DOPE historical, 12 inputs, 512 steps	n	Linear	92	92	0.982	[0.957, 0.990]
DOPE historical, 12 inputs, 512 steps	n	MLP	87	87	0.960	[0.831, 0.988]
DOPE historical, 12 inputs, 512 steps	4n	CatBoost	97	97	0.801	[0.445, 0.958]
DOPE historical, 12 inputs, 512 steps	4n	Linear	92	92	0.984	[0.966, 0.994]
DOPE historical, 12 inputs, 512 steps	4n	MLP	87	87	1.046	[0.998, 1.201]
DOPE historical, 24 inputs, 512 steps	n	CatBoost	93	93	0.762	[0.416, 0.924]
DOPE historical, 24 inputs, 512 steps	n	Linear	88	88	0.982	[0.964, 0.991]
DOPE historical, 24 inputs, 512 steps	n	MLP	84	84	0.964	[0.868, 0.993]
DOPE historical, 24 inputs, 512 steps	4n	CatBoost	93	93	0.802	[0.457, 0.958]
DOPE historical, 24 inputs, 512 steps	4n	Linear	88	88	0.986	[0.968, 0.996]
DOPE historical, 24 inputs, 512 steps	4n	MLP	84	84	1.035	[0.999, 1.204]
DOPE historical, 24 inputs, 2048 steps	n	CatBoost	93	93	0.884	[0.747, 0.971]
DOPE historical, 24 inputs, 2048 steps	n	Linear	88	88	0.977	[0.953, 0.995]
DOPE historical, 24 inputs, 2048 steps	n	MLP	84	84	0.990	[0.949, 0.999]
DOPE historical, 24 inputs, 2048 steps	4n	CatBoost	93	93	0.931	[0.814, 0.985]
DOPE historical, 24 inputs, 2048 steps	4n	Linear	88	88	0.990	[0.967, 0.997]
DOPE historical, 24 inputs, 2048 steps	4n	MLP	84	84	1.124	[1.045, 1.266]
Forest-Flow default (historical)	n	CatBoost	6	6	0.977	[0.807, 0.999]
Forest-Flow default (historical)	n	Linear	5	5	0.981	[0.754, 1.083]
Forest-Flow default (historical)	n	MLP	6	6	0.885	[0.578, 1.000]
Forest-Flow default (historical)	4n	CatBoost	6	6	1.011	[0.905, 1.041]
Forest-Flow default (historical)	4n	Linear	5	5	0.985	[0.977, 1.001]
Forest-Flow default (historical)	4n	MLP	6	6	1.054	[0.929, 2.732]
Gaussian copula (default)	n	CatBoost	99	99	0.565	[0.260, 0.809]
Gaussian copula (default)	n	Linear	94	94	0.921	[0.742, 0.962]
Gaussian copula (default)	n	MLP	88	88	0.786	[0.573, 0.931]
Gaussian copula (default)	4n	CatBoost	99	99	0.605	[0.305, 0.879]
Gaussian copula (default)	4n	Linear	94	94	0.956	[0.878, 0.983]
Gaussian copula (default)	4n	MLP	88	88	0.935	[0.851, 1.033]
Independent marginals (default)	n	CatBoost	99	99	-0.044	[-0.077, -0.028]
Independent marginals (default)	n	Linear	94	94	-0.049	[-0.092, -0.005]
Independent marginals (default)	n	MLP	88	88	-0.184	[-0.449, -0.053]
Independent marginals (default)	4n	CatBoost	99	99	-0.016	[-0.040, -0.006]
Independent marginals (default)	4n	Linear	94	94	0.001	[-0.027, 0.021]
Independent marginals (default)	4n	MLP	88	88	-0.047	[-0.132, -0.018]
Predictor-only, split 2027	n	CatBoost	99	99	0.530	[0.322, 0.729]
Predictor-only, split 2027	n	Linear	97	97	0.990	[0.951, 0.998]
Predictor-only, split 2027	n	MLP	84	84	0.828	[0.731, 0.910]
Predictor-only, split 2027	4n	CatBoost	99	99	0.600	[0.365, 0.778]
Predictor-only, split 2027	4n	Linear	97	97	0.998	[0.987, 1.001]
Predictor-only, split 2027	4n	MLP	84	84	0.950	[0.883, 1.031]
Predictor-only, split 2999	n	CatBoost	100	100	0.570	[0.339, 0.722]
Predictor-only, split 2999	n	Linear	92	92	0.993	[0.972, 0.999]
Predictor-only, split 2999	n	MLP	86	86	0.844	[0.764, 0.937]
Predictor-only, split 2999	4n	CatBoost	100	100	0.647	[0.368, 0.777]

Arm	Size	Auditor	Lineages	Fits	Median	95% nested interval
Predictor-only, split 2999	4n	Linear	92	92	0.998	[0.979, 1.002]
Predictor-only, split 2999	4n	MLP	86	86	0.963	[0.882, 1.065]
Predictor-only, split 4099	n	CatBoost	100	100	0.562	[0.308, 0.716]
Predictor-only, split 4099	n	Linear	95	95	0.987	[0.961, 0.995]
Predictor-only, split 4099	n	MLP	86	86	0.837	[0.732, 0.942]
Predictor-only, split 4099	4n	CatBoost	100	100	0.640	[0.362, 0.777]
Predictor-only, split 4099	4n	Linear	95	95	0.993	[0.974, 0.999]
Predictor-only, split 4099	4n	MLP	86	86	0.965	[0.917, 1.021]
Predictor-only, split 8191	n	CatBoost	100	100	0.611	[0.337, 0.752]
Predictor-only, split 8191	n	Linear	96	96	0.978	[0.960, 0.996]
Predictor-only, split 8191	n	MLP	85	85	0.810	[0.727, 0.903]
Predictor-only, split 8191	4n	CatBoost	100	100	0.677	[0.359, 0.787]
Predictor-only, split 8191	4n	Linear	96	96	0.996	[0.969, 0.999]
Predictor-only, split 8191	4n	MLP	85	85	0.951	[0.884, 1.018]
TVAE (author default)	n	CatBoost	18	18	0.300	[0.072, 0.547]
TVAE (author default)	n	Linear	18	18	0.354	[-0.480, 0.657]
TVAE (author default)	n	MLP	17	17	0.179	[-1.050, 0.569]
TVAE (author default)	4n	CatBoost	18	18	0.283	[0.175, 0.596]
TVAE (author default)	4n	Linear	18	18	0.429	[0.156, 0.682]
TVAE (author default)	4n	MLP	17	17	0.333	[-0.417, 0.612]

Table 11: **All DOPE-minus-comparator contrasts.** Nested bootstrap 95% intervals. The Holm family size is in parentheses; families are declared per size in predeclare.v2. HL eq. is the pre-declared equivalence rule (90% interval of the Hodges–Lehmann difference inside ± 0.02); Yuen eq. is the trimmed-mean sensitivity.

Size	Auditor	Comparator	n	Median diff.	Hodges–Lehmann	W/T/L	Holm p (family)	HL eq.	Yuen eq.
n	CatBoost	TabSyn (scaled schedule)	98	0.000 [-0.040, 0.033]	-0.003 [-0.053, 0.050]	48/0/50	1.000 (12)	no	no
n	Linear	TabSyn (scaled schedule)	93	0.030 [0.006, 0.155]	0.107 [0.012, 0.273]	70/0/23	6.39×10^{-8} (12)	no	no
n	MLP	TabSyn (scaled schedule)	87	0.041 [-0.002, 0.141]	0.073 [0.003, 0.214]	56/0/31	0.003 (12)	no	no
n	CatBoost	ARF (author default)	99	0.084 [0.004, 0.165]	0.090 [0.006, 0.190]	71/0/28	0.001 (12)	no	no
n	Linear	ARF (author default)	94	0.065 [0.003, 0.190]	0.090 [0.008, 0.251]	74/0/20	3.53×10^{-6} (12)	no	no
n	MLP	ARF (author default)	88	0.111 [0.050, 0.226]	0.144 [0.061, 0.332]	74/0/14	6.04×10^{-8} (12)	no	no
n	CatBoost	ARF (native selected)	99	0.083 [0.009, 0.201]	0.109 [0.025, 0.220]	71/0/28	1.71×10^{-4} (9)	no	no
n	Linear	ARF (native selected)	94	0.077 [0.010, 0.198]	0.104 [0.019, 0.256]	76/0/18	1.10×10^{-7} (9)	no	no
n	MLP	ARF (native selected)	88	0.142 [0.063, 0.274]	0.175 [0.081, 0.363]	76/0/12	3.06×10^{-10} (9)	no	no
n	CatBoost	Gaussian copula	99	0.196 [0.090, 0.243]	0.202 [0.110, 0.266]	88/0/11	1.41×10^{-14} (9)	no	no
n	Linear	Gaussian copula	94	0.008 [-0.005, 0.038]	0.020 [-0.005, 0.049]	59/0/35	0.010 (9)	no	no
n	MLP	Gaussian copula	88	0.132 [0.090, 0.185]	0.157 [0.097, 0.242]	73/0/15	2.68×10^{-9} (9)	no	no
n	CatBoost	Chow-Liu tree	99	0.313 [0.240, 0.513]	0.369 [0.268, 0.660]	87/0/12	1.24×10^{-15} (9)	no	no
n	Linear	Chow-Liu tree	94	0.342 [0.226, 0.624]	0.481 [0.282, 0.819]	91/0/3	1.97×10^{-15} (9)	no	no
n	MLP	Chow-Liu tree	88	0.489 [0.340, 0.755]	0.620 [0.406, 1.079]	83/0/5	4.89×10^{-14} (9)	no	no
n	CatBoost	Independent marginals	99	0.866 [0.767, 0.982]	0.849 [0.745, 0.964]	99/0/0	5.13×10^{-17} (9)	no	no
n	Linear	Independent marginals	94	1.002 [0.974, 1.065]	1.023 [0.975, 1.097]	94/0/0	3.05×10^{-16} (9)	no	no
n	MLP	Independent marginals	88	1.164 [1.039, 1.473]	1.265 [1.078, 1.549]	87/0/1	1.17×10^{-14} (9)	no	no
n	CatBoost	TVAE (native selected)	21	0.195 [0.089, 0.395]	0.221 [0.115, 0.347]	19/0/2	—	no	no
n	Linear	TVAE (native selected)	21	0.250 [0.128, 0.427]	0.284 [0.158, 0.549]	20/0/1	—	no	no
n	MLP	TVAE (native selected)	20	0.221 [0.109, 0.507]	0.254 [0.118, 0.596]	17/0/3	—	no	no
n	CatBoost	CTGAN (native selected)	21	0.785 [0.706, 1.342]	0.880 [0.731, 1.204]	21/0/0	—	no	no
n	Linear	CTGAN (native selected)	21	1.052 [0.849, 1.342]	1.104 [0.833, 1.732]	21/0/0	—	no	no
n	MLP	CTGAN (native selected)	20	1.135 [0.772, 1.776]	1.295 [0.789, 2.228]	19/0/1	—	no	no

Size	Auditor	Comparator	n	Median diff.	Hodges-Lehmann	W/T/L	Holm p (family)	HL eq.	Yuen eq.
n	CatBoost	Forest-Flow (historical)	6	-0.045 [-0.463, 0.020]	-0.206 [-0.385, 0.003]	1/0/5	—	no	no
n	Linear	Forest-Flow (historical)	5	0.024 [-0.005, 0.074]	0.037 [-0.034, 0.074]	5/0/0	—	no	no
n	MLP	Forest-Flow (historical)	6	0.005 [-0.041, 0.123]	0.010 [-0.086, 0.123]	3/0/3	—	no	no
n	CatBoost	DOPE, untuned default	99	0.101 [0.035, 0.155]	0.127 [0.052, 0.186]	89/0/10	—	no	no
n	Linear	DOPE, untuned default	94	-0.002 [-0.016, 0.003]	-0.007 [-0.027, 0.002]	42/0/52	—	no	yes
n	MLP	DOPE, untuned default	88	0.034 [0.003, 0.069]	0.050 [0.008, 0.082]	63/0/25	—	no	no
n	CatBoost	Predictor-only control	99	0.215 [0.128, 0.271]	0.214 [0.148, 0.269]	93/0/6	4.36×10^{-15} (6)	no	no
n	Linear	Predictor-only control	94	0.003 [-0.014, 0.041]	0.016 [-0.022, 0.054]	55/0/39	0.142 (6)	no	no
n	MLP	Predictor-only control	88	0.106 [0.044, 0.168]	0.134 [0.067, 0.195]	73/0/15	1.96×10^{-8} (6)	no	no
n	CatBoost	Chow-Liu tree (default)	99	0.239 [0.175, 0.313]	0.261 [0.185, 0.352]	86/0/13	—	no	no
n	Linear	Chow-Liu tree (default)	94	0.282 [0.177, 0.346]	0.327 [0.202, 0.481]	90/0/4	—	no	no
n	MLP	Chow-Liu tree (default)	88	0.389 [0.242, 0.502]	0.490 [0.311, 0.769]	82/0/6	—	no	no
n	CatBoost	CTGAN (author default)	17	1.093 [0.787, 1.446]	1.108 [0.861, 1.707]	17/0/0	1.53×10^{-4} (12)	no	no
n	Linear	CTGAN (author default)	17	1.314 [0.981, 2.569]	1.610 [1.151, 2.510]	17/0/0	1.53×10^{-4} (12)	no	no
n	MLP	CTGAN (author default)	16	1.314 [0.891, 1.790]	1.378 [1.030, 2.528]	15/0/1	0.001 (12)	no	no
n	CatBoost	DOPE, later binary, seed 11	99	-0.020 [-0.058, -0.003]	-0.033 [-0.074, -0.009]	14/10/75	—	no	no
n	Linear	DOPE, later binary, seed 11	94	0.000 [-0.001, 0.003]	0.002 [-0.006, 0.013]	40/19/35	—	yes	yes
n	MLP	DOPE, later binary, seed 11	88	-0.003 [-0.010, 0.000]	-0.007 [-0.021, 0.000]	25/16/47	—	yes	yes
n	CatBoost	DOPE, later binary, seeds 23-71	99	0.008 [0.000, 0.011]	0.010 [0.000, 0.020]	79/0/20	—	yes	yes
n	Linear	DOPE, later binary, seeds 23-71	94	0.000 [-0.002, 0.002]	0.000 [-0.005, 0.006]	44/0/50	—	yes	yes
n	MLP	DOPE, later binary, seeds 23-71	88	0.003 [-0.001, 0.006]	0.003 [-0.005, 0.012]	54/0/34	—	yes	yes
n	CatBoost	DOPE, historical seed 11	97	-0.022 [-0.060, -0.003]	-0.034 [-0.076, -0.009]	15/8/74	—	no	no
n	Linear	DOPE, historical seed 11	92	0.000 [-0.002, 0.002]	0.001 [-0.006, 0.011]	38/16/38	—	yes	yes
n	MLP	DOPE, historical seed 11	87	-0.003 [-0.011, 0.000]	-0.007 [-0.023, 0.000]	25/15/47	—	yes	yes
n	CatBoost	DOPE historical, 12 inputs, 512 steps	97	0.032 [0.006, 0.092]	0.060 [0.012, 0.143]	72/0/25	—	no	no
n	Linear	DOPE historical, 12 inputs, 512 steps	92	0.002 [-0.006, 0.007]	0.001 [-0.010, 0.012]	49/0/43	—	yes	yes
n	MLP	DOPE historical, 12 inputs, 512 steps	87	0.007 [-0.004, 0.016]	0.015 [-0.005, 0.038]	50/0/37	—	no	no
n	CatBoost	DOPE historical, 24 inputs, 512 steps	93	0.032 [0.006, 0.087]	0.053 [0.012, 0.126]	69/0/24	—	no	no
n	Linear	DOPE historical, 24 inputs, 512 steps	88	0.003 [-0.005, 0.008]	0.003 [-0.009, 0.012]	48/0/40	—	yes	yes
n	MLP	DOPE historical, 24 inputs, 512 steps	84	0.002 [-0.015, 0.014]	0.010 [-0.015, 0.029]	45/0/39	—	no	no
n	CatBoost	DOPE historical, 24 inputs, 2048 steps	93	-0.013 [-0.030, -0.001]	-0.019 [-0.041, -0.006]	24/6/63	—	no	no
n	Linear	DOPE historical, 24 inputs, 2048 steps	88	0.000 [-0.005, 0.001]	0.000 [-0.012, 0.012]	35/14/39	—	yes	yes
n	MLP	DOPE historical, 24 inputs, 2048 steps	84	-0.005 [-0.020, 0.000]	-0.014 [-0.040, -0.003]	21/13/50	—	no	no
n	CatBoost	Forest-Flow default (historical)	6	-0.045 [-0.463, 0.020]	-0.206 [-0.366, 0.003]	1/0/5	—	no	no
n	Linear	Forest-Flow default (historical)	5	0.024 [-0.005, 0.074]	0.037 [-0.043, 0.074]	5/0/0	—	no	no
n	MLP	Forest-Flow default (historical)	6	0.005 [-0.041, 0.123]	0.010 [-0.074, 0.123]	3/0/3	—	no	no
n	CatBoost	Gaussian copula (default)	99	0.222 [0.092, 0.283]	0.229 [0.121, 0.303]	89/0/10	—	no	no
n	Linear	Gaussian copula (default)	94	0.048 [0.013, 0.150]	0.104 [0.025, 0.224]	71/0/23	—	no	no
n	MLP	Gaussian copula (default)	88	0.129 [0.069, 0.231]	0.178 [0.085, 0.309]	74/0/14	—	no	no

Size	Auditor	Comparator	n	Median diff.	Hodges–Lehmann	W/T/L	Holm p (family)	HL eq.	Yuen eq.
n	CatBoost	Independent marginals (default)	99	0.875 [0.757, 0.981]	0.842 [0.739, 0.962]	99/0/0	—	no	no
n	Linear	Independent marginals (default)	94	1.011 [0.969, 1.061]	1.017 [0.965, 1.089]	93/0/1	—	no	no
n	MLP	Independent marginals (default)	88	1.121 [1.039, 1.441]	1.229 [1.061, 1.449]	86/0/2	—	no	no
n	CatBoost	Predictor-only, split 2027	98	0.214 [0.147, 0.281]	0.230 [0.160, 0.301]	92/0/6	—	no	no
n	Linear	Predictor-only, split 2027	93	−0.006 [−0.032, 0.009]	−0.010 [−0.043, 0.012]	40/0/53	—	no	no
n	MLP	Predictor-only, split 2027	78	0.104 [0.020, 0.162]	0.128 [0.040, 0.198]	61/0/17	—	no	no
n	CatBoost	Predictor-only, split 2999	99	0.200 [0.111, 0.263]	0.216 [0.137, 0.287]	86/0/13	—	no	no
n	Linear	Predictor-only, split 2999	89	−0.007 [−0.023, 0.008]	−0.006 [−0.034, 0.015]	35/0/54	—	no	yes
n	MLP	Predictor-only, split 2999	82	0.074 [0.011, 0.159]	0.121 [0.030, 0.191]	61/0/21	—	no	no
n	CatBoost	Predictor-only, split 4099	99	0.223 [0.126, 0.275]	0.226 [0.152, 0.299]	86/0/13	—	no	no
n	Linear	Predictor-only, split 4099	91	−0.005 [−0.016, 0.012]	−0.003 [−0.033, 0.025]	40/0/51	—	no	yes
n	MLP	Predictor-only, split 4099	84	0.098 [0.005, 0.159]	0.119 [0.041, 0.197]	58/0/26	—	no	no
n	CatBoost	Predictor-only, split 8191	99	0.179 [0.091, 0.252]	0.199 [0.121, 0.263]	85/0/14	—	no	no
n	Linear	Predictor-only, split 8191	91	0.004 [−0.010, 0.027]	0.007 [−0.016, 0.037]	50/0/41	—	no	yes
n	MLP	Predictor-only, split 8191	81	0.101 [0.024, 0.170]	0.133 [0.044, 0.201]	67/0/14	—	no	no
n	CatBoost	TVAE (author default)	18	0.452 [0.329, 0.778]	0.504 [0.365, 0.747]	18/0/0	9.16×10^{-5} (12)	no	no
n	Linear	TVAE (author default)	18	0.593 [0.326, 1.897]	0.628 [0.432, 12.069]	18/0/0	9.16×10^{-5} (12)	no	no
n	MLP	TVAE (author default)	17	0.632 [0.270, 1.242]	0.748 [0.378, 1.925]	16/0/1	2.44×10^{-4} (12)	no	no
$4n$	CatBoost	TabSyn (scaled schedule)	98	−0.003 [−0.060, 0.039]	−0.008 [−0.057, 0.042]	47/0/51	1.000 (12)	no	no
$4n$	Linear	TabSyn (scaled schedule)	93	0.019 [0.002, 0.130]	0.072 [0.004, 0.248]	70/0/23	6.36×10^{-6} (12)	no	no
$4n$	MLP	TabSyn (scaled schedule)	87	0.000 [−0.013, 0.166]	0.069 [−0.014, 0.392]	43/0/44	0.341 (12)	no	no
$4n$	CatBoost	ARF (author default)	99	0.035 [−0.021, 0.130]	0.048 [−0.034, 0.149]	60/0/39	0.258 (12)	no	no
$4n$	Linear	ARF (author default)	94	0.045 [−0.002, 0.155]	0.062 [−0.011, 0.193]	70/0/24	2.87×10^{-4} (12)	no	no
$4n$	MLP	ARF (author default)	88	0.117 [0.020, 0.237]	0.137 [0.041, 0.397]	70/0/18	3.26×10^{-7} (12)	no	no
$4n$	CatBoost	ARF (native selected)	99	0.051 [−0.005, 0.162]	0.083 [0.006, 0.191]	66/0/33	0.004 (9)	no	no
$4n$	Linear	ARF (native selected)	94	0.073 [0.003, 0.164]	0.083 [0.006, 0.217]	73/0/21	1.89×10^{-6} (9)	no	no
$4n$	MLP	ARF (native selected)	88	0.106 [0.034, 0.306]	0.191 [0.060, 0.431]	71/0/17	2.48×10^{-8} (9)	no	no
$4n$	CatBoost	Gaussian copula	99	0.164 [0.062, 0.242]	0.191 [0.089, 0.273]	83/0/16	1.46×10^{-13} (9)	no	no
$4n$	Linear	Gaussian copula	94	0.002 [−0.014, 0.015]	0.007 [−0.023, 0.030]	53/0/41	0.183 (9)	no	yes
$4n$	MLP	Gaussian copula	88	0.170 [0.047, 0.264]	0.234 [0.126, 0.390]	77/0/11	9.26×10^{-13} (9)	no	no
$4n$	CatBoost	Chow-Liu tree	99	0.270 [0.201, 0.517]	0.356 [0.242, 0.728]	91/0/8	5.59×10^{-15} (9)	no	no
$4n$	Linear	Chow-Liu tree	94	0.333 [0.209, 0.659]	0.495 [0.266, 0.816]	91/0/3	1.91×10^{-15} (9)	no	no
$4n$	MLP	Chow-Liu tree	88	0.566 [0.353, 0.795]	0.658 [0.435, 1.166]	88/0/0	2.24×10^{-15} (9)	no	no
$4n$	CatBoost	Independent marginals	99	0.913 [0.826, 0.978]	0.881 [0.790, 0.964]	99/0/0	5.13×10^{-17} (9)	no	no
$4n$	Linear	Independent marginals	94	0.982 [0.937, 1.014]	0.982 [0.926, 1.008]	94/0/0	3.05×10^{-16} (9)	no	no
$4n$	MLP	Independent marginals	88	1.220 [1.094, 1.467]	1.327 [1.143, 1.629]	88/0/0	2.24×10^{-15} (9)	no	no
$4n$	CatBoost	TVAE (native selected)	21	0.201 [0.080, 0.351]	0.201 [0.097, 0.339]	19/0/2	—	no	no
$4n$	Linear	TVAE (native selected)	21	0.213 [0.130, 0.414]	0.262 [0.145, 0.473]	20/0/1	—	no	no
$4n$	MLP	TVAE (native selected)	20	0.286 [0.141, 0.694]	0.426 [0.189, 1.231]	19/0/1	—	no	no

Size	Auditor	Comparator	n	Median diff.	Hodges-Lehmann	W/T/L	Holm p (family)	HL eq.	Yuen eq.
4n	CatBoost	CTGAN (native selected)	21	0.792 [0.705, 1.302]	0.912 [0.731, 1.298]	21/0/0	—	no	no
4n	Linear	CTGAN (native selected)	21	0.983 [0.843, 1.426]	1.066 [0.812, 1.695]	21/0/0	—	no	no
4n	MLP	CTGAN (native selected)	20	1.086 [0.707, 2.268]	1.439 [0.788, 2.272]	20/0/0	—	no	no
4n	CatBoost	Forest-Flow (historical)	6	−0.035 [−0.477, −0.004]	−0.231 [−0.382, −0.005]	0/0/6	—	no	no
4n	Linear	Forest-Flow (historical)	5	0.000 [−0.060, 0.037]	0.002 [−0.051, 0.117]	3/0/2	—	no	no
4n	MLP	Forest-Flow (historical)	6	−0.027 [−0.411, −0.001]	−0.080 [−0.562, 0.004]	1/0/5	—	no	no
4n	CatBoost	DOPE, untuned default	99	0.110 [0.033, 0.166]	0.134 [0.053, 0.208]	87/0/12	—	no	no
4n	Linear	DOPE, untuned default	94	−0.001 [−0.017, 0.006]	−0.003 [−0.026, 0.005]	43/0/51	—	no	yes
4n	MLP	DOPE, untuned default	88	0.080 [0.028, 0.132]	0.138 [0.052, 0.251]	69/0/19	—	no	no
4n	CatBoost	Predictor-only control	99	0.194 [0.069, 0.252]	0.197 [0.112, 0.267]	92/0/7	2.80×10^{-14} (6)	no	no
4n	Linear	Predictor-only control	94	0.001 [−0.015, 0.031]	0.008 [−0.027, 0.046]	51/0/43	0.230 (6)	no	yes
4n	MLP	Predictor-only control	88	0.165 [0.077, 0.248]	0.227 [0.118, 0.378]	79/0/9	4.73×10^{-13} (6)	no	no
4n	CatBoost	Real rows resampled	99	−0.132 [−0.222, −0.050]	−0.179 [−0.276, −0.074]	10/0/89	2.02×10^{-13} (6)	no	no
4n	Linear	Real rows resampled	94	−0.008 [−0.023, 0.003]	−0.008 [−0.046, 0.008]	33/0/61	0.050 (6)	no	yes
4n	MLP	Real rows resampled	88	−0.032 [−0.067, 0.000]	−0.037 [−0.101, 0.070]	25/0/63	0.046 (6)	no	no
4n	CatBoost	Chow-Liu tree (default)	99	0.211 [0.096, 0.283]	0.230 [0.125, 0.348]	82/0/17	—	no	no
4n	Linear	Chow-Liu tree (default)	94	0.259 [0.136, 0.355]	0.317 [0.183, 0.466]	87/0/7	—	no	no
4n	MLP	Chow-Liu tree (default)	88	0.394 [0.209, 0.570]	0.489 [0.284, 0.799]	84/0/4	—	no	no
4n	CatBoost	CTGAN (author default)	17	0.978 [0.810, 1.489]	1.097 [0.863, 1.760]	17/0/0	1.53×10^{-4} (12)	no	no
4n	Linear	CTGAN (author default)	17	1.277 [0.987, 2.628]	1.463 [1.093, 2.148]	17/0/0	1.53×10^{-4} (12)	no	no
4n	MLP	CTGAN (author default)	16	1.454 [1.030, 2.301]	1.541 [1.148, 2.552]	16/0/0	2.14×10^{-4} (12)	no	no
4n	CatBoost	DOPE, later binary, seed 11	99	−0.021 [−0.061, 0.000]	−0.037 [−0.079, −0.008]	19/15/65	—	no	no
4n	Linear	DOPE, later binary, seed 11	94	0.000 [−0.002, 0.000]	0.000 [−0.010, 0.007]	39/16/39	—	yes	yes
4n	MLP	DOPE, later binary, seed 11	88	−0.012 [−0.018, 0.000]	−0.020 [−0.042, −0.002]	24/12/52	—	no	no
4n	CatBoost	DOPE, later binary, seeds 23-71	99	0.007 [0.000, 0.012]	0.011 [0.000, 0.021]	76/0/23	—	yes	yes
4n	Linear	DOPE, later binary, seeds 23-71	94	0.000 [−0.001, 0.003]	0.000 [−0.003, 0.007]	45/0/49	—	yes	yes
4n	MLP	DOPE, later binary, seeds 23-71	88	0.005 [0.000, 0.008]	0.007 [−0.004, 0.019]	56/0/32	—	yes	yes
4n	CatBoost	DOPE, historical seed 11	97	−0.028 [−0.061, 0.000]	−0.037 [−0.081, −0.008]	20/12/65	—	no	no
4n	Linear	DOPE, historical seed 11	92	0.000 [−0.003, 0.001]	0.000 [−0.011, 0.006]	38/14/40	—	yes	yes
4n	MLP	DOPE, historical seed 11	87	−0.012 [−0.021, 0.000]	−0.020 [−0.042, −0.001]	24/11/52	—	no	no
4n	CatBoost	DOPE historical, 12 inputs, 512 steps	97	0.036 [0.008, 0.107]	0.061 [0.013, 0.162]	73/0/24	—	no	no
4n	Linear	DOPE historical, 12 inputs, 512 steps	92	0.001 [−0.005, 0.007]	0.003 [−0.008, 0.013]	49/0/43	—	yes	yes
4n	MLP	DOPE historical, 12 inputs, 512 steps	87	0.025 [0.004, 0.071]	0.058 [0.015, 0.142]	62/0/25	—	no	no
4n	CatBoost	DOPE historical, 24 inputs, 512 steps	93	0.033 [0.006, 0.075]	0.051 [0.013, 0.137]	70/0/23	—	no	no
4n	Linear	DOPE historical, 24 inputs, 512 steps	88	0.001 [−0.005, 0.009]	0.003 [−0.010, 0.016]	48/0/40	—	yes	yes
4n	MLP	DOPE historical, 24 inputs, 512 steps	84	0.024 [0.002, 0.064]	0.047 [0.005, 0.124]	57/0/27	—	no	no
4n	CatBoost	DOPE historical, 24 inputs, 2048 steps	93	−0.012 [−0.033, 0.000]	−0.020 [−0.045, −0.004]	27/11/55	—	no	no
4n	Linear	DOPE historical, 24 inputs, 2048 steps	88	0.000 [−0.006, 0.004]	0.000 [−0.015, 0.012]	38/11/39	—	yes	yes
4n	MLP	DOPE historical, 24 inputs, 2048 steps	84	−0.011 [−0.033, 0.000]	−0.023 [−0.050, 0.000]	24/9/51	—	no	no

Size	Auditor	Comparator	n	Median diff.	Hodges-Lehmann	W/T/L	Holm p (family)	HL eq.	Yuen eq.
4n	CatBoost	Forest-Flow default (historical)	6	-0.035 [-0.477, -0.004]	-0.231 [-0.408, -0.005]	0/0/6	—	no	no
4n	Linear	Forest-Flow default (historical)	5	0.000 [-0.033, 0.037]	0.002 [-0.051, 0.117]	3/0/2	—	no	no
4n	MLP	Forest-Flow default (historical)	6	-0.027 [-0.411, -0.001]	-0.080 [-0.562, 0.004]	1/0/5	—	no	no
4n	CatBoost	Gaussian copula (default)	99	0.201 [0.062, 0.264]	0.215 [0.091, 0.302]	84/0/15	—	no	no
4n	Linear	Gaussian copula (default)	94	0.034 [0.003, 0.079]	0.061 [0.008, 0.169]	68/0/26	—	no	no
4n	MLP	Gaussian copula (default)	88	0.185 [0.100, 0.351]	0.299 [0.144, 0.618]	78/0/10	—	no	no
4n	CatBoost	Independent marginals (default)	99	0.917 [0.819, 0.975]	0.884 [0.790, 0.966]	99/0/0	—	no	no
4n	Linear	Independent marginals (default)	94	0.978 [0.947, 1.012]	0.982 [0.924, 1.008]	92/0/2	—	no	no
4n	MLP	Independent marginals (default)	88	1.188 [1.056, 1.411]	1.305 [1.121, 1.580]	88/0/0	—	no	no
4n	CatBoost	Predictor-only, split 2027	98	0.215 [0.097, 0.296]	0.229 [0.141, 0.300]	92/0/6	—	no	no
4n	Linear	Predictor-only, split 2027	93	-0.009 [-0.034, 0.002]	-0.015 [-0.060, 0.001]	35/0/58	—	no	no
4n	MLP	Predictor-only, split 2027	78	0.103 [0.034, 0.205]	0.194 [0.067, 0.362]	60/0/18	—	no	no
4n	CatBoost	Predictor-only, split 2999	99	0.177 [0.062, 0.266]	0.198 [0.107, 0.280]	83/0/16	—	no	no
4n	Linear	Predictor-only, split 2999	89	-0.007 [-0.031, 0.009]	-0.010 [-0.060, 0.013]	34/0/55	—	no	no
4n	MLP	Predictor-only, split 2999	82	0.144 [0.045, 0.263]	0.202 [0.075, 0.413]	64/0/18	—	no	no
4n	CatBoost	Predictor-only, split 4099	99	0.188 [0.088, 0.262]	0.209 [0.130, 0.291]	89/0/10	—	no	no
4n	Linear	Predictor-only, split 4099	91	-0.003 [-0.018, 0.014]	-0.002 [-0.028, 0.025]	40/0/51	—	no	yes
4n	MLP	Predictor-only, split 4099	84	0.113 [0.011, 0.199]	0.154 [0.041, 0.290]	61/0/23	—	no	no
4n	CatBoost	Predictor-only, split 8191	99	0.172 [0.086, 0.245]	0.190 [0.099, 0.269]	85/0/14	—	no	no
4n	Linear	Predictor-only, split 8191	91	-0.004 [-0.020, 0.011]	-0.004 [-0.031, 0.022]	39/0/52	—	no	yes
4n	MLP	Predictor-only, split 8191	81	0.122 [0.040, 0.219]	0.178 [0.064, 0.319]	67/0/14	—	no	no
4n	CatBoost	TVAE (author default)	18	0.470 [0.299, 0.785]	0.505 [0.357, 0.739]	18/0/0	9.16×10^{-5} (12)	no	no
4n	Linear	TVAE (author default)	18	0.562 [0.321, 0.832]	0.575 [0.393, 7.242]	18/0/0	9.16×10^{-5} (12)	no	no
4n	MLP	TVAE (author default)	17	0.761 [0.375, 1.880]	0.932 [0.557, 2.021]	17/0/0	1.53×10^{-4} (12)	no	no

E.1 FIT-SEED-11 REVIEW PANELS

The tables below are the fit-seed-11 review panels that preceded the five-seed refit. A method with two configurations keeps one row per configuration. Their intervals are the family-cluster bootstrap of the median, their equivalence column is the v1 mean-difference test, and their Holm families are split by size: nine density tests, six neural tests, six ARF tests, and three Forest-Flow tests per size. Forest-Flow cells store no null loss, so its loss-path columns in Table 18 are empty while its informative count matches the size table. The sensitivity column of that table is the capped cohort, a different denominator from the headline.

Table 13: **DOPE minus comparator at both sizes, fit seed 11.** Median paired retention difference with its family-cluster bootstrap 95% interval, the mean difference, and the v1 mean-difference equivalence test at ± 0.02 . W/T/L counts lineages. The Wilcoxon and equivalence tests treat lineages as independent. Holm families are separate per size: density 9, neural 6, ARF 6, and Forest-Flow 3 tests.

Size	Auditor	Comparator	Configuration	n	Median diff.	Hierarchical CI	Mean	TOST	W/T/L	Holm p
1n	catboost	GaussianCopula	native_selected	97	0.2178	[0.1028, 0.3340]	0.2483	no	90/0/7	0.000
1n	catboost	Chow-Liu	native_selected	97	0.3635	[0.2986, 0.5314]	7.7460	no	93/0/4	0.000
1n	catboost	independent_marginals	native_selected	97	0.9238	[0.8575, 1.0178]	1.1580	no	97/0/0	0.000
1n	linear	GaussianCopula	native_selected	92	0.0088	[-0.0118, 0.0329]	0.2382	no	53/0/39	0.071
1n	linear	Chow-Liu	native_selected	92	0.3309	[0.2275, 0.6482]	1.5301	no	86/0/6	0.000
1n	linear	independent_marginals	native_selected	92	0.9955	[0.9691, 1.0693]	1.1054	no	92/0/0	0.000

Size	Auditor	Comparator	Configuration	n	Median diff.	Hierarchical CI	Mean	TOST	W/T/L	Holm p
1n	mlp	GaussianCopula	native_selected	75	0.1453	[0.0479, 0.2445]	0.1971	no	64/0/11	0.000
1n	mlp	Chow-Liu	native_selected	75	0.5339	[0.3343, 0.8171]	1.1319	no	75/0/0	0.000
1n	mlp	independent_marginals	native_selected	75	1.0462	[0.9834, 1.2772]	1.3698	no	75/0/0	0.000
4n	catboost	GaussianCopula	native_selected	97	0.2237	[0.0723, 0.3530]	0.2613	no	86/0/11	0.000
4n	catboost	Chow-Liu	native_selected	97	0.3370	[0.2823, 0.5048]	8.0937	no	94/0/3	0.000
4n	catboost	independent_marginals	native_selected	97	0.9556	[0.8911, 0.9875]	1.2876	no	96/0/1	0.000
4n	linear	GaussianCopula	native_selected	92	0.0037	[-0.0144, 0.0138]	0.1832	no	49/0/43	0.373
4n	linear	Chow-Liu	native_selected	92	0.3399	[0.1947, 0.6595]	1.4415	no	86/0/6	0.000
4n	linear	independent_marginals	native_selected	92	0.9879	[0.9427, 1.0142]	1.0440	no	92/0/0	0.000
4n	mlp	GaussianCopula	native_selected	75	0.1850	[0.0811, 0.2624]	0.3295	no	65/0/10	0.000
4n	mlp	Chow-Liu	native_selected	75	0.5301	[0.3183, 0.8832]	0.9903	no	74/0/1	0.000
4n	mlp	independent_marginals	native_selected	75	1.1043	[1.0367, 1.3283]	1.3100	no	75/0/0	0.000
1n	catboost	CTGAN	native_selected	21	0.8816	[0.7166, 1.4415]	1.0872	no	21/0/0	0.000
1n	catboost	TVAE	native_selected	21	0.2381	[0.1133, 0.4819]	0.2867	no	19/0/2	0.000
1n	linear	CTGAN	native_selected	21	1.0516	[0.8722, 1.3740]	1.4244	no	21/0/0	0.000
1n	linear	TVAE	native_selected	21	0.2510	[0.1501, 0.4434]	1.3010	no	20/0/1	0.000
1n	mlp	CTGAN	native_selected	19	0.9336	[0.6818, 2.0568]	2.1780	no	19/0/0	0.000
1n	mlp	TVAE	native_selected	19	0.2725	[0.1706, 0.7092]	0.5985	no	18/0/1	0.000
4n	catboost	CTGAN	native_selected	21	0.8807	[0.7020, 1.3326]	1.0962	no	21/0/0	0.000
4n	catboost	TVAE	native_selected	21	0.2361	[0.0868, 0.4471]	0.2673	no	19/0/2	0.000
4n	linear	CTGAN	native_selected	21	1.0485	[0.8718, 1.4642]	1.3489	no	21/0/0	0.000
4n	linear	TVAE	native_selected	21	0.2629	[0.1298, 0.4183]	1.0124	no	20/0/1	0.000
4n	mlp	CTGAN	native_selected	19	1.0788	[0.8605, 2.0296]	1.8412	no	19/0/0	0.000
4n	mlp	TVAE	native_selected	19	0.2842	[0.1227, 0.8120]	0.7771	no	17/0/2	0.000
1n	catboost	ARF	author_default	97	0.1086	[0.0026, 0.2394]	0.1144	no	73/0/24	0.000
1n	catboost	ARF	native_selected	97	0.1010	[0.0236, 0.2404]	0.1354	no	76/0/21	0.000
1n	linear	ARF	author_default	92	0.0409	[0.0049, 0.1752]	0.1786	no	64/0/28	0.000
1n	linear	ARF	native_selected	92	0.0649	[0.0059, 0.2012]	0.2123	no	68/0/24	0.000
1n	mlp	ARF	author_default	75	0.1371	[0.0311, 0.2860]	0.2632	no	61/0/14	0.000
1n	mlp	ARF	native_selected	75	0.1423	[0.0464, 0.2643]	0.3097	no	64/0/11	0.000
4n	catboost	ARF	author_default	97	0.0767	[0.0103, 0.1957]	0.2059	no	69/0/28	0.000
4n	catboost	ARF	native_selected	97	0.0954	[0.0166, 0.2137]	0.1737	no	70/0/27	0.000
4n	linear	ARF	author_default	92	0.0274	[-0.0020, 0.1195]	0.0924	no	61/0/31	0.000
4n	linear	ARF	native_selected	92	0.0369	[0.0008, 0.1639]	0.1249	no	65/0/27	0.000
4n	mlp	ARF	author_default	75	0.1153	[0.0340, 0.3709]	0.2894	no	57/0/18	0.000
4n	mlp	ARF	native_selected	75	0.1429	[0.0283, 0.3673]	0.3089	no	58/0/17	0.000
1n	catboost	ForestDiffusion/Forest-Flow	native_selected	6	-0.0198	[-0.4355, 0.0060]	-0.0917	no	1/0/5	0.469
1n	linear	ForestDiffusion/Forest-Flow	native_selected	5	0.0275	[-0.1396, 0.0723]	-0.0040	no	3/0/2	1.000
1n	mlp	ForestDiffusion/Forest-Flow	native_selected	4	0.0196	[-0.0405, 0.2906]	0.0723	no	2/0/2	1.000
4n	catboost	ForestDiffusion/Forest-Flow	native_selected	6	-0.0269	[-0.4550, 0.0080]	-0.1089	no	1/0/5	0.375
4n	linear	ForestDiffusion/Forest-Flow	native_selected	5	0.0115	[-0.2599, 0.3016]	0.0118	no	3/0/2	0.812
4n	mlp	ForestDiffusion/Forest-Flow	native_selected	4	-0.0317	[-0.2275, -0.0102]	-0.0753	no	0/0/4	0.375

Table 14: **TabSyn and DOPE retention levels, fit seed 11.** TabSyn runs the scaled schedule (200 VAE and 1,000 diffusion epochs). A lineage median needs three informative sample seeds at size n or $4n$; own-lineage rows use every informative TabSyn lineage, and matched rows use the lineages shared with informative DOPE features12_steps2048. Intervals are the family-cluster bootstrap. Lineages absent from the admitted cohort: 02a45777900441ce; absence is no win. No Holm family.

Method	Configuration	Cohort	Size	Auditor	n lineages	Median	Hierarchical CI
TabSyn	scaled_200_vae_1000_diffusion	own lineages	1n	catboost	98	0.8557	[0.6955, 0.9459]
TabSyn	scaled_200_vae_1000_diffusion	own lineages	1n	linear	93	0.9107	[0.7705, 0.9890]
TabSyn	scaled_200_vae_1000_diffusion	own lineages	1n	mlp	87	0.9170	[0.7414, 0.9844]
TabSyn	scaled_200_vae_1000_diffusion	own lineages	4n	catboost	98	0.8966	[0.8188, 0.9779]
TabSyn	scaled_200_vae_1000_diffusion	own lineages	4n	linear	93	0.9585	[0.8297, 0.9897]
TabSyn	scaled_200_vae_1000_diffusion	own lineages	4n	mlp	87	1.0314	[0.9689, 1.1865]
TabSyn	scaled_200_vae_1000_diffusion	matched lineages	1n	catboost	96	0.8599	[0.7076, 0.9459]
TabSyn	scaled_200_vae_1000_diffusion	matched lineages	1n	linear	91	0.9280	[0.7705, 0.9896]
TabSyn	scaled_200_vae_1000_diffusion	matched lineages	1n	mlp	72	0.9438	[0.8201, 0.9857]
TabSyn	scaled_200_vae_1000_diffusion	matched lineages	4n	catboost	96	0.9009	[0.8161, 0.9788]
TabSyn	scaled_200_vae_1000_diffusion	matched lineages	4n	linear	91	0.9595	[0.8310, 0.9897]
TabSyn	scaled_200_vae_1000_diffusion	matched lineages	4n	mlp	72	1.0355	[0.9992, 1.1968]
DOPE	features12_steps2048	matched lineages	1n	catboost	96	0.8852	[0.8224, 0.9690]
DOPE	features12_steps2048	matched lineages	1n	linear	91	0.9812	[0.9659, 0.9942]
DOPE	features12_steps2048	matched lineages	1n	mlp	72	0.9732	[0.9490, 0.9898]
DOPE	features12_steps2048	matched lineages	4n	catboost	96	0.9386	[0.8922, 0.9831]
DOPE	features12_steps2048	matched lineages	4n	linear	91	0.9940	[0.9718, 0.9967]
DOPE	features12_steps2048	matched lineages	4n	mlp	72	1.0684	[1.0048, 1.2714]

Table 12: **Seed exchangeability and the binary-by-seed diagnostic, CatBoost.** The count of lineages where seed 11 is the best of five is descriptive: a seed acts on every lineage at once, so the lineages are not independent trials and no p -value is given. The 2×2 cells are medians on the lineages common to all cells, with the median on each cell’s own lineages beside it. No Holm family.

Diagnostic	Count	Value
Seed 11 highest of five (CatBoost, $4n$)	47 of 99	19.8 if seeds were exchangeable; descriptive, no p -value
Seeds 23–71 only, CatBoost, n	99 lineages	0.791 [0.692, 0.930]
Seeds 23–71 only, Linear, n	94 lineages	0.984 [0.961, 0.992]
Seeds 23–71 only, MLP, n	88 lineages	0.958 [0.875, 0.990]
Seeds 23–71 only, CatBoost, $4n$	99 lineages	0.862 [0.771, 0.958]
Seeds 23–71 only, Linear, $4n$	94 lineages	0.989 [0.970, 0.996]
Seeds 23–71 only, MLP, $4n$	88 lineages	1.125 [1.025, 1.235]
2×2 cell: historical binary, seed 11	97 common (99 own)	0.949 (own lineages 0.944)
2×2 cell: historical binary, seeds 23–71	97 common (99 own)	0.862 (own lineages 0.862)
2×2 cell: later binary, seed 11	97 common (99 own)	0.949 (own lineages 0.944)
2×2 cell: later binary, seeds 23–71	97 common (99 own)	0.862 (own lineages 0.862)
2×2 cell: historical seed-11 population	97 common (97 own)	0.949 (own lineages 0.949)

Table 15: **DOPE minus TabSyn, fit seed 11.** Median paired retention difference on the matched lineages with its family-cluster bootstrap interval, the mean difference, the $v1$ mean-difference equivalence test at ± 0.02 , and W/T/L. The Holm family has three tests per size (one comparator, three auditors), separate from every other comparator block; the Wilcoxon and equivalence tests treat lineages as independent.

Size	Auditor	Comparator	Configuration	n	Median diff.	Hierarchical CI	Mean	TOST	W/T/L	Holm p
1n	catboost	TabSyn	scaled_200.vae.1000.diffusion	96	0.0122	[-0.0033, 0.0851]	0.0338	no	59/0/37	0.042
1n	linear	TabSyn	scaled_200.vae.1000.diffusion	91	0.0348	[-0.0008, 0.1265]	58.8632	no	61/0/30	0.000
1n	mlp	TabSyn	scaled_200.vae.1000.diffusion	72	0.0333	[-0.0078, 0.1272]	0.2127	no	43/0/29	0.011
4n	catboost	TabSyn	scaled_200.vae.1000.diffusion	96	0.0157	[-0.0033, 0.0890]	0.0737	no	56/0/40	0.056
4n	linear	TabSyn	scaled_200.vae.1000.diffusion	91	0.0199	[-0.0018, 0.1364]	68.9447	no	60/0/31	0.000
4n	mlp	TabSyn	scaled_200.vae.1000.diffusion	72	-0.0041	[-0.0384, 0.0738]	0.1739	no	33/0/39	0.682

Table 16: **$v1$ mean-difference equivalence, fit seed 11.** Mean DOPE-minus-comparator retention difference, its 90% t interval, and the larger of the two one-sided p -values against ± 0.02 . The test treats lineages as independent, and outliers move the mean; it is kept for continuity with the first version. Holm families as in the paired table; the equivalence p is unadjusted.

Size	Auditor	Comparator	Configuration	Mean diff.	90% CI	TOST p	Equivalent at ± 0.02
1n	catboost	GaussianCopula	native_selected	0.2483	[0.1779, 0.3188]	1.000	no
1n	catboost	Chow-Liu	native_selected	7.7460	[-3.7576, 19.2497]	0.866	no
1n	catboost	independent_marginals	native_selected	1.1580	[0.7097, 1.6062]	1.000	no
1n	linear	GaussianCopula	native_selected	0.2382	[0.0337, 0.4427]	0.960	no
1n	linear	Chow-Liu	native_selected	1.5301	[0.7702, 2.2899]	0.999	no
1n	linear	independent_marginals	native_selected	1.1054	[0.9906, 1.2202]	1.000	no
1n	mlp	GaussianCopula	native_selected	0.1971	[0.1217, 0.2725]	1.000	no
1n	mlp	Chow-Liu	native_selected	1.1319	[0.7996, 1.4642]	1.000	no
1n	mlp	independent_marginals	native_selected	1.3698	[1.1741, 1.5655]	1.000	no
4n	catboost	GaussianCopula	native_selected	0.2613	[0.2007, 0.3220]	1.000	no
4n	catboost	Chow-Liu	native_selected	8.0937	[-3.9560, 20.1434]	0.866	no
4n	catboost	independent_marginals	native_selected	1.2876	[0.6528, 1.9224]	0.999	no
4n	linear	GaussianCopula	native_selected	0.1832	[0.0184, 0.3480]	0.948	no
4n	linear	Chow-Liu	native_selected	1.4415	[0.7544, 2.1285]	1.000	no
4n	linear	independent_marginals	native_selected	1.0440	[0.9496, 1.1384]	1.000	no
4n	mlp	GaussianCopula	native_selected	0.3295	[0.2309, 0.4282]	1.000	no
4n	mlp	Chow-Liu	native_selected	0.9903	[0.6981, 1.2825]	1.000	no
4n	mlp	independent_marginals	native_selected	1.3100	[1.2007, 1.4194]	1.000	no
1n	catboost	CTGAN	native_selected	1.0872	[0.8392, 1.3352]	1.000	no
1n	catboost	TVAE	native_selected	0.2867	[0.1907, 0.3827]	1.000	no
1n	linear	CTGAN	native_selected	1.4244	[0.9705, 1.8783]	1.000	no
1n	linear	TVAE	native_selected	1.3010	[-0.2842, 2.8862]	0.911	no
1n	mlp	CTGAN	native_selected	2.1780	[1.0282, 3.3279]	0.998	no

Size	Auditor	Comparator	Configuration	Mean diff.	90% CI	TOST p	Equivalent at ± 0.02
1n	mlp	TVAE	native_selected	0.5985	[0.2704, 0.9265]	0.997	no
4n	catboost	CTGAN	native_selected	1.0962	[0.8335, 1.3588]	1.000	no
4n	catboost	TVAE	native_selected	0.2673	[0.1704, 0.3642]	1.000	no
4n	linear	CTGAN	native_selected	1.3489	[0.9775, 1.7204]	1.000	no
4n	linear	TVAE	native_selected	1.0124	[-0.1398, 2.1646]	0.923	no
4n	mlp	CTGAN	native_selected	1.8412	[1.1388, 2.5436]	1.000	no
4n	mlp	TVAE	native_selected	0.7771	[0.2190, 1.3353]	0.985	no
1n	catboost	ARF	author_default	0.1144	[0.0426, 0.1862]	0.984	no
1n	catboost	ARF	native_selected	0.1354	[0.0651, 0.2057]	0.996	no
1n	linear	ARF	author_default	0.1786	[0.0606, 0.2965]	0.986	no
1n	linear	ARF	native_selected	0.2123	[0.0869, 0.3378]	0.994	no
1n	mlp	ARF	author_default	0.2632	[0.1518, 0.3746]	1.000	no
1n	mlp	ARF	native_selected	0.3097	[0.1406, 0.4789]	0.997	no
4n	catboost	ARF	author_default	0.2059	[0.0480, 0.3638]	0.973	no
4n	catboost	ARF	native_selected	0.1737	[0.1098, 0.2376]	1.000	no
4n	linear	ARF	author_default	0.0924	[0.0269, 0.1579]	0.965	no
4n	linear	ARF	native_selected	0.1249	[0.0554, 0.1943]	0.993	no
4n	mlp	ARF	author_default	0.2894	[0.1767, 0.4022]	1.000	no
4n	mlp	ARF	native_selected	0.3089	[0.1650, 0.4527]	0.999	no
1n	catboost	ForestDiffusion/Forest-Flow	native_selected	-0.0917	[-0.2333, 0.0500]	0.823	no
1n	linear	ForestDiffusion/Forest-Flow	native_selected	-0.0040	[-0.0820, 0.0740]	0.342	no
1n	mlp	ForestDiffusion/Forest-Flow	native_selected	0.0723	[-0.1050, 0.2497]	0.731	no
4n	catboost	ForestDiffusion/Forest-Flow	native_selected	-0.1089	[-0.2579, 0.0401]	0.859	no
4n	linear	ForestDiffusion/Forest-Flow	native_selected	0.0118	[-0.1779, 0.2015]	0.466	no
4n	mlp	ForestDiffusion/Forest-Flow	native_selected	-0.0753	[-0.1953, 0.0447]	0.821	no

Table 17: **Source-family sensitivity, fit seed 11.** A DOPE row is the median DOPE retention within one source family; a comparator row is the median DOPE-minus-comparator difference on that family, under CatBoost. The excluding-simulated rows drop the four simulated families. Intervals are the family-cluster bootstrap. These rows are a sensitivity, printed without Holm p -values.

Family	Size	Auditor	Contrast	n	Median	Hierarchical CI
feynman	1n	catboost	DOPE	20	0.9703	[0.9454, 0.9906]
feynman	1n	catboost	GaussianCopula	20	0.2002	[0.1097, 0.2549]
feynman	1n	catboost	Chow-Liu	20	0.3253	[0.2200, 0.4252]
feynman	1n	catboost	independent_marginals	20	1.0201	[0.9776, 1.0254]
feynman	1n	linear	DOPE	20	0.9938	[0.9806, 0.9975]
feynman	1n	mlp	DOPE	19	0.9713	[0.9609, 0.9885]
feynman	4n	catboost	DOPE	20	0.9894	[0.9713, 0.9944]
feynman	4n	catboost	GaussianCopula	20	0.1894	[0.0993, 0.2840]
feynman	4n	catboost	Chow-Liu	20	0.2829	[0.2346, 0.3576]
feynman	4n	catboost	independent_marginals	20	0.9813	[0.9777, 0.9947]
feynman	4n	linear	DOPE	20	0.9941	[0.9854, 0.9982]
feynman	4n	mlp	DOPE	19	0.9965	[0.9860, 1.0067]
strogatz	1n	catboost	DOPE	14	0.7819	[0.4900, 0.9793]
strogatz	1n	catboost	GaussianCopula	14	0.1925	[-0.0086, 0.3450]
strogatz	1n	catboost	Chow-Liu	14	0.8451	[0.2765, 2.3318]
strogatz	1n	catboost	independent_marginals	14	0.8779	[0.6389, 1.0237]
strogatz	1n	linear	DOPE	13	0.9687	[0.8345, 0.9960]
strogatz	1n	mlp	DOPE	9	0.9967	[0.8086, 1.0923]
strogatz	4n	catboost	DOPE	14	0.8854	[0.5626, 0.9845]
strogatz	4n	catboost	GaussianCopula	14	0.1303	[-0.0086, 0.3896]
strogatz	4n	catboost	Chow-Liu	14	0.8739	[0.2276, 2.9286]
strogatz	4n	catboost	independent_marginals	14	0.9465	[0.6883, 1.0851]
strogatz	4n	linear	DOPE	13	0.9718	[0.7502, 0.9969]
strogatz	4n	mlp	DOPE	9	1.0815	[0.9862, 1.4876]
fri	1n	catboost	DOPE	24	0.8292	[0.7550, 0.8791]
fri	1n	catboost	GaussianCopula	24	0.4997	[0.3340, 0.5624]
fri	1n	catboost	Chow-Liu	24	0.3401	[0.2105, 0.4416]
fri	1n	catboost	independent_marginals	24	0.8660	[0.7905, 0.8976]
fri	1n	linear	DOPE	23	0.9675	[0.8787, 1.0009]
fri	1n	mlp	DOPE	15	0.9889	[0.8894, 1.0310]
fri	4n	catboost	DOPE	24	0.9054	[0.8485, 0.9276]
fri	4n	catboost	GaussianCopula	24	0.5583	[0.3530, 0.5913]
fri	4n	catboost	Chow-Liu	24	0.2964	[0.1981, 0.4561]
fri	4n	catboost	independent_marginals	24	0.9152	[0.8394, 0.9537]
fri	4n	linear	DOPE	23	0.9960	[0.9668, 1.0359]
fri	4n	mlp	DOPE	15	1.5792	[1.0250, 1.6797]
bng	1n	catboost	DOPE	6	0.9118	[0.0978, 1.1020]
bng	1n	catboost	GaussianCopula	6	0.1041	[0.0472, 0.1614]
bng	1n	catboost	Chow-Liu	6	0.1781	[0.0312, 0.6353]
bng	1n	catboost	independent_marginals	6	1.0364	[0.5939, 1.3518]
bng	1n	linear	DOPE	5	0.9718	[0.9305, 1.0382]
bng	1n	mlp	DOPE	4	1.0475	[0.9751, 1.6088]
bng	4n	catboost	DOPE	6	0.9334	[0.2662, 1.0918]
bng	4n	catboost	GaussianCopula	6	0.0387	[-0.0122, 0.1742]

Family	Size	Auditor	Contrast	n	Median	Hierarchical CI
bng	4n	catboost	Chow-Liu	6	0.0761	[0.0430, 0.5724]
bng	4n	catboost	independent_marginals	6	0.9060	[0.4007, 1.2632]
bng	4n	linear	DOPE	5	0.9941	[0.9173, 1.0458]
bng	4n	mlp	DOPE	4	1.2962	[0.9871, 1.7049]
other	1n	catboost	DOPE	33	0.9020	[0.8417, 0.9725]
other	1n	catboost	GaussianCopula	33	0.1651	[0.0633, 0.2597]
other	1n	catboost	Chow-Liu	33	0.4339	[0.2997, 0.6052]
other	1n	catboost	independent_marginals	33	0.9217	[0.8303, 0.9973]
other	1n	linear	DOPE	31	0.9812	[0.9713, 0.9952]
other	1n	mlp	DOPE	28	0.9322	[0.8545, 0.9864]
other	4n	catboost	DOPE	33	0.9373	[0.8908, 0.9878]
other	4n	catboost	GaussianCopula	33	0.1912	[0.0762, 0.2495]
other	4n	catboost	Chow-Liu	33	0.4514	[0.3214, 0.6687]
other	4n	catboost	independent_marginals	33	0.9293	[0.8699, 1.0026]
other	4n	linear	DOPE	31	0.9951	[0.9662, 1.0017]
other	4n	mlp	DOPE	28	1.0955	[1.0357, 1.1953]
excluding simulated	1n	catboost	GaussianCopula	33	0.1651	[0.0698, 0.2637]
excluding simulated	1n	catboost	Chow-Liu	33	0.4339	[0.2997, 0.6052]
excluding simulated	1n	catboost	independent_marginals	33	0.9217	[0.8575, 0.9973]
excluding simulated	4n	catboost	GaussianCopula	33	0.1912	[0.0762, 0.2495]
excluding simulated	4n	catboost	Chow-Liu	33	0.4514	[0.2957, 0.6687]
excluding simulated	4n	catboost	independent_marginals	33	0.9293	[0.8699, 1.0026]

Table 18: **Informative-lineage denominators, fit seed 11, CatBoost.** Informative is the lineage count behind each retention level. Complete-loss rows have three finite stored null, TRTR, and TSTR losses, and the noninformative, undefined, capped, and sensitivity counts are taken among them. No Holm family.

Method	Configuration	Size	Informative	Complete-loss rows	Non-informative	Undefined	Capped	Sensitivity n	Capped median
DOPE	features12_steps2048	1n	97	98	1	0	1	98	0.8862
DOPE	features12_steps2048	4n	97	98	1	0	1	98	0.9359
GaussianCopula	native_selected	1n	99	100	1	0	1	100	0.5996
GaussianCopula	native_selected	4n	99	100	1	0	1	100	0.6658
Chow-Liu	native_selected	1n	99	100	1	0	12	100	0.4809
Chow-Liu	native_selected	4n	99	100	1	0	11	100	0.5782
independent_marginals	native_selected	1n	99	100	1	0	3	100	-0.0440
independent_marginals	native_selected	4n	99	100	1	0	1	100	-0.0120
CTGAN	native_selected	1n	21	22	1	0	2	22	-0.0889
CTGAN	native_selected	4n	21	22	1	0	2	22	-0.0801
TVAE	native_selected	1n	21	22	1	0	0	22	0.6094
TVAE	native_selected	4n	21	22	1	0	0	22	0.6582
ARF	author_default	1n	99	100	1	0	1	100	0.7505
ARF	author_default	4n	99	100	1	0	1	100	0.8314
ARF	native_selected	1n	99	100	1	0	1	100	0.7410
ARF	native_selected	4n	99	100	1	0	1	100	0.8237
ForestDiffusion/Forest-Flow	native_selected	1n	6	0	0	0	0	0	—
ForestDiffusion/Forest-Flow	native_selected	4n	6	0	0	0	0	0	—

Table 19: **CatBoost auditor RMSE, fit seed 11.** Median over regression lineages of the validation RMSE of the auditor trained on synthetic rows and of the auditor trained on the real fit rows. RMSE is in each lineage's target units, so a median mixes scales and is no cross-method ranking. No Holm family.

Method	Configuration	Size	Regression lineages	Synthetic RMSE	Real TRTR RMSE
DOPE	features12_steps2048	1n	98	0.0887	0.0522
DOPE	features12_steps2048	4n	98	0.0786	0.0522
GaussianCopula	native_selected	1n	100	0.1166	0.0529
GaussianCopula	native_selected	4n	100	0.1124	0.0529
Chow-Liu	native_selected	1n	100	0.1359	0.0529
Chow-Liu	native_selected	4n	100	0.1280	0.0529
independent_marginals	native_selected	1n	100	0.1796	0.0529
independent_marginals	native_selected	4n	100	0.1760	0.0529
CTGAN	native_selected	1n	22	0.1862	0.0678
CTGAN	native_selected	4n	22	0.1780	0.0678
TVAE	native_selected	1n	22	0.1314	0.0678
TVAE	native_selected	4n	22	0.1257	0.0678
ARF	author_default	1n	100	0.0971	0.0529
ARF	author_default	4n	100	0.0866	0.0529
ARF	native_selected	1n	100	0.1019	0.0529
ARF	native_selected	4n	100	0.0958	0.0529
ForestDiffusion/Forest-Flow	native_selected	1n	0	—	—

Table 20: **Density block at $4n$, fit seed 11.** Each cell is a retention; the denominator is the paired informative lineages in that row. Holm family size is nine. W/T/L counts lineages; a tie is a difference of zero.

Auditor	Comparator	n	DOPE	Other	Difference	W/T/L	Holm p
CatBoost	Gaussian copula	97	0.940 [0.914, 0.976]	0.656 [0.431, 0.803]	0.224 [0.149, 0.290]	86/0/11	1.35×10^{-13}
CatBoost	Chow-Liu	97	0.940 [0.914, 0.976]	0.586 [0.465, 0.637]	0.337 [0.282, 0.451]	94/0/3	4.85×10^{-16}
CatBoost	Indep. marginals	97	0.940 [0.914, 0.976]	-0.012 [-0.020, -0.00649]	0.956 [0.918, 0.980]	96/0/1	1.17×10^{-16}
Linear	Gaussian copula	92	0.994 [0.981, 0.996]	0.989 [0.968, 0.994]	0.00366 [-0.00438, 0.012]	49/0/43	0.373
Linear	Chow-Liu	92	0.994 [0.981, 0.996]	0.661 [0.356, 0.738]	0.340 [0.244, 0.425]	86/0/6	7.12×10^{-15}
Linear	Indep. marginals	92	0.994 [0.981, 0.996]	0.00206 [-0.023, 0.00583]	0.988 [0.975, 1.001]	92/0/0	5.70×10^{-16}
MLP	Gaussian copula	75	1.061 [1.024, 1.125]	0.928 [0.865, 0.953]	0.185 [0.128, 0.228]	65/0/10	4.96×10^{-11}
MLP	Chow-Liu	75	1.061 [1.024, 1.125]	0.692 [0.582, 0.754]	0.530 [0.350, 0.651]	74/0/1	2.11×10^{-13}
MLP	Indep. marginals	75	1.061 [1.024, 1.125]	-0.028 [-0.047, -0.013]	1.104 [1.047, 1.179]	75/0/0	2.11×10^{-13}

Method	Configuration	Size	Regression lineages	Synthetic RMSE	Real TRTR RMSE
ForestDiffusion/Forest-Flow	native_selected	$4n$	0	—	—

E.2 HISTORICAL BLOCK TABLES

The density, neural, and Forest-Flow block tables below are the fit-seed-11 receipts behind the first version of this paper; their intervals are lineage-bootstrap intervals of the median, and they are kept for continuity. The paired median difference is distinct from the difference of two medians. The two DOPE artifacts above the cap in that population are 4544_GeographicalOriginalofMusic at 22,233 bytes and 588_fri_c4_1000_100 at 19,019 bytes, and 678_visualizing_environmental has a fit whose CatBoost group is uninformative, so the CatBoost median uses 97 lineages while the byte count uses 100. The displayed byte range runs from 758 to 22,233. CatBoost reads RMSE 0.053 for the real-trained auditor and 0.081 for the synthetic-trained one, with training-mean skill 0.881 and 0.792; RMSE is the square root of the stored mean squared error, and conventional held-out R^2 is unstored for these rows.

F CONTROLS

The real-row control draws the published training rows with replacement until the table has four times the training row count. The predictor-only control fits a ridge regression on the training rows, samples each feature independently from its training marginal, and adds Gaussian residual noise; neither control uses the DOPE encoding, and the predictor-only gap bundles DOPE’s dependence model with DOPE’s nonlinear residual. The paired column is DOPE minus the control. In the fit-seed-11 control table below, the size- n Holm family is the three predictor-only auditor tests, because the real-row control is defined at $4n$ only, and the size- $4n$ family has six tests; the five-seed panel keeps all six declared control slots at each size. Extra predictor fit seeds are reported and are outside the Holm family, and split-seed rows are a separate sensitivity.

Table 26: **DOPE minus each control, fit seed 11.** A paired row is the median DOPE-minus-control retention difference with its family-cluster bootstrap 95% interval, the mean difference, the v1 mean-difference equivalence test at ± 0.02 , and W/T/L over lineages; the Wilcoxon and equivalence tests treat lineages as independent. Predictor fit-seed and split-seed rows give the control’s own CatBoost retention level. The size- n Holm family is the three predictor-only tests, because the real-row control exists at $4n$ only; the size- $4n$ family has six tests. Fit-seed and split-seed rows are outside both Holm families.

Size	Auditor	Control	n	Median diff. or level	Hierarchical CI	Mean	TOST	W/T/L	Holm p
1n	catboost	predictor_only	97	0.2528	[0.1449, 0.3903]	0.2286	no	87/0/10	0.000
1n	linear	predictor_only	92	-0.0026	[-0.0176, 0.0137]	0.0156	no	42/0/50	0.523
1n	mlp	predictor_only	73	0.0739	[-0.0009, 0.1498]	0.1331	no	52/0/21	0.000
4n	catboost	real_bootstrap_4n	97	-0.0806	[-0.1276, -0.0290]	-0.1922	no	14/0/83	0.000
4n	catboost	predictor_only	97	0.2680	[0.0784, 0.3687]	0.2265	no	84/0/13	0.000
4n	linear	real_bootstrap_4n	92	-0.0058	[-0.0203, 0.0028]	0.0028	no	34/0/58	0.147
4n	linear	predictor_only	92	-0.0051	[-0.0181, 0.0077]	-0.0023	no	39/0/53	0.229
4n	mlp	real_bootstrap_4n	73	-0.0522	[-0.1495, 0.0057]	-0.0256	no	23/0/50	0.087
4n	mlp	predictor_only	73	0.1296	[0.0356, 0.1867]	0.2230	no	55/0/18	0.000
1n	catboost	predictor_fit_11	99	0.5962	[0.3397, 0.7479]	—	—	—	—
4n	catboost	predictor_fit_11	99	0.6374	[0.3756, 0.7905]	—	—	—	—
1n	catboost	predictor_fit_23	99	0.5664	[0.3364, 0.7239]	—	—	—	—
4n	catboost	predictor_fit_23	99	0.6363	[0.3854, 0.7705]	—	—	—	—

Size	Auditor	Control	n	Median diff. or level	Hierarchical CI	Mean	TOST	W/T/L	Holm p
1n	catboost	predictor fit 37	99	0.5497	[0.3611, 0.7305]	—	—	—	—
4n	catboost	predictor fit 37	99	0.6263	[0.3547, 0.7705]	—	—	—	—
1n	catboost	predictor fit 53	99	0.6074	[0.3341, 0.7333]	—	—	—	—
4n	catboost	predictor fit 53	99	0.6164	[0.3525, 0.7823]	—	—	—	—
1n	catboost	predictor fit 71	99	0.5933	[0.3386, 0.7163]	—	—	—	—
4n	catboost	predictor fit 71	99	0.6722	[0.3677, 0.8007]	—	—	—	—
1n	catboost	split 2027	99	0.5300	[0.3215, 0.7287]	—	—	—	—
4n	catboost	split 2027	99	0.5997	[0.3654, 0.7784]	—	—	—	—
1n	catboost	split 2999	100	0.5705	[0.3371, 0.7224]	—	—	—	—
4n	catboost	split 2999	100	0.6475	[0.3684, 0.7774]	—	—	—	—
1n	catboost	split 4099	100	0.5623	[0.3075, 0.7163]	—	—	—	—
4n	catboost	split 4099	100	0.6399	[0.3621, 0.7742]	—	—	—	—
1n	catboost	split 8191	100	0.6108	[0.3366, 0.7559]	—	—	—	—
4n	catboost	split 8191	100	0.6767	[0.3594, 0.7694]	—	—	—	—

G FIDELITY AND EMPIRICAL PRIVACY

Table 30 is the grouped-validation fidelity ledger of the fit-seed-11 population. On the density cells, DOPE has median logistic two-sample AUC 0.517 (0.504 to 0.532, $n = 98$), mean marginal Kolmogorov–Smirnov distance 0.114 (0.103 to 0.135, $n = 98$), and pairwise correlation fidelity 0.953 (0.945 to 0.961, $n = 98$); each count is the lineages with a complete three-seed fidelity group. A logistic detector near one half can still miss nonlinear differences, which is why the main paper reports a CatBoost detector. The DCR column of that ledger is the fraction of synthetic queries whose nearest reference is in the fit view, with ties scored one half. Under an independent draw, the unbalanced share sits near the fit-view fraction of the pooled references. The grouped split keeps about four fifths of those rows in the fit view, so the baseline is near four fifths, not one half. The DOMIAS column is the isolated kernel density-ratio equation from the pinned author expression, evaluated on this numeric projection, which differs from the original DOMIAS training pipeline.

Table 28 is the review privacy panel behind Table 3 of the main paper: for every method on the fit-seed-11 samples, a CatBoost two-sample detector with at most 200 rows per class, the median distance to the closest fit row and to the closest validation row, the nearest-neighbour distance ratio against the fit rows, and a distance-based membership AUC. The fit reference holds about four times as many rows as the validation reference, so the distance to the fit rows is smaller by construction. None of these numbers is a differential-privacy guarantee or a HIPAA de-identification claim. Shadow-model attacks, attribute inference, and an attack on the shipped artifact are defined in predeclare_v2 and are scheduled. Table 27 compares every arm’s marginal, pairwise, and logistic-detector fidelity on the five-seed cells, and Table 29 adds the stored fidelity of the fit-seed-11 population ledgers.

Table 27: **Fidelity of every arm at 4n, five-seed cells.** Median over lineages with the 95% nested bootstrap interval of the mean marginal Kolmogorov–Smirnov distance (lower is closer), the pairwise correlation fidelity (higher is closer), and a logistic two-sample AUC (one half is indistinguishable); n counts lineages with a complete three-seed group. A five-seed arm needs three complete fits per lineage. No Holm family.

Arm	n	KS distance	Pair correlation	Logistic C2ST
DOPE (ours)	100	0.117 [0.089, 0.165]	0.953 [0.930, 0.964]	0.531 [0.516, 0.548]
DOPE, untuned default	100	0.116 [0.090, 0.165]	0.952 [0.932, 0.962]	0.532 [0.512, 0.547]
TabSyn (scaled schedule)	99	0.120 [0.083, 0.135]	0.958 [0.946, 0.963]	0.549 [0.523, 0.587]
ARF (author default)	100	0.103 [0.084, 0.140]	0.962 [0.948, 0.965]	0.521 [0.505, 0.538]
ARF (native selected)	100	0.108 [0.084, 0.138]	0.962 [0.946, 0.964]	0.521 [0.506, 0.539]
Gaussian copula	100	0.134 [0.114, 0.211]	0.963 [0.950, 0.969]	0.520 [0.502, 0.563]
Chow-Liu tree	100	0.127 [0.085, 0.229]	0.928 [0.891, 0.945]	0.555 [0.516, 0.625]
Independent marginals	100	0.130 [0.081, 0.267]	0.899 [0.826, 0.934]	0.517 [0.494, 0.568]
TVAE (native selected)	22	0.138 [0.117, 0.243]	0.926 [0.893, 0.945]	0.582 [0.574, 0.629]
CTGAN (native selected)	22	0.230 [0.184, 0.301]	0.912 [0.837, 0.943]	0.770 [0.667, 0.858]
Forest-Flow (historical)	6	0.123 [0.056, 0.206]	0.969 [0.932, 0.981]	0.492 [0.462, 0.540]
Predictor-only control	100	0.105 [0.084, 0.141]	0.945 [0.882, 0.961]	0.510 [0.497, 0.527]
Real rows resampled	100	0.080 [0.069, 0.106]	0.964 [0.954, 0.970]	0.503 [0.482, 0.530]

Table 21: **Neural block at $4n$, fit seed 11.** The denominator is the paired informative lineages in the cell; CatBoost uses 21. Holm family size is six. W/T/L counts lineages; a tie is a difference of zero.

Auditor	Comparator	n	DOPE	Other	Difference	W/T/L	Holm p
CatBoost	CTGAN	21	0.953 [0.896, 0.990]	-0.081 [-0.202, 0.103]	0.881 [0.702, 1.150]	21/0/0	5.72×10^{-6}
CatBoost	TVAE	21	0.953 [0.896, 0.990]	0.669 [0.539, 0.787]	0.236 [0.087, 0.431]	19/0/2	1.57×10^{-4}
Linear	CTGAN	21	0.999 [0.996, 1.015]	-0.033 [-0.225, 0.080]	1.049 [0.916, 1.236]	21/0/0	5.72×10^{-6}
Linear	TVAE	21	0.999 [0.996, 1.015]	0.796 [0.605, 0.867]	0.263 [0.166, 0.410]	20/0/1	1.57×10^{-4}
MLP	CTGAN	19	1.026 [0.998, 1.263]	-0.149 [-0.471, 0.083]	1.079 [0.877, 2.030]	19/0/0	1.53×10^{-5}
MLP	TVAE	19	1.026 [0.998, 1.263]	0.753 [0.478, 0.852]	0.284 [0.153, 0.734]	17/0/2	1.57×10^{-4}

Table 22: **Forest-Flow confirmation at $4n$, CatBoost, fit seed 11.** The denominator is 6 lineages. Holm family size is three.

Lineage	DOPE bytes	DOPE	Forest-Flow bytes	Forest-Flow
1193_BNG_lowbwt	2,436	0.970	416,921,273	1.004
503_wind	3,846	0.996	637,056,689	1.016
582_fri.c1_500.25	5,670	0.909	1,085,155,062	0.883
596_fri.c2_250.5	1,698	0.807	234,776,554	0.968
feynman.L12_1	1,032	0.991	123,784,681	1.001
strogatz_vdp1	984	0.581	115,202,253	1.036

Table 28: **Holdout privacy and CatBoost detection, fit seed 11.** Each value is the median over lineages of each lineage’s median over three sample seeds, with the family-cluster bootstrap interval. Own rows use each method’s lineages whose synthetic row count equals fit rows times size; matched-density rows use the lineages shared by DOPE, the Gaussian copula, Chow–Liu, and independent marginals. Fit DCR and NDR use the fit rows as the reference set and holdout DCR the validation rows, about a quarter as many. Distance MIA AUC is an empirical membership attack, and c2st.catboost_auc is the CatBoost detector. These are empirical attack scores, not formal differential privacy and not HIPAA de-identification. Unavailable cells are no wins: CTGAN 468 (sample.csv_unresolved); DOPE 12 (sample.csv_unresolved); ForestDiffusion/Forest-Flow 36 (sample.csv_absent); TVAE 468 (sample.csv_unresolved). No Holm family.

Method	Configuration	Size	Metric	n lineages	Median	Hierarchical CI
ARF	author_default	1n	dcr_validation_median	100	0.1095	[0.0532, 0.1639]
ARF	author_default	1n	dcr_fit_median	100	0.0826	[0.0332, 0.1257]
ARF	author_default	1n	nndr_fit_median	100	0.8983	[0.8617, 0.9278]
ARF	author_default	1n	distance_mia_auc	100	0.5452	[0.5298, 0.5580]
ARF	author_default	1n	c2st_catboost_auc	100	0.6159	[0.5751, 0.6928]
ARF	author_default	4n	dcr_validation_median	100	0.1088	[0.0517, 0.1617]
ARF	author_default	4n	dcr_fit_median	100	0.0820	[0.0331, 0.1253]
ARF	author_default	4n	nndr_fit_median	100	0.9069	[0.8639, 0.9282]
ARF	author_default	4n	distance_mia_auc	100	0.5563	[0.5358, 0.5754]
ARF	author_default	4n	c2st_catboost_auc	100	0.6120	[0.5696, 0.6877]
ARF	native_selected	1n	dcr_validation_median	100	0.1108	[0.0544, 0.1637]
ARF	native_selected	1n	dcr_fit_median	100	0.0823	[0.0361, 0.1316]
ARF	native_selected	1n	nndr_fit_median	100	0.9079	[0.8687, 0.9271]
ARF	native_selected	1n	distance_mia_auc	100	0.5422	[0.5309, 0.5576]
ARF	native_selected	1n	c2st_catboost_auc	100	0.6100	[0.5789, 0.7089]
ARF	native_selected	4n	dcr_validation_median	100	0.1109	[0.0547, 0.1635]
ARF	native_selected	4n	dcr_fit_median	100	0.0825	[0.0389, 0.1295]
ARF	native_selected	4n	nndr_fit_median	100	0.9107	[0.8647, 0.9281]
ARF	native_selected	4n	distance_mia_auc	100	0.5408	[0.5303, 0.5562]
ARF	native_selected	4n	c2st_catboost_auc	100	0.6091	[0.5766, 0.7075]
CTGAN	native_selected	1n	dcr_validation_median	22	0.1863	[0.1163, 0.2214]
CTGAN	native_selected	1n	dcr_fit_median	22	0.1538	[0.0903, 0.1899]
CTGAN	native_selected	1n	nndr_fit_median	22	0.9395	[0.9106, 0.9610]
CTGAN	native_selected	1n	distance_mia_auc	22	0.4804	[0.4514, 0.5101]
CTGAN	native_selected	1n	c2st_catboost_auc	21	0.9522	[0.8778, 0.9841]
CTGAN	native_selected	4n	dcr_validation_median	22	0.1869	[0.1157, 0.2225]
CTGAN	native_selected	4n	dcr_fit_median	22	0.1529	[0.0867, 0.1906]
CTGAN	native_selected	4n	nndr_fit_median	22	0.9413	[0.9117, 0.9581]
CTGAN	native_selected	4n	distance_mia_auc	22	0.5024	[0.4721, 0.5168]
CTGAN	native_selected	4n	c2st_catboost_auc	22	0.9535	[0.8487, 0.9777]
Chow-Liu	native_selected	1n	dcr_validation_median	100	0.1334	[0.0976, 0.2006]
Chow-Liu	native_selected	1n	dcr_fit_median	100	0.1078	[0.0786, 0.1671]
Chow-Liu	native_selected	1n	nndr_fit_median	100	0.9228	[0.8889, 0.9638]
Chow-Liu	native_selected	1n	distance_mia_auc	100	0.5077	[0.4962, 0.5202]
Chow-Liu	native_selected	1n	c2st_catboost_auc	100	0.6868	[0.6025, 0.8986]
Chow-Liu	native_selected	4n	dcr_validation_median	100	0.1291	[0.1053, 0.1901]
Chow-Liu	native_selected	4n	dcr_fit_median	100	0.1080	[0.0858, 0.1661]
Chow-Liu	native_selected	4n	nndr_fit_median	100	0.9238	[0.8883, 0.9644]
Chow-Liu	native_selected	4n	distance_mia_auc	100	0.5196	[0.5037, 0.5350]

Method	Configuration	Size	Metric	n lineages	Median	Hierarchical CI
Chow-Liu	native_selected	4n	c2st_catboost_auc	99	0.7061	[0.5910, 0.9369]
DOPE	features12_steps2048	1n	dcr_validation_median	98	0.1054	[0.0529, 0.1560]
DOPE	features12_steps2048	1n	dcr_fit_median	98	0.0800	[0.0423, 0.1254]
DOPE	features12_steps2048	1n	nndr_fit_median	98	0.9040	[0.8465, 0.9368]
DOPE	features12_steps2048	1n	distance_mia_auc	98	0.5132	[0.5012, 0.5249]
DOPE	features12_steps2048	1n	c2st_catboost_auc	98	0.6344	[0.5705, 0.8368]
DOPE	features12_steps2048	4n	dcr_validation_median	98	0.1086	[0.0528, 0.1521]
DOPE	features12_steps2048	4n	dcr_fit_median	98	0.0797	[0.0421, 0.1280]
DOPE	features12_steps2048	4n	nndr_fit_median	98	0.9053	[0.8482, 0.9360]
DOPE	features12_steps2048	4n	distance_mia_auc	98	0.5052	[0.4921, 0.5160]
DOPE	features12_steps2048	4n	c2st_catboost_auc	98	0.6405	[0.5830, 0.8312]
GaussianCopula	native_selected	1n	dcr_validation_median	100	0.1094	[0.0780, 0.1543]
GaussianCopula	native_selected	1n	dcr_fit_median	100	0.0868	[0.0621, 0.1268]
GaussianCopula	native_selected	1n	nndr_fit_median	100	0.9177	[0.8814, 0.9603]
GaussianCopula	native_selected	1n	distance_mia_auc	100	0.5050	[0.4960, 0.5293]
GaussianCopula	native_selected	1n	c2st_catboost_auc	100	0.7343	[0.6789, 0.8981]
GaussianCopula	native_selected	4n	dcr_validation_median	100	0.1074	[0.0774, 0.1563]
GaussianCopula	native_selected	4n	dcr_fit_median	100	0.0863	[0.0627, 0.1280]
GaussianCopula	native_selected	4n	nndr_fit_median	100	0.9201	[0.8867, 0.9610]
GaussianCopula	native_selected	4n	distance_mia_auc	100	0.5022	[0.4851, 0.5158]
GaussianCopula	native_selected	4n	c2st_catboost_auc	99	0.7346	[0.6902, 0.8689]
TVAE	native_selected	1n	dcr_validation_median	22	0.1122	[0.0853, 0.1943]
TVAE	native_selected	1n	dcr_fit_median	22	0.0815	[0.0533, 0.1535]
TVAE	native_selected	1n	nndr_fit_median	22	0.9056	[0.8479, 0.9264]
TVAE	native_selected	1n	distance_mia_auc	22	0.5312	[0.4992, 0.5585]
TVAE	native_selected	1n	c2st_catboost_auc	22	0.7479	[0.7221, 0.8859]
TVAE	native_selected	4n	dcr_validation_median	22	0.1118	[0.0855, 0.1938]
TVAE	native_selected	4n	dcr_fit_median	22	0.0799	[0.0535, 0.1516]
TVAE	native_selected	4n	nndr_fit_median	22	0.9058	[0.8489, 0.9312]
TVAE	native_selected	4n	distance_mia_auc	22	0.5315	[0.4915, 0.5424]
TVAE	native_selected	4n	c2st_catboost_auc	22	0.7309	[0.7032, 0.8893]
independent_marginals	native_selected	1n	dcr_validation_median	100	0.1360	[0.1010, 0.1877]
independent_marginals	native_selected	1n	dcr_fit_median	100	0.1132	[0.0890, 0.1587]
independent_marginals	native_selected	1n	nndr_fit_median	100	0.9457	[0.9120, 0.9757]
independent_marginals	native_selected	1n	distance_mia_auc	100	0.4982	[0.4883, 0.5096]
independent_marginals	native_selected	1n	c2st_catboost_auc	99	0.8694	[0.7602, 0.9749]
independent_marginals	native_selected	4n	dcr_validation_median	100	0.1319	[0.0995, 0.1848]
independent_marginals	native_selected	4n	dcr_fit_median	100	0.1134	[0.0886, 0.1580]
independent_marginals	native_selected	4n	nndr_fit_median	100	0.9453	[0.9141, 0.9702]
independent_marginals	native_selected	4n	distance_mia_auc	100	0.5081	[0.4906, 0.5180]
independent_marginals	native_selected	4n	c2st_catboost_auc	100	0.8578	[0.7553, 0.9727]
DOPE	matched density	1n	dcr_validation_median	98	0.1054	[0.0529, 0.1560]
GaussianCopula	matched density	1n	dcr_validation_median	98	0.1094	[0.0770, 0.1543]
Chow-Liu	matched density	1n	dcr_validation_median	98	0.1304	[0.0976, 0.1913]
independent_marginals	matched density	1n	dcr_validation_median	98	0.1360	[0.1000, 0.1872]
DOPE	matched density	1n	dcr_fit_median	98	0.0800	[0.0423, 0.1254]
GaussianCopula	matched density	1n	dcr_fit_median	98	0.0859	[0.0617, 0.1268]
Chow-Liu	matched density	1n	dcr_fit_median	98	0.1075	[0.0786, 0.1613]
independent_marginals	matched density	1n	dcr_fit_median	98	0.1128	[0.0890, 0.1578]
DOPE	matched density	1n	nndr_fit_median	98	0.9040	[0.8495, 0.9368]
GaussianCopula	matched density	1n	nndr_fit_median	98	0.9166	[0.8814, 0.9603]
Chow-Liu	matched density	1n	nndr_fit_median	98	0.9219	[0.8888, 0.9638]
independent_marginals	matched density	1n	nndr_fit_median	98	0.9422	[0.9120, 0.9729]
DOPE	matched density	1n	distance_mia_auc	98	0.5132	[0.5012, 0.5249]
GaussianCopula	matched density	1n	distance_mia_auc	98	0.5035	[0.4949, 0.5235]
Chow-Liu	matched density	1n	distance_mia_auc	98	0.5077	[0.4958, 0.5203]
independent_marginals	matched density	1n	distance_mia_auc	98	0.4982	[0.4883, 0.5070]
DOPE	matched density	1n	c2st_catboost_auc	97	0.6234	[0.5705, 0.8368]
GaussianCopula	matched density	1n	c2st_catboost_auc	97	0.7296	[0.6596, 0.9162]
Chow-Liu	matched density	1n	c2st_catboost_auc	97	0.6866	[0.6055, 0.8986]
independent_marginals	matched density	1n	c2st_catboost_auc	97	0.8694	[0.7656, 0.9749]
DOPE	matched density	4n	dcr_validation_median	98	0.1086	[0.0528, 0.1521]
GaussianCopula	matched density	4n	dcr_validation_median	98	0.1074	[0.0766, 0.1554]
Chow-Liu	matched density	4n	dcr_validation_median	98	0.1262	[0.1020, 0.1874]
independent_marginals	matched density	4n	dcr_validation_median	98	0.1308	[0.0995, 0.1841]
DOPE	matched density	4n	dcr_fit_median	98	0.0797	[0.0404, 0.1280]
GaussianCopula	matched density	4n	dcr_fit_median	98	0.0859	[0.0627, 0.1273]
Chow-Liu	matched density	4n	dcr_fit_median	98	0.1042	[0.0853, 0.1642]
independent_marginals	matched density	4n	dcr_fit_median	98	0.1124	[0.0884, 0.1569]
DOPE	matched density	4n	nndr_fit_median	98	0.9053	[0.8492, 0.9360]
GaussianCopula	matched density	4n	nndr_fit_median	98	0.9198	[0.8858, 0.9603]
Chow-Liu	matched density	4n	nndr_fit_median	98	0.9221	[0.8883, 0.9644]
independent_marginals	matched density	4n	nndr_fit_median	98	0.9435	[0.9119, 0.9702]
DOPE	matched density	4n	distance_mia_auc	98	0.5052	[0.4919, 0.5160]
GaussianCopula	matched density	4n	distance_mia_auc	98	0.5004	[0.4850, 0.5152]
Chow-Liu	matched density	4n	distance_mia_auc	98	0.5196	[0.5023, 0.5380]
independent_marginals	matched density	4n	distance_mia_auc	98	0.5077	[0.4891, 0.5176]
DOPE	matched density	4n	c2st_catboost_auc	96	0.6465	[0.5863, 0.8508]
GaussianCopula	matched density	4n	c2st_catboost_auc	96	0.7348	[0.6902, 0.8804]

Method	Configuration	Size	Metric	n lineages	Median	Hierarchical CI
Chow-Liu	matched density	$4n$	c2st_catboost_auc	96	0.7062	[0.6017, 0.9302]
independent_marginals	matched density	$4n$	c2st_catboost_auc	96	0.8578	[0.7555, 0.9727]

Table 29: **Stored fidelity from the population ledgers, fit seed 11.** Logistic two-sample AUC, mean marginal Kolmogorov–Smirnov distance, and pair-correlation fidelity; n counts lineages with three finite stored values, with the family-cluster bootstrap interval. Cohorts differ across methods, so the rows are no ranking. The stored detector is logistic, not CatBoost, and TabSyn is absent from these ledgers. No Holm family.

Method	Configuration	Size	Metric	n lineages	Median	Hierarchical CI
DOPE	features12_steps2048	$1n$	stored logistic C2ST	98	0.5172	[0.5000, 0.5376]
DOPE	features12_steps2048	$1n$	marginal KS	98	0.1216	[0.0950, 0.1680]
DOPE	features12_steps2048	$1n$	pair correlation	98	0.9532	[0.9317, 0.9610]
DOPE	features12_steps2048	$4n$	stored logistic C2ST	98	0.5172	[0.5000, 0.5376]
DOPE	features12_steps2048	$4n$	marginal KS	98	0.1138	[0.0917, 0.1583]
DOPE	features12_steps2048	$4n$	pair correlation	98	0.9531	[0.9346, 0.9645]
GaussianCopula	native_selected	$1n$	stored logistic C2ST	100	0.5246	[0.5100, 0.5518]
GaussianCopula	native_selected	$1n$	marginal KS	100	0.1393	[0.1171, 0.2132]
GaussianCopula	native_selected	$1n$	pair correlation	100	0.9614	[0.9516, 0.9652]
GaussianCopula	native_selected	$4n$	stored logistic C2ST	100	0.5246	[0.5100, 0.5537]
GaussianCopula	native_selected	$4n$	marginal KS	100	0.1336	[0.1137, 0.2115]
GaussianCopula	native_selected	$4n$	pair correlation	100	0.9635	[0.9488, 0.9687]
Chow-Liu	native_selected	$1n$	stored logistic C2ST	100	0.5423	[0.5040, 0.6292]
Chow-Liu	native_selected	$1n$	marginal KS	100	0.1293	[0.0892, 0.2361]
Chow-Liu	native_selected	$1n$	pair correlation	100	0.9273	[0.8834, 0.9418]
Chow-Liu	native_selected	$4n$	stored logistic C2ST	100	0.5404	[0.5185, 0.6303]
Chow-Liu	native_selected	$4n$	marginal KS	100	0.1266	[0.0852, 0.2292]
Chow-Liu	native_selected	$4n$	pair correlation	100	0.9283	[0.8906, 0.9465]
independent_marginals	native_selected	$1n$	stored logistic C2ST	100	0.5213	[0.5000, 0.5347]
independent_marginals	native_selected	$1n$	marginal KS	100	0.1363	[0.0866, 0.2847]
independent_marginals	native_selected	$1n$	pair correlation	100	0.8994	[0.8099, 0.9289]
independent_marginals	native_selected	$4n$	stored logistic C2ST	100	0.5166	[0.4988, 0.5488]
independent_marginals	native_selected	$4n$	marginal KS	100	0.1297	[0.0813, 0.2672]
independent_marginals	native_selected	$4n$	pair correlation	100	0.8986	[0.8258, 0.9329]
CTGAN	native_selected	$1n$	stored logistic C2ST	22	0.7839	[0.6997, 0.8819]
CTGAN	native_selected	$1n$	marginal KS	22	0.2319	[0.1887, 0.2943]
CTGAN	native_selected	$1n$	pair correlation	22	0.9091	[0.8313, 0.9406]
CTGAN	native_selected	$4n$	stored logistic C2ST	22	0.7839	[0.6997, 0.8821]
CTGAN	native_selected	$4n$	marginal KS	22	0.2300	[0.1839, 0.2947]
CTGAN	native_selected	$4n$	pair correlation	22	0.9123	[0.8367, 0.9431]
TVAE	native_selected	$1n$	stored logistic C2ST	22	0.5918	[0.5616, 0.6392]
TVAE	native_selected	$1n$	marginal KS	22	0.1500	[0.1199, 0.2598]
TVAE	native_selected	$1n$	pair correlation	22	0.9240	[0.8905, 0.9418]
TVAE	native_selected	$4n$	stored logistic C2ST	22	0.6033	[0.5544, 0.6681]
TVAE	native_selected	$4n$	marginal KS	22	0.1382	[0.1174, 0.2431]
TVAE	native_selected	$4n$	pair correlation	22	0.9258	[0.8913, 0.9451]
ARF	author_default	$1n$	stored logistic C2ST	100	0.5321	[0.5156, 0.5442]
ARF	author_default	$1n$	marginal KS	100	0.1128	[0.0895, 0.1433]
ARF	author_default	$1n$	pair correlation	100	0.9595	[0.9471, 0.9644]
ARF	author_default	$4n$	stored logistic C2ST	100	0.5268	[0.5069, 0.5388]
ARF	author_default	$4n$	marginal KS	100	0.1035	[0.0846, 0.1340]
ARF	author_default	$4n$	pair correlation	100	0.9618	[0.9487, 0.9653]
ARF	native_selected	$1n$	stored logistic C2ST	100	0.5255	[0.5072, 0.5449]
ARF	native_selected	$1n$	marginal KS	100	0.1122	[0.0869, 0.1445]
ARF	native_selected	$1n$	pair correlation	100	0.9599	[0.9418, 0.9626]
ARF	native_selected	$4n$	stored logistic C2ST	100	0.5296	[0.5094, 0.5375]
ARF	native_selected	$4n$	marginal KS	100	0.1046	[0.0845, 0.1382]
ARF	native_selected	$4n$	pair correlation	100	0.9620	[0.9464, 0.9643]
ForestDiffusion/Forest-Flow	native_selected	$1n$	stored logistic C2ST	6	0.5187	[0.4139, 0.7000]
ForestDiffusion/Forest-Flow	native_selected	$1n$	marginal KS	6	0.1254	[0.0612, 0.1992]
ForestDiffusion/Forest-Flow	native_selected	$1n$	pair correlation	6	0.9548	[0.9242, 0.9799]
ForestDiffusion/Forest-Flow	native_selected	$4n$	stored logistic C2ST	6	0.5187	[0.4139, 0.7000]
ForestDiffusion/Forest-Flow	native_selected	$4n$	marginal KS	6	0.1231	[0.0557, 0.2062]
ForestDiffusion/Forest-Flow	native_selected	$4n$	pair correlation	6	0.9691	[0.9321, 0.9811]

G.1 CONFIGURATION-AWARE FIDELITY AND DETECTORS

The side-by-side columns use the common numeric projection and the published pilot KS and pair-fidelity formulas, including the projected target. Alpha-precision and beta-recall use the fixed unembedded equations. The stored logistic detector remains distinct from the shared grouped CatBoost detector: five disjoint row-group folds, the fixed iteration and depth budget, and the

Table 23: **Published ARF matched receipts at $4n$, CatBoost.** The denominator is the paired complete informative lineages in the generated row. No Holm family in this table.

Publication	Lineages	DOPE	Receipt	Difference
ARF, author_default	97	0.940	0.840	0.077
ARF, native_selected	97	0.940	0.826	0.095

Table 24: **Raw auditor medians on the informative lineages, fit seed 11.** The denominator is the informative lineages for that auditor. No Holm family.

Auditor	n	RMSE real	RMSE synth.	training-mean	training-mean
				skill, real	skill, synth.
CatBoost	97	0.053	0.081	0.881	0.792
Linear	92	0.105	0.112	0.585	0.551
MLP	75	0.098	0.083	0.666	0.785

predeclared equal per-class cap. Every metric gives its own complete-three-seed lineage count and family-cluster interval. These marginal summaries use different availability cohorts and establish no cross-method ranking. ARF author-default and native-selected configurations stay separate. Forest-Flow and originally unmeasured size- n alpha/beta entries retain unavailable cells. The retained mixed-type TabSyn alpha/beta cohort remains separate from the common-numeric cohort. No Holm family is applied to these descriptive tables.

H ABLATIONS AND REPLAY

Step budget. `features12_steps8192` is the same residual with four times as many AdamW steps. Where both budgets have a complete three-seed group, the CatBoost median paired difference is 0.034 (68 lineages); the linear medians are 0.990 and 0.987 (64 lineages; difference -0.000283) and the MLP medians 1.065 and 1.057 (52 lineages; difference 0.00451). Appendix J.2 lists every lineage. **Profile and default.** The main paper’s untuned trainer default is a second axis (width and step budget together). The full width, step, readout, screening-size, and byte-budget grid, and the interaction stress suite (XOR, parity, products, many relevant features), are defined in `predeclare_v2` and `scheduled`.

Replay. The scored runs did not store per-step losses; a later logged replay of both step budgets, before the readout refit, did. All 100 lineages finished. Median training loss starts at 0.321 and ends at 0.00274 and 0.000763; median validation loss starts at 0.271 and ends at 0.00443 and 0.00181. The replay files are distinct from the scored artifacts.

I BASELINE COVERAGE AND RETAINED EVIDENCE

Table 33 records which baseline publications exist and on how many lineages. The ARF follow-on supplies complete five-fit measurements on the same validation lineages. TabSyn’s full scaled-schedule panel is in the main paper; the earlier eight-lineage TabSyn subset below is a different cohort. Forest-Flow confirmation is distinct from the Forest-Diffusion pilot modes, and no full-population diffusion result is inferred from the Flow subset. Receipt paths and digests are in `docs/whitepaper/generated/baseline-coverage.json` and `docs/whitepaper/generated/retained-evidence.json`.

Retained validation. Tables 34 and 35 add the merged retained-sample evidence. Each lineage averages all three sample seeds before a lineage-bootstrap median, which differs from the median-of-seeds rule of the main tables. Table 34 covers ARF and the copula and Chow–Liu defaults and native configurations; Table 35 includes only the eight complete scaled-schedule TabSyn lineages of an earlier run and the classical methods on identical inputs. Table 36 gives exploratory TabSyn-minus-comparator differences on those lineages, with conservative simultaneous intervals (Bonferroni over four metrics times six configurations), which are unbounded at this count. Table 37 extends the published DOPE–ARF comparison with lineage intervals under its six-test family.

Table 25: **Charged bytes, fit seed 11.** For DOPE artifacts the charge is the encoded model file plus `projection.json`. The denominator is the artifacts with a stored charge. No Holm family.

Method	Median bytes	Within cap
DOPE reference	1,895.5	98/100
Gaussian copula	5,464.5	62/100
Chow-Liu	29,865	31/100
Independent marginals	2,549.5	94/100
CTGAN	741,834	0/22
TVAE	317,141.5	0/22
Forest-Flow	325,848,913.5	0/6
ARF (native selected)	1,627,718.5	0/100




Figure 4: **CatBoost retention at $4n$, fit seed 11.** Each bar is a method’s paired-lineage retention median and its lineage-bootstrap interval. Density uses 97 lineages and the neural panel 21; the panels are separate cohorts. No Holm family.

First-batch diagnostics. The canonical-job batch below is incomplete at population and sample-group level, so its values are individual physical evaluations without a population interval or ranking. DOPE jobs in this batch use frozen research profiles that differ from the reported profile. Pearson entries are the pinned source mean, alpha and beta use its pinned unembedded variant, and MIA is an empirical distance attack. Per-column marginals, detection folds, DCR and NNDR, costs, and job digests are in the first-batch source listed in `docs/whitepaper/generated/retained-evidence.json`.

BeyondArena (`2ecfe882ccfb814fc27c4de10a64ceefd5d7655c`) is a separate fit check outside the 100 PMLB lineages. 5 of the 12 families in that check finished inside the cap, and 7 did not prepare (4 outside the projected range, 3 overlapping); BeyondArena medians are 5.00 and 14.0. The inventory of all 142 families is in Appendix J.3. No retention auditor ran on BeyondArena.

J PER-LINEAGE TABLES

The tables below list every prepared lineage of the fit-seed-11 population, the sufficiency budget, and the BeyondArena inventory. They are the rows behind the medians of the historical block tables, and a blank cell is an unavailable fit or an uninformative auditor rather than a zero. The five-seed v2 cells are in the committed scalar ledger `research/benchmark/results/review-fixes-v2/scalar-cells.jsonl`, which holds one row per fit, sample seed, size, and auditor for every arm.

J.1 REPORTED PROFILE, ALL 100 LINEAGES (FIT SEED 11)

These rows are the historical fit-seed-11 population of the reported profile. The cap column marks charged bytes at or under 10,240; it is a byte check, not a release certificate. Retention columns are




Figure 5: **Density-block retention at $4n$ against charged bytes, per lineage, fit seed 11.** Each point is one lineage’s retention level (vertical) against its charged bytes on a log axis. Values outside $[-1.5, 1.6]$ are counted in the panel title and omitted from the frame. The dashed line is the 10,240-byte cap. No Holm family.




Figure 6: **CatBoost paired differences, sorted, and the Nemenyi scale, fit seed 11.** Rank 1 is the higher retention; the segment length is the critical difference. This display is outside the nine-test Holm family.

CatBoost at n and $4n$, linear at $4n$, and the MLP at $4n$, each the median of the three sample seeds; the identifier is the lineage’s dataset hash prefix and the name is its PMLB name.

Table 40: Displayed profile, all 100 lineages. The cap column is charged bytes at or under 10,240. It is not a release certificate. No Holm family.

id	name	rows	feat.	bytes	cap	CB n	CB $4n$	lin $4n$	MLP $4n$
02a45777900441ce	215_2dplanes	800	10	2,583	pass	0.9930	1.0001	0.9990	1.0477
047e62b52a3f4ede	695_chatfield_4	156	12	3,311	pass	1.0041	1.0155	1.0048	—
04c80944e971cb3e	first_principles_supernovae.zr	157	1	762	pass	0.8753	0.8846	0.9959	1.2627
0561d673cb73c27e	666_rmftsa_ladata	339	10	2,921	pass	1.0095	1.1026	1.0338	1.1394
060b9ff8092db17	589_fri_c2_1000_25	667	25	5,717	pass	0.8893	0.9534	0.9911	1.5794
06b0312a585b436e	1196_BNG_pharynx	800	10	2,631	pass	0.8601	0.8964	1.0458	1.7049
0779d3fd5b3133a2	nikuradse_2	241	1	758	pass	0.9096	0.9261	2.1723	2.2082
07f38946316c039b	593_fri.c1_1000_10	667	10	2,876	pass	0.5650	0.6068	0.9617	1.6005

id	name	rows	feat.	bytes	cap	CB n	CB $4n$	lin $4n$	MLP $4n$
0a8428ef25a69175	537_houses	800	8	2,609	pass	0.6257	0.7036	1.0150	1.0256
0d482c1dd7cc6592	210_cloud	72	5	1,844	pass	1.0374	0.9772	1.0944	1.0239
0db59fd2579600d0	665_sleuth.case2002	98	6	1,962	pass	1.0371	1.1324	1.1345	—
0f86a714ec5f3ac6	1191_BNG_pbc	800	18	4,183	pass	0.8097	0.8761	0.9173	—
167578f5deba4977	547_no2	333	7	2,443	pass	0.8073	0.8513	1.0097	0.9898
19f2c3c2d67defd7	feynman.II.11_27	800	4	1,494	pass	0.9704	0.9890	0.9960	1.0048
19f4780b53b3fa41	678_visualizing_environmental	74	3	1,515	pass	—	—	—	0.4431
1a4456d1cc71a26a	feynman.test.1	800	7	2,191	pass	1.0608	1.1671	0.9816	—
1a803b27c2400dc4	feynman.III.10_19	800	4	1,520	pass	0.9882	0.9952	0.9942	0.9965
1cb97b1d7314d43d	229_pwLinear	133	10	2,587	pass	0.8417	0.9878	0.9104	—
205f4b4982b9ea49	608_fri.c3_1000_10	667	10	2,845	pass	0.5263	0.5671	1.0451	0.8182
26eff843be35d831	581_fri.c3_500_25	333	25	5,647	pass	0.7602	0.9020	1.1297	1.5792
282f8ddb43b5e677	595_fri.c0_1000_10	667	10	2,904	pass	0.8974	0.9234	0.9962	1.0250
297b7688a045c583	603_fri.c0_250_50	167	50	10,171	pass	0.8848	1.1562	1.0700	2.6883
2c8d098f5357d2b	feynman.II.11_28	800	2	1,026	pass	0.9831	0.9899	1.0031	0.9976
30ca2887550f0d8d	201_pol	800	26	8,145	pass	0.5742	0.5976	0.7015	0.6923
3408a6a513c627a3	strogatz_barmag1	267	2	1,011	pass	0.9750	0.9775	0.9969	1.0058
35a9aa9b88f0e4d1	344_mv	800	10	2,737	pass	0.9985	1.0016	0.9962	1.0347
3d3321e27d14f9aa	599_fri.c2_1000_5	667	5	1,707	pass	0.8359	0.8611	0.9668	1.6797
44eaabe4cb923766	613_fri.c3_250_5	167	5	1,712	pass	0.8671	0.9276	0.8599	—
46b25db9396c8209	227_cpu_small	800	12	3,327	pass	0.8569	0.9072	0.9018	0.9364
46c5a59e5ae42037	1193_BNG_lowbwt	800	9	2,436	pass	0.9635	0.9703	1.0030	0.9871
471f11cc2d38853a	feynman.test.11	800	4	1,481	pass	0.9645	0.9760	0.9975	0.9826
47c581af5de1136b	582_fri.c1_500_25	333	25	5,670	pass	0.7510	0.9088	0.7173	—
48b7a3bf81e5f3e5	605_fri.c2_250_25	167	25	5,692	pass	0.4046	0.6465	1.9138	—
4a354d98c965b01f	597_fri.c2_500_5	333	5	1,713	pass	0.7550	0.8055	0.9814	1.8381
4aa84ff8ad0cd26	feynman.III.13_18	800	4	1,501	pass	0.8748	0.9137	0.8850	0.9348
51730b348b1bbd14	557_analcatdata.apnea1	316	3	1,237	pass	-0.0032	0.0086	—	—
54b221668673164c	feynman.III.14_14	800	5	1,726	pass	1.0346	1.1164	1.0011	1.0611
552665868f8f72ed	strogatz_vdp2	267	2	989	pass	0.9945	0.9963	0.9960	0.9969
585fef2a1cb6b33f	strogatz_lv2	267	2	1,018	pass	0.8286	0.9859	0.7502	—
642738a361f1c63a	feynman.I.12_1	800	2	1,032	pass	0.9892	0.9910	0.9892	0.9933
64c64f0f9983c48f	strogatz_bacres2	267	2	998	pass	0.9793	0.9832	0.9783	0.9862
654566a1f91408aa	4544_GeographicalOriginalofMusic	706	117	22,233	fail	—	—	—	—
693890aae9dc860b	first_principles_supernovae.zg	162	1	761	pass	0.8492	0.8908	1.0066	1.1997
6cacb11611345a24	feynman.I.11_19	800	6	1,947	pass	0.9524	0.9812	0.9780	0.9800
7291b347fe0238fb	feynman.II.11_20	800	5	1,757	pass	0.9928	1.0432	1.0016	1.0674
74237837fe38bb9b	strogatz_shearflow1	267	2	1,021	pass	0.4900	0.5614	-0.3978	—
7a3d123dd555110d	598_fri.c0_1000_25	667	25	5,668	pass	0.8613	0.9166	1.0046	1.0034
7c8943b160464add	601_fri.c1_250_5	167	5	1,719	pass	0.7859	0.8552	1.0773	—
7e51711baa1b1fb1	feynman.I.12_11	800	5	1,718	pass	0.9267	0.9320	1.0015	0.9860
80589368fa78f21a	218_house.8L	800	8	2,392	pass	0.5835	0.6708	0.9616	1.1040
8164950b155f7f56	feynman.I.12_2	800	4	1,511	pass	0.9702	0.9986	0.9485	0.9753
85a31ae409944a1b	feynman.II.10_9	800	3	1,269	pass	0.9330	0.9508	0.9589	0.9515
8d859cfab7b51a351	feynman.III.15_12	800	3	1,248	pass	0.5172	0.5542	1.0228	1.0749
923483635406b213	1203_BNG_pwLinear	800	10	2,581	pass	1.2079	1.1750	0.9941	1.5230
9312d813229df035	690_visualizing_galaxy	216	4	1,458	pass	0.9725	0.9755	0.9853	1.0870
94a00ea67f49bf84	596_fri.c2_250_5	167	5	1,698	pass	0.7546	0.8072	—	—
94bd4520b8ceb3f0	feynman.test.12	800	5	1,747	pass	0.9688	0.9804	0.9913	1.0067
956d34f13e7af830	586_fri.c3_1000_25	667	25	5,701	pass	0.9325	0.9849	1.2339	2.1457
9a10c15885022034	_deprecated.195_auto_price	106	15	5,025	pass	0.8717	0.9485	0.9598	1.3107
a04f964bc53281e0	strogatz_lv1	267	2	985	pass	-3.8784	-3.4212	—	—
a23f4bd626b76590	feynman.test.13	800	5	1,740	pass	0.9178	0.9605	0.9900	1.0991
a32a69e7c701eb14	588_fri.c4_1000_100	667	100	19,019	fail	—	—	—	—
a5ba682572f53add	602_fri.c3_250_10	167	10	2,900	pass	0.5355	0.8109	0.7988	—
a863cd42be4d05d4	663_rabe_266	80	2	1,014	pass	0.9781	0.9933	0.9967	1.0521
a97065df07a47eb0	592_fri.c4_1000_25	667	25	5,655	pass	0.8791	0.9340	0.9578	1.2714
ab279eded3b3c52b	574_house.16H	800	16	4,035	pass	0.1181	0.1796	0.1795	0.3705
aebf5d861a8365b4	522_pm10	333	7	2,409	pass	0.5082	0.6149	0.4125	2.3076
b6adf79118a58cea	562_cpu_small	800	12	3,353	pass	0.6968	0.7747	0.9834	0.9307
b78d323566930bd7	583_fri.c1_1000_50	667	50	10,098	pass	0.8607	0.9444	0.9960	1.6211
b7f9774e2142a758	strogatz_shearflow2	267	2	1,019	pass	0.4707	0.5439	0.7450	1.8955
bc85856c012bcc1b	607_fri.c4_1000_50	667	50	10,209	pass	0.8856	0.9640	3.9394	—
c107ee39eb333a9d	584_fri.c4_500_25	333	25	5,638	pass	0.8062	0.8619	0.9417	—
c30ff9405c916718	579_fri.c0_250_5	167	5	1,692	pass	0.9475	1.0923	0.9961	0.9783
c4f8e81b70d33a13	590_fri.c0_1000_50	667	50	10,180	pass	0.8839	0.9172	1.0359	1.1762
c649a4cee87c4321	529_pollen	800	4	1,506	pass	1.0356	1.1012	0.9717	1.2459
cb5bf66ef2ee09ed	nikuradse.1	241	2	992	pass	0.9346	0.9373	0.9997	1.1909
cba8393879afd1dc	230_machine_cpu	140	6	2,338	pass	1.0026	1.0891	1.1718	—
ce1a79e2d658eb8a	strogatz_glider2	267	2	1,003	pass	0.6382	0.6625	1.0720	1.3756
ce837ccd24a3a18c	strogatz_bacres1	267	2	1,004	pass	0.7352	0.8309	0.9014	0.9826
cee6f811d0bf199d	feynman.II.11_3	800	5	1,728	pass	0.9385	0.9666	0.8950	1.0149
d255d3a2e1c8dbf4	feynman.III.12_43	800	2	1,028	pass	0.9920	0.9935	0.9966	0.9940
d52bfd2b6ab1c850	feynman.I.10_7	800	3	1,250	pass	0.9920	0.9932	0.9940	0.9959
d9a3a408a840c9f48	560_bodyfat	168	14	4,364	pass	1.0247	1.0436	0.9930	1.3651
da290fae0d6c072f	feynman.test.10	800	3	1,267	pass	0.9955	0.9990	0.9990	1.0015
da78606950f100cd	1199_BNG_echoMonths	800	9	2,544	pass	0.9962	1.0086	0.9655	1.0694
dd16ac872ffb8673	strogatz_glider1	267	2	991	pass	0.1145	0.1723	0.9123	—
de2868ac837699fd	564_fried	800	10	2,861	pass	0.9477	0.9829	0.9543	1.4707

id	name	rows	feat.	bytes	cap	CB n	CB $4n$	lin $4n$	MLP $4n$
e08f0b9931f93686	225_puma8NH	800	8	2,439	pass	1.0639	1.1043	0.9951	2.8498
e18350fb51a9d1b0	strogatz_barmag2	267	2	1,027	pass	0.8897	0.9399	0.2855	—
e5a3a7b3668a1cc2	556_analcatdata_apnea2	316	3	1,238	pass	0.0339	0.0341	—	—
e68fdd96cc63271b	573_cpu_act	800	21	4,954	pass	0.8080	0.9045	0.6593	1.1465
ed2efb82fce4b098	strogatz_vdp1	267	2	984	pass	0.5228	0.5812	1.2867	1.2401
edc9d41a298fd60	612_fri.c1_1000_5	667	5	1,705	pass	0.8224	0.8596	0.9670	1.2891
efd53e9c0f36d608	_deprecated.207_autoPrice	106	15	5,017	pass	0.9020	0.9357	1.0899	1.0398
f0a3d6881458f990	604_fri.c4_500_10	333	10	2,913	pass	0.6868	0.8360	0.8552	—
f24e9383331c43e2	503_wind	800	14	3,846	pass	0.9690	0.9962	1.0017	1.0367
f652ea96eff62654	strogatz_predprey1	267	2	989	pass	0.9964	1.1345	1.0218	1.4876
f8cf362e371e8960	strogatz_predprey2	267	2	1,030	pass	0.9876	0.9847	0.9718	1.0815
fc3812567b5fdd64	197_cpu_act	800	21	5,011	pass	0.7680	0.8922	0.6669	1.1253
fe6d7ac3fcd3129e	1201_BNG_breastTumor	800	9	2,427	pass	-0.6141	-0.3437	—	—

J.2 SUFFICIENCY BUDGET (FIT SEED 11)

These rows are the `features12_steps8192` budget check on the lineages where it produced a charged byte count or a retention; the remaining lineages have neither and are omitted. The columns match the reported-profile table, so the two tables can be read side by side for the step-budget axis of Appendix H. The longer budget is a check on the reported profile and is not the reported model.

Table 41: Sufficiency budget. 69 lineages have a charged byte count or a retention. 31 lineages have neither and are omitted. The cap column is charged bytes at or under 10,240. It is not a release certificate. No Holm family.

id	name	rows	feat.	bytes	cap	CB n	CB $4n$	lin $4n$	MLP $4n$
047e62b52a3f4ede	695_chatfield_4	156	12	3,353	pass	1.0106	1.0287	1.0110	—
04c80944e971cb3e	first_principles_supernovae_zr	157	1	750	pass	0.9372	0.9454	1.0346	1.3271
060b9fff8092db17	589_fri.c2_1000_25	667	25	5,734	pass	0.9237	0.9886	1.0258	1.5430
06b0312a585b436e	1196_BNG_pharynx	800	10	2,681	pass	0.8188	0.8768	1.0632	0.7968
07f38946316c039b	593_fri.c1_1000_10	667	10	2,869	pass	0.9780	1.0222	0.9929	1.8097
0a8428ef25a69175	537_houses	800	8	2,636	pass	0.7243	0.7841	1.0073	1.0334
0d482c1dd7cc6592	210_cloud	72	5	1,855	pass	0.9561	0.8815	1.0056	0.9224
0f86a714ec5f3ac6	1191_BNG_pbc	800	18	4,218	pass	0.8069	0.8576	0.8783	—
167578f5deba4977	547_no2	333	7	2,474	pass	0.8821	0.9276	0.9850	1.0276
1a4456d1cc71a26a	feynman_test_1	800	7	2,185	pass	1.1162	1.3026	0.6721	—
1a803b27c2400dc4	feynman_III_10_19	800	4	1,522	pass	0.9698	0.9777	0.9732	0.9807
1cb97b1d7314d43d	229_pwLinear	133	10	2,598	pass	0.8275	1.0471	0.9727	—
26eff843be35d831	581_fri.c3_500_25	333	25	5,655	pass	0.8660	0.9960	1.0900	1.6067
282f8db43b5e677	595_fri.c0_1000_10	667	10	2,913	pass	0.9645	1.0058	1.0029	1.1269
297b7688a045c583	603_fri.c0_250_50	167	50	10,183	pass	0.8257	1.1292	1.0240	2.8999
30ca2887550fd08d	201_pol	800	26	8,215	pass	0.6214	0.6436	0.7350	0.6603
3d3321e27d14f9aa	599_fri.c2_1000_5	667	5	1,713	pass	0.9820	1.0044	1.0441	1.6148
44eaabe4cb923766	613_fri.c3_250_5	167	5	1,708	pass	0.8605	0.9611	0.9427	—
46b25db9396c8209	227_cpu_small	800	12	3,338	pass	0.7033	0.8730	0.9142	0.9908
46c5a59e5ae42037	1193_BNG_lowbwt	800	9	2,468	pass	0.8393	0.8390	0.8798	0.8103
471f11cc2d38853a	feynman_test_11	800	4	1,472	pass	0.9340	0.9481	0.9730	0.9578
47c581af5de1136b	582_fri.c1_500_25	333	25	5,676	pass	0.8232	0.9820	0.8086	—
48b7a3bf81e5f3e5	605_fri.c2_250_25	167	25	5,696	pass	0.2608	0.5247	1.4268	—
4a354d98c965b01f	597_fri.c2_500_5	333	5	1,715	pass	0.9578	1.0001	1.0224	1.3200
51730b348b1bbd14	557_analcatdata_apnea1	316	3	1,246	pass	0.6556	0.6654	—	—
54b221668673164c	feynman_III_14_14	800	5	1,724	pass	1.0175	1.1226	1.0021	1.0727
585fef2a1cb6b33f	strogatz_lv2	267	2	1,021	pass	0.9239	1.0451	0.8041	—
642738a361f1c63a	feynman_III_12_1	800	2	1,034	pass	0.9936	0.9961	0.9993	0.9989
64c64f0f9983c48f	strogatz_bacres2	267	2	996	pass	0.9743	0.9826	0.9816	0.9923
693890aae9dc860b	first_principles_supernovae_zg	162	1	766	pass	0.9415	0.9642	1.0277	1.1876
6cacb11611345a24	feynman_III_11_19	800	6	1,952	pass	0.9911	1.0176	1.0047	1.0089
7423783f7e38bb9b	strogatz_shearflow1	267	2	1,023	pass	0.4860	0.5973	-4.3352	—
7a3d123dd555110d	598_fri.c0_1000_25	667	25	5,672	pass	0.8570	0.9206	1.0010	1.0168
7e51711baa1bf1b1	feynman_III_12_11	800	5	1,720	pass	0.9780	0.9927	0.9910	1.0559
80589368fa78f21a	218_house_8L	800	8	2,410	pass	0.6661	0.7528	0.6830	1.0728
85a31ae40994a1b	feynman_III_10_9	800	3	1,264	pass	0.9803	0.9961	0.9881	1.0061
8d859cfa7b51a351	feynman_III_15_12	800	3	1,280	pass	0.7428	0.8572	0.8771	1.2306
9312d813229df035	690_visualizing_galaxy	216	4	1,470	pass	0.9770	0.9875	0.9843	1.0661
94a00ea67f49bf84	596_fri.c2_250_5	167	5	1,718	pass	0.8262	0.8934	—	—
94bd4520b8ceb3f0	feynman_test_12	800	5	1,745	pass	0.9677	0.9845	0.9937	1.0157
9a10c15885022034	_deprecated.195_auto_price	106	15	5,038	pass	0.8965	0.9544	0.9901	1.3165
a04f964bc53281e0	strogatz_lv1	267	2	986	pass	-3.0531	-2.4502	—	—
a32a69e7c701eb14	588_fri.c4_1000_100	667	100	19,041	fail	—	—	—	—
a5ba682572f53add	602_fri.c3_250_10	167	10	2,902	pass	0.7774	0.9773	0.9471	—
a863cd42be4d05d4	663_rabe_266	80	2	1,014	pass	0.9849	0.9996	1.0018	1.0573
ab279ded3b3c52b	574_house_16H	800	16	4,076	pass	0.0805	0.1211	-0.7472	-0.3199

id	name	rows	feat.	bytes	cap	CB n	CB $4n$	lin $4n$	MLP $4n$
aebf5d861a8365b4	522_pm10	333	7	2,450	pass	0.6372	0.7959	0.6203	3.4124
b78d323566930bd7	583_fri_c1_1000_50	667	50	10,098	pass	0.8817	0.9737	0.8733	0.9465
b7f9774e2142a758	strogatz_shearflow2	267	2	1,022	pass	0.7388	0.8339	-0.4873	1.5681
c107ee39eb333a9d	584_fri_c4_500_25	333	25	5,641	pass	0.8307	0.9024	0.9623	—
c30ff9405c916718	579_fri_c0_250_5	167	5	1,700	pass	0.9081	1.0750	0.9949	0.9830
c649a4cee87c4321	529_pollen	800	4	1,506	pass	1.0619	1.1198	0.9983	1.2801
cb5bf66ef2ee09ed	nikuradse_1	241	2	990	pass	0.9662	0.9677	0.9970	1.1953
ce1a79e2d658eb8a	strogatz_glider2	267	2	1,004	pass	0.5396	0.6099	1.1474	1.2017
ce837ccd24a3a18c	strogatz_bacres1	267	2	1,004	pass	0.7807	0.8568	0.8858	0.9570
d255d3a2e1c8dbf4	feynman.III.12_43	800	2	1,023	pass	0.9947	0.9964	0.9984	0.9967
d52bfd2b6ab1c850	feynman.L.10_7	800	3	1,248	pass	0.9978	0.9988	0.9994	1.0005
da290fae0d6c072f	feynman_test_10	800	3	1,257	pass	0.9681	0.9718	0.9737	0.9744
da78606950f100cd	1199_BNG_echoMonths	800	9	2,578	pass	0.9473	0.9673	0.9447	1.0028
de2868ac837699fd	564_fried	800	10	2,864	pass	0.9826	1.0201	0.9548	1.5094
e08f0b9931f93686	225_puma8NH	800	8	2,440	pass	1.0168	1.0664	0.9433	2.7529
e5a3a7b3668a1cc2	556_analcatdata_apnea2	316	3	1,241	pass	0.6282	0.6707	—	—
e68fdd96cc63271b	573_cpu_act	800	21	4,958	pass	0.8507	0.9689	0.5661	1.1177
ed2efb82fce4b098	strogatz_vdp1	267	2	991	pass	0.5397	0.6472	-0.1705	1.1154
efd53e9c0f36d608	_deprecated_207_autoPrice	106	15	5,024	pass	0.9035	0.9215	1.0781	1.0481
f0a3d6881458f990	604_fri_c4_500_10	333	10	2,918	pass	0.8895	1.0064	1.0421	—
f24e9383331c43e2	503_wind	800	14	3,877	pass	0.9992	0.9953	0.9918	1.0034
f8cf362e371e8960	strogatz_predprey2	267	2	1,030	pass	0.9865	0.9827	0.9658	1.0750
fc3812567b5fdd64	197_cpu_act	800	21	5,006	pass	0.8164	0.9529	0.6632	1.1726

J.3 BEYONDARENA FAMILIES IN THE DATASET BUCKET

The inventory below lists every BeyondArena family name in the dataset bucket with its task, row count, feature count, recorded license, and split type. Only the families with a cleared license entered the fit check of Appendix I, and none of them was scored for retention. The inventory is a list of candidate tables for later panels and makes no claim about any of them.

Table 42: All 142 BeyondArena families. No Holm family.

name	task	rows	cols	license	split
5g_energy_consumption	regression	92,629	22	MIT	grouped
acquire_valued_shoppers_challenge	binary	160,057	112	not cleared	temporal
airfoil_self_noise	regression	1,503	6	CC-BY-4.0	iid
allstate_claims_severity	regression	188,317	131	not cleared	iid
amazon_employee_access	binary	32,769	10	public-domain	iid
amex_non_iid	binary	1,249,605	191	not cleared	grouped
anes_voting_2026	binary	48,587	319	not cleared	temporal
aps_failure	binary	76,000	171	CC-BY-4.0	iid
asp_potassco_classification	multiclass	1,212	138	not cleared	grouped
audiology_diagnosis	multiclass	199	69	CC-BY-4.0	iid
bad_customer_detection	binary	1,723	14	public-domain	iid
bank_customer_churn	binary	10,000	11	public-domain	iid
bank_marketing	binary	45,211	14	CC-BY-4.0	iid
biogeographical_ancestry_prediction	multiclass	635	105	not cleared	iid
biomechanical_orthopaedic_prediction	multiclass	310	7	CC-BY-4.0	iid
bioresponse	binary	3,751	1777	public-domain	iid
blood_tests_drink_prediction	regression	345	6	CC-BY-4.0	iid
blood_transfusion	binary	748	5	CC-BY-4.0	iid
body_density_prediction	regression	252	14	not cleared	iid
california_house_prices_2020	regression	41,528	42	not cleared	temporal
cardiotocography	multiclass	2,126	24	CC-BY-4.0	grouped
churn	binary	5,000	20	public-domain	iid
cirrhosis_patient_survival_prediction	regression	161	18	not cleared	iid
climate_model_weather_forecasting	regression	1,250,000	101	not cleared	temporal
clock_protein_toxicity	binary	171	1118	CC-BY-4.0	iid
coffee_rating_prediction	regression	2,369	13	not cleared	temporal
coil_2000	binary	9,822	86	CC-BY-4.0	iid
concrete_compressive_strength	regression	1,030	9	CC-BY-4.0	iid
consumer_complaints	multiclass	1,226,140	13	us-gov	temporal
cooking_time	regression	1,250,000	197	not cleared	temporal
covertype	multiclass	512,625	15	CC-BY-4.0	grouped
credit_approval	binary	690	16	CC-BY-4.0	iid
credit_card_clients_default	binary	30,000	24	CC-BY-4.0	iid
credit_g	binary	1,000	21	CC-BY-4.0	iid
customer_satisfaction_in_airline	binary	129,880	22	public-domain	iid
delivery_eta	regression	1,250,000	226	not cleared	temporal
dementia_prediction	multiclass	370	10	not cleared	grouped
diabetes_130_us	binary	71,518	45	CC-BY-4.0	iid
diamonds	regression	53,940	10	MIT	iid
drug_induced_autoimmunity_prediction	binary	597	178	CC-BY-4.0	iid

Table 42, continued

name	task	rows	cols	license	split
early_learning_predictors	regression	18,874	745	CC-BY-4.0	grouped
early_stage_diabetes_risk_prediction	binary	251	17	CC-BY-4.0	iid
ecoli_proteins	multiclass	327	7	CC-BY-4.0	iid
ecommerce_shipping	binary	10,999	11	public-domain	iid
electric_motor_temperature_prediction	regression	1,296,316	111	not cleared	grouped
emscad	binary	17,460	18	CC0-1.0	iid
eryhemato_squamous_disease	multiclass	366	35	CC-BY-4.0	iid
fiat_500	regression	1,538	8	CC0-1.0	iid
fitness_club	binary	1,500	7	public-domain	iid
food_delivery_time	regression	45,451	10	not cleared	iid
forensic_glass_identification	multiclass	214	10	CC-BY-4.0	iid
forest_fires	regression	517	13	CC-BY-4.0	iid
gallstone_disease	binary	319	39	CC-BY-4.0	iid
garments_worker_productivity	regression	1,197	16	CC-BY-4.0	temporal
ghanas_indigenous_intel	multiclass	10,928	11	not cleared	temporal
give_me_some_credit	binary	150,000	11	not cleared	iid
hazelnut_spread_contaminant_detection	binary	2,400	31	not cleared	iid
healthcare_insurance_expenses	regression	1,338	7	not cleared	iid
heart_disease_cleveland	binary	303	14	CC-BY-4.0	iid
heart_disease_hungary	binary	294	14	CC-BY-4.0	iid
heart_disease_va_long_beach	binary	200	14	CC-BY-4.0	iid
heart_failure_followup_survival	binary	299	13	CC-BY-4.0	iid
heloc	binary	10,459	24	not cleared	iid
hepatitis_c_prediction	multiclass	608	13	CC-BY-4.0	iid
hepatitis_survival_prediction	binary	155	20	CC-BY-4.0	iid
hiva_agnostic	binary	3,845	1519	public-domain	iid
home_credit_default_risk	binary	307,507	505	not cleared	iid
home_credit_default_stability	binary	1,224,927	712	not cleared	temporal
homeq_default_prediction	binary	5,708	13	not cleared	iid
homesite_quote_conversion	binary	260,753	296	not cleared	iid
horse_colic_survival	multiclass	344	21	CC-BY-4.0	iid
hotel_booking_demand	binary	81,418	32	CC-BY-4.0	temporal
houses	regression	19,675	9	not cleared	iid
hr_analytics	binary	19,158	13	CC0-1.0	iid
ieee_fraud_detection	binary	590,540	436	not cleared	temporal
immoscout_german_house_prices	regression	10,317	24	not cleared	iid
in_vehicle_coupon_recommendation	binary	12,684	25	CC-BY-4.0	iid
indian_liver_patient_dataset	binary	583	11	CC-BY-4.0	iid
iranian_churn	binary	2,850	14	CC-BY-4.0	iid
jml	binary	10,885	22	public-domain	iid
kdd_cup_09_appetency	binary	50,000	213	public-domain	iid
kick	binary	72,983	33	public-domain	temporal
kickstarter	binary	187,118	16	not cleared	temporal
labour_inspection_compliance	binary	63,634	377	CC0-1.0	iid
lending_club	binary	1,064,751	97	CC0-1.0	temporal
ljubljana_breast_cancer	binary	286	10	CC-BY-4.0	iid
ljubljana_primary_tumor	multiclass	302	18	CC-BY-4.0	iid
lung_cancer	multiclass	197	12601	not cleared	iid
lung_cancer_epithelial_genexp	binary	187	22216	public-domain	iid
maps_router_eta	regression	1,250,000	989	not cleared	temporal
marketing_campaign	binary	2,240	26	CC-BY-4.0	iid
maternal_health_risk	multiclass	1,014	7	CC-BY-4.0	iid
mercari_price_suggestion	regression	1,250,000	7	not cleared	iid
mercedes_benz_greener_manufacturing	regression	4,204	372	not cleared	temporal
miami_housing	regression	13,776	16	not cleared	iid
mic	multiclass	1,699	112	CC-BY-4.0	iid
mice_protein_trisomy_discriminant	multiclass	1,080	78	CC-BY-4.0	grouped
micro_mass	multiclass	571	1084	CC-BY-4.0	grouped
musk	binary	6,598	168	CC-BY-4.0	grouped
mutual_funds_india	regression	793	13	CC0-1.0	iid
naticusdroid_android_permissions_dataset	binary	7,491	86	CC-BY-4.0	iid
obesity_estimation	regression	498	15	CC-BY-4.0	iid
online_shoppers_purchasing_intention_dataset	binary	12,330	18	CC-BY-4.0	iid
otto_group_product_classification_challenge	multiclass	61,878	94	not cleared	iid
pancreatic_cancer_mouse_detection	binary	181	6773	not cleared	grouped
parkinsons_biomedical_voice_measurements	binary	195	25	CC-BY-4.0	grouped
physiochemical_protein	regression	45,730	10	CC-BY-4.0	iid
polish_companies_bankruptcy	binary	5,790	65	CC-BY-4.0	iid
porto_seguro	binary	595,206	38	not cleared	iid
predict_students_dropout_and_academic_success	multiclass	4,424	37	CC-BY-4.0	iid
prostate_cancer_detection	binary	322	15155	not cleared	iid
pva_revenue_prediction_kddcup98	binary	144,095	478	not cleared	iid
qsar_aquatic_toxicity	regression	546	9	CC-BY-4.0	iid
qsar_biodeg	binary	1,054	42	CC-BY-4.0	iid
qsar_fish_toxicity	regression	908	7	CC-BY-4.0	iid
qsar_tid_11	regression	5,741	1025	public-domain	iid
regensburg_pediatric_appendicitis	binary	763	52	not cleared	iid
rossmann_store_sales	regression	844,392	16	not cleared	temporal

Table 42, continued

name	task	rows	cols	license	split
santander_customer_satisfaction	binary	71,080	308	not cleared	iid
santander_customer_transaction_prediction	binary	200,000	601	not cleared	iid
santander_transaction_value	regression	4,447	541	not cleared	iid
sat11_hand_algo_runtime	regression	2,960	171	not cleared	grouped
sberbank_housing_market_forecasting	regression	27,195	387	not cleared	temporal
sdss_17	multiclass	78,053	12	public-domain	iid
seismic_bumps	binary	2,584	16	CC-BY-4.0	iid
sepsis_prediction	binary	1,228,686	44	not cleared	grouped
sepsis_survival_minimal_clinical_records	binary	110,204	4	CC-BY-4.0	iid
sf_permit_time	regression	116,954	38	not cleared	temporal
south_africa_coronary_heart_disease	binary	462	10	not cleared	iid
splice	multiclass	3,190	61	CC-BY-4.0	iid
student_portuguese_performance	regression	649	31	CC-BY-4.0	iid
superconductivity	regression	21,263	82	CC-BY-4.0	iid
taiwanese_bankruptcy_prediction	binary	6,819	93	CC-BY-4.0	iid
telemonitoring_parkinsons_biomedical_voice_measurements	regression	502	21	CC-BY-4.0	grouped
thyroid_discordant	binary	3,711	27	CC-BY-4.0	iid
tour_travels_churn	binary	954	7	CC0-1.0	iid
video_game_fps_prediction	regression	12,288	40	CC-BY-4.0	grouped
video_transcoding_time_prediction	regression	68,784	20	CC-BY-4.0	grouped
website_phishing	multiclass	1,353	10	CC-BY-4.0	iid
wids_diabetes_mellitus	binary	127,358	182	not cleared	iid
wine_quality	regression	6,497	13	CC-BY-4.0	iid
wine_world_cost	regression	1,279	15	CC0-1.0	iid

K EXTENDED RELATED WORK

Generators by family. Parametric copulas (Patki et al., 2016) and Chow–Liu trees (Chow & Liu, 1968) are the classical density models, and pair-copula constructions (Aas et al., 2009) generalize them. Conditional GANs and variational autoencoders form a second family: CTGAN, TVAE, TableGAN, CTAB-GAN, and CTAB-GAN+ (Xu et al., 2019; Park et al., 2018; Zhao et al., 2021; 2024). Diffusion and score-based models form a third: TabDDPM, TabSyn, STASy, CoDi, TabDiff, CDTD, Forest-Diffusion, and GATD (Kotelnikov et al., 2023; Zhang et al., 2024; Kim et al., 2023; Lee et al., 2023; Shi et al., 2025; Mueller et al., 2025; Jolicoeur-Martineau et al., 2024; Zagardo, 2026). Autoregressive and language-model generators, and prior-data fitted generators, form a fourth (Borisov et al., 2023; Solatorio & Dupriez, 2023; Ma et al., 2023; Hollmann et al., 2025). Graphical neural generators are separate (Liu et al., 2023); ARF estimates density with a forest (Watson et al., 2023); CART synthesis is the synthpop line (Nowok et al., 2016); SMOTE interpolates between neighbouring rows (Chawla et al., 2002); and a survey collects these design choices (Borisov et al., 2024). Probabilistic-program synthesis (Saad et al., 2019) and CrossCat (Mansinghka et al., 2016) produce small, inspectable joint models.

How DOPE differs from each line. A parametric copula or a Chow–Liu tree optimizes a joint density of a fixed form; DOPE adds an explicit model of the target’s conditional mean and charges the result against a byte cap. A conditional GAN optimizes adversarial agreement and a VAE reconstruction plus a latent regularizer; both ship weights orders of magnitude larger than DOPE’s cap. A diffusion model optimizes a denoising or score objective over the joint table; DOPE optimizes the target’s conditional mean directly. A language-model generator optimizes serialized-row likelihood with a pretrained backbone; DOPE stores a screened residual of fixed width. A forest density model such as ARF optimizes axis-aligned splits of the joint table; CART synthesis optimizes sequential conditional draws, which also model the target given the features but store the trees. A prior-data fitted network optimizes a pretraining objective on other tables, whereas DOPE is fit on the lineage alone. A probabilistic-program synthesizer targets posterior inference over the joint distribution. A differentially private synthesizer optimizes under a privacy budget; DOPE’s privacy evidence is empirical, and its release gate is the MFS contract.

What this paper measured. The measured generators are the Gaussian copula, a Chow–Liu baseline and independent marginals written for this comparison, CTGAN, TVAE, ARF at five fit seeds, TabSyn on the full panel with a scaled schedule, and a historical Forest-Flow block; TabDDPM, TabDiff, a language-model generator, synthpop CART, and SMOTE are scheduled under a common budget in predeclare_v2. This panel scores regression retention.

Table 30: **Grouped-validation fidelity and empirical privacy, fit seed 11.** The fidelity denominator is the complete three-seed group count, distinct from the 97 retention lineages. Holm families are 12 density tests and 8 neural tests, separate from retention.

Method	α -prec. <i>Median [interval] (lineages).</i>	β -recall	DCR share (unbalanced)	DOMIAS
DOPE	0.905 [0.889, 0.918] (98)	0.437 [0.398, 0.473] (98)	0.811 [0.806, 0.814] (98)	0.508 [0.498, 0.520] (94)
Gaussian	0.812 [0.786, 0.832] (100)	0.430 [0.378, 0.453] (100)	0.806 [0.802, 0.811] (100)	0.513 [0.494, 0.527] (94)
Chow-Liu	0.858 [0.825, 0.904] (100)	0.359 [0.312, 0.400] (100)	0.810 [0.806, 0.816] (100)	0.512 [0.495, 0.522] (94)
Indep.	0.885 [0.868, 0.897] (100)	0.284 [0.197, 0.347] (100)	0.801 [0.797, 0.806] (100)	0.507 [0.496, 0.521] (94)
DOPE	0.902 [0.882, 0.940] (22)	0.441 [0.381, 0.489] (22)	0.811 [0.802, 0.819] (22)	0.508 [0.480, 0.538] (22)
CTGAN	0.719 [0.520, 0.856] (22)	0.170 [0.142, 0.267] (22)	0.820 [0.800, 0.842] (22)	0.502 [0.463, 0.534] (22)
TVAE	0.850 [0.731, 0.896] (22)	0.356 [0.283, 0.401] (22)	0.825 [0.813, 0.843] (22)	0.527 [0.487, 0.563] (22)

Contrast	α -prec.	β -recall	DCR share (unbalanced)	DOMIAS
<i>Line 2: Holm p; rank-biserial r; n. DOPE minus comparator.</i>				
	0.083 [0.045, 0.113]	0.024 [0.011, 0.039]	-0.0025 [-0.00625, 0.005]	0.00157 [-0.0069, 0.0074]
DOPE-Gaussian	1.86×10^{-10} ; 0.781; 98 0.019 [0.00782, 0.056]	0.394; 0.215; 98 0.048 [0.029, 0.072]	1.000; 0.000206; 98	1.000; -0.085; 94
DOPE-Chow-Liu	2.00×10^{-5} ; 0.551; 98 0.018 [0.0072, 0.028]	1.32×10^{-10} ; 0.789; 98 0.104 [0.068, 0.138]	0.806; -0.169; 98 0.00625 [0.0025, 0.011]	0.000581 [-0.00802, 0.011] 1.000; 0.00572; 94
DOPE-Indep.	5.47×10^{-5} ; 0.523; 98	9.71×10^{-15} ; 0.937; 98	0.005; 0.398; 98	0.00244 [-0.00595, 0.015] 1.000; 0.080; 94
DOPE-CTGAN	0.168 [0.054, 0.368] 0.004; 0.794; 22 0.071 [0.00713, 0.141]	0.219 [0.131, 0.256] 3.81×10^{-6} ; 1.000; 22 0.062 [0.00546, 0.135]	-0.00812 [-0.029, 0.015] 0.532; -0.277; 22 -0.025 [-0.030, -0.00641]	-0.00028 [-0.014, 0.031] 0.610; 0.130; 22 -0.021 [-0.036, -0.00972]
DOPE-TVAE	0.040; 0.644; 22	0.069; 0.573; 22	0.059; -0.613; 22	0.108; -0.510; 22




Figure 7: **Logged replay of both step budgets before the readout refit.** The solid line is the training median and the dashed line the validation median; each band is a 95% lineage-bootstrap interval. The longer budget is not the reported model.

Utility ratios, fidelity, and privacy. Retention is a null-normalized TSTR/TRTR ratio; Esteban, Hyland, and Rättsch introduced train-on-synthetic, test-on-real and train-on-real, test-on-real (Esteban et al., 2017). Utility reviews separate marginal fidelity from downstream prediction (Snok et al., 2018; Raab et al., 2018); alpha-precision, beta-recall, SDMetrics, and synthcity are the usual fidelity reports (Alaa et al., 2022; Zhang et al., 2020; Qian et al., 2023). Differentially private Bayesian networks and marginal synthesizers optimize a privacy budget (Zhang et al., 2014; McKenna et al., 2021; 2022; Cai et al., 2021). Membership inference, shadow-model attacks, and holdout-referenced distances form the empirical privacy literature (Shokri et al., 2017; Hayes et al., 2018; van Breugel et al., 2023; Stadler et al., 2022; Platzer & Reutterer, 2021); formal differential privacy is a different claim from an empirical resemblance score (Dwork & Roth, 2014; Abadi et al., 2016).

Auditors and rank tests. Tree models remain strong tabular predictors, which is why CatBoost is one of the three auditors (Grinsztajn et al., 2022; Prokhorenkova et al., 2018; Schwartz-Ziv & Armon, 2022). TabZilla and TabArena show how brittle a single split can be (McElfresh et al., 2023; Erickson et al., 2026). The regression lineages come from PMLB (Olson et al., 2017;




Figure 8: **Earlier eight-lineage TabSyn subset: medians and marginal lineage-bootstrap intervals.** No Holm family. Privacy AUC is an empirical diagnostic.

Romano et al., 2022). Comparing several methods on the same lineages needs a rank test and a critical difference, or a Bayesian signed-rank alternative (Demšar, 2006; Benavoli et al., 2017).

REFERENCES

- Kjersti Aas, Claudia Czado, Arnoldo Frigessi, and Henrik Bakken. Pair-copula constructions of multiple dependence. *Insurance: Mathematics and Economics*, 44(2):182–198, 2009. doi: 10.1016/j.insmatheco.2007.02.001.
- Martin Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, pp. 308–318, 2016. doi: 10.1145/2976749.2978318.
- Ahmed Alaa, Boris Van Breugel, Evgeny S. Saveliev, and Mihaela van der Schaar. How faithful is your synthetic data? sample-level metrics for evaluating and auditing generative models. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pp. 290–306. PMLR, 2022. doi: 10.48550/ARXIV.2102.08921. URL <https://proceedings.mlr.press/v162/aaa22a.html>.
- Alessio Benavoli, Giorgio Corani, Janez Demšar, and Marco Zaffalon. Time for a change: a tutorial for comparing multiple classifiers through bayesian analysis. *Journal of Machine Learning Research*, 18(77):1–36, 2017. URL <https://jmlr.org/papers/v18/16-305.html>.
- Vadim Borisov, Kathrin Seßler, Tobias Leemann, Martin Pawelczyk, and Gjergji Kasneci. Language models are realistic tabular data generators. In *International Conference on Learning Representations*, 2023. doi: 10.48550/ARXIV.2210.06280. URL <https://openreview.net/forum?id=cEygmQNOeI>.
- Vadim Borisov, Tobias Leemann, Kathrin Seßler, Johannes Haug, Martin Pawelczyk, and Gjergji Kasneci. Deep neural networks and tabular data: A survey. *IEEE Transactions on Neural Networks and Learning Systems*, 35(6):7499–7519, 2024. doi: 10.1109/tnnls.2022.3229161.
- Kuntai Cai, Xiaoyu Lei, Jianxin Wei, and Xiaokui Xiao. Data synthesis via differentially private markov random fields. *Proceedings of the VLDB Endowment*, 14(11):2190–2202, 2021. doi: 10.14778/3476249.3476272.
- N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer. SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16:321–357, 2002. doi: 10.1613/jair.953.

- C. K. Chow and C. N. Liu. Approximating discrete probability distributions with dependence trees. *IEEE Transactions on Information Theory*, 14(3):462–467, May 1968. doi: 10.1109/TIT.1968.1054142.
- Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7(1):1–30, 2006. URL <https://jmlr.org/papers/v7/demsar06a.html>.
- Cynthia Dwork and Aaron Roth. The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4):211–487, 2014. doi: 10.1561/04000000042.
- Nick Erickson, Lennart Purucker, Andrej Tschalzev, David Holzmüller, Prateek Mutalik Desai, David Salinas, and Frank Hutter. TabArena: A living benchmark for machine learning on tabular data. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, pp. 17285–17350, 2026. doi: 10.52202/085713-0519. URL <https://openreview.net/forum?id=jZqCqpCLdU>.
- Cristóbal Esteban, Stephanie L. Hyland, and Gunnar Rätsch. Real-valued (medical) time series generation with recurrent conditional GANs. *arXiv preprint arXiv:1706.02633*, 2017. doi: 10.48550/ARXIV.1706.02633.
- Leo Grinsztajn, Edouard Oyallon, and Gael Varoquaux. Why do tree-based models still outperform deep learning on typical tabular data? In *Advances in Neural Information Processing Systems*, volume 35, pp. 507–520, 2022. doi: 10.52202/068431-0037.
- Jamie Hayes, Luca Melis, George Danezis, and Emiliano De Cristofaro. LOGAN: Membership inference attacks against generative models. *Proceedings on Privacy Enhancing Technologies*, 2019(1):133–152, 2018. doi: 10.2478/popets-2019-0008.
- Noah Hollmann, Samuel Müller, Lennart Purucker, Arjun Krishnakumar, Max Körfer, Shi Bin Hoo, Robin Tibor Schirrmeyer, and Frank Hutter. Accurate predictions on small data with a tabular foundation model. *Nature*, 637(8045):319–326, 2025. doi: 10.1038/s41586-024-08328-6.
- Alexia Jolicoeur-Martineau, Kilian Fatras, and Tal Kachman. Generating and imputing tabular data via diffusion and flow-based gradient-boosted trees. In *Proceedings of The 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, pp. 1288–1296. PMLR, 2024. doi: 10.48550/ARXIV.2309.09968. URL <https://proceedings.mlr.press/v238/jolicoeur-martineau24a.html>.
- Jayoung Kim, Chaejeong Lee, and Noseong Park. STaSy: Score-based tabular data synthesis. In *International Conference on Learning Representations*, 2023. doi: 10.48550/ARXIV.2210.04018. URL https://openreview.net/forum?id=lmNssCWt_v.
- Akim Kotelnikov, Dmitry Baranchuk, Ivan Rubachev, and Artem Babenko. TabDDPM: Modelling tabular data with diffusion models. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 17564–17579, 2023. doi: 10.48550/ARXIV.2209.15421. URL <https://proceedings.mlr.press/v202/kotelnikov23a.html>.
- Chaejeong Lee, Jayoung Kim, and Noseong Park. CoDi: Co-evolving contrastive diffusion models for mixed-type tabular synthesis. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 18940–18956. PMLR, 2023. doi: 10.48550/ARXIV.2304.12654. URL <https://proceedings.mlr.press/v202/lee23i.html>.
- Tennison Liu, Zhaozhi Qian, Jeroen Berrevoets, and Mihaela van der Schaar. GOGGLE: Generative modelling for tabular data by learning relational structure. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=fPVRcJqspu>.

- Junwei Ma, Apoorv Dankar, George Stein, Guangwei Yu, and Anthony Caterini. TabPFGGen – tabular data generation with TabPFN. In *NeurIPS 2023 Second Table Representation Learning Workshop*, 2023. URL <https://openreview.net/forum?id=4MkNsAEmO>.
- Vikash Mansinghka, Patrick Shafto, Eric Jonas, Cap Petschulat, Max Gasner, and Joshua B. Tenenbaum. CrossCat: A fully Bayesian nonparametric method for analyzing heterogeneous, high dimensional data. *Journal of Machine Learning Research*, 17(138):1–49, 2016. URL <https://jmlr.org/papers/v17/11-392.html>.
- Duncan McElfresh, Sujay Khandagale, Jonathan Valverde, Vishak Prasad C, Ganesh Ramakrishnan, Micah Goldblum, and Colin White. When do neural nets outperform boosted trees on tabular data? In *Advances in Neural Information Processing Systems*, volume 36, pp. 76336–76369, 2023. doi: 10.52202/075280-3337. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/f06d5ebd4ff40b40dd97e30cee632123-Abstract-Datasets_and_Benchmarks.html. Datasets and Benchmarks Track.
- Ryan McKenna, Gerome Miklau, and Daniel Sheldon. Winning the NIST contest: A scalable and general approach to differentially private synthetic data. *Journal of Privacy and Confidentiality*, 11(3), 2021. doi: 10.29012/jpc.778.
- Ryan McKenna, Brett Mullins, Daniel Sheldon, and Gerome Miklau. AIM: An adaptive and iterative mechanism for differentially private synthetic data. *Proceedings of the VLDB Endowment*, 15(11):2599–2612, 2022. doi: 10.14778/3551793.3551817.
- Markus Mueller, Kathrin Gruber, and Dennis Fok. Continuous diffusion for mixed-type tabular data. In *International Conference on Learning Representations*, 2025. doi: 10.48550/ARXIV.2312.10431. URL <https://openreview.net/forum?id=QPTobPn4lZ>.
- Beata Nowok, Gillian M. Raab, and Chris Dibben. synthpop: Bespoke creation of synthetic data in R. *Journal of Statistical Software*, 74(11):1–26, 2016. doi: 10.18637/jss.v074.i11.
- Randal S. Olson, William La Cava, Patryk Orzechowski, Ryan J. Urbanowicz, and Jason H. Moore. PMLB: a large benchmark suite for machine learning evaluation and comparison. *BioData Mining*, 10(1):36, 2017. doi: 10.1186/s13040-017-0154-4.
- Noseong Park, Mahmoud Mohammadi, Kshitij Gorde, Sushil Jajodia, Hongkyu Park, and Youngmin Kim. Data synthesis based on generative adversarial networks. *Proceedings of the VLDB Endowment*, 11(10):1071–1083, 2018. doi: 10.14778/3231751.3231757.
- Neha Patki, Roy Wedge, and Kalyan Veeramachaneni. The synthetic data vault. In *2016 IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, pp. 399–410, 2016. doi: 10.1109/DSAA.2016.49.
- Michael Platzer and Thomas Reutterer. Holdout-based empirical assessment of mixed-type synthetic data. *Frontiers in Big Data*, 4:679939, 2021. doi: 10.3389/fdata.2021.679939.
- Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. CatBoost: unbiased boosting with categorical features. In *Advances in Neural Information Processing Systems*, volume 31, 2018. doi: 10.48550/ARXIV.1706.09516. URL <https://papers.nips.cc/paper/2018/hash/14491b756b3a51daac41c24863285549-Abstract.html>.
- Zhaozhi Qian, Rob Davis, and Mihaela van der Schaar. Synthcity: A benchmark framework for diverse use cases of tabular synthetic data. In *Advances in Neural Information Processing Systems*, volume 36, pp. 3173–3188, 2023. doi: 10.52202/075280-0140. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/09723c9f291f6056fd1885081859c186-Abstract-Datasets_and_Benchmarks.html. Datasets and Benchmarks Track.
- Gillian M. Raab, Beata Nowok, and Chris Dibben. Practical data synthesis for large samples. *Journal of Privacy and Confidentiality*, 7(3):67–97, 2018. doi: 10.29012/jpc.v7i3.407.

- Joseph D. Romano, Trang T. Le, William La Cava, John T. Gregg, Daniel J. Goldberg, Praneel Chakraborty, Natasha L. Ray, Daniel Himmelstein, Weixuan Fu, and Jason H. Moore. PMLB v1.0: an open-source dataset collection for benchmarking machine learning methods. *Bioinformatics*, 38(3):878–880, 2022. doi: 10.1093/bioinformatics/btab727.
- Feras A. Saad, Marco F. Cusumano-Towner, Ulrich Schaechtle, Martin C. Rinard, and Vikash K. Mansinghka. Bayesian synthesis of probabilistic programs for automatic data modeling. *Proceedings of the ACM on Programming Languages*, 3(POPL):1–32, 2019. doi: 10.1145/3290350.
- Juntong Shi, Minkai Xu, Harper Hua, Hengrui Zhang, Stefano Ermon, and Jure Leskovec. TabDiff: a mixed-type diffusion model for tabular data generation. In *International Conference on Learning Representations*, 2025. doi: 10.48550/ARXIV.2410.20626. URL <https://openreview.net/forum?id=swvURjrt8z>.
- Reza Shokri, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. Membership inference attacks against machine learning models. In *2017 IEEE Symposium on Security and Privacy (SP)*, pp. 3–18, 2017. doi: 10.1109/sp.2017.41.
- Ravid Shwartz-Ziv and Amitai Armon. Tabular data: Deep learning is not all you need. *Information Fusion*, 81:84–90, 2022. doi: 10.1016/j.inffus.2021.11.011.
- Joshua Snok, Gillian M. Raab, Beata Nowok, Chris Dibben, and Aleksandra Slavkovic. General and specific utility measures for synthetic data. *Journal of the Royal Statistical Society Series A: Statistics in Society*, 181(3):663–688, 2018. doi: 10.1111/rssa.12358.
- Aivin V. Solatorio and Olivier Dupriez. REaLTabFormer: Generating realistic relational and tabular data using transformers, 2023.
- Theresa Stadler, Bristena Oprisanu, and Carmela Troncoso. Synthetic data – anonymisation groundhog day. In *31st USENIX Security Symposium (USENIX Security 22)*, pp. 1451–1468, 2022. URL <https://www.usenix.org/conference/usenixsecurity22/presentation/stadler>.
- Boris van Breugel, Hao Sun, Zhaozhi Qian, and Mihaela van der Schaar. Membership inference attacks against synthetic data through overfitting detection. In *Proceedings of The 26th International Conference on Artificial Intelligence and Statistics*, volume 206 of *Proceedings of Machine Learning Research*, pp. 3493–3514. PMLR, 2023. doi: 10.48550/ARXIV.2302.12580. URL <https://proceedings.mlr.press/v206/breugel23a.html>.
- David S. Watson, Kristin Blesch, Jan Kapar, and Marvin N. Wright. Adversarial random forests for density estimation and generative modeling. In *Proceedings of the 26th International Conference on Artificial Intelligence and Statistics*, volume 206 of *Proceedings of Machine Learning Research*, pp. 5357–5375, 2023. doi: 10.48550/ARXIV.2205.09435. URL <https://proceedings.mlr.press/v206/watson23a.html>.
- Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. Modeling tabular data using conditional GAN. In *Advances in Neural Information Processing Systems*, volume 32, 2019. doi: 10.48550/ARXIV.1907.00503. URL <https://papers.nips.cc/paper/2019/hash/254ed7d2de3b23ab10936522dd547b78-Abstract.html>.
- David Turtora Zagardo. Geometry-aware tabular diffusion. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pp. 152812–152835. PMLR, 2026. doi: 10.48550/ARXIV.2606.02607. URL <https://proceedings.mlr.press/v306/zagardo26a.html>.
- Hengrui Zhang, Jiani Zhang, Balasubramaniam Srinivasan, Zhengyuan Shen, Xiao Qin, Christos Faloutsos, Huzefa Rangwala, and George Karypis. Mixed-type tabular data synthesis with score-based diffusion in latent space. In *International Conference on Learning Representations*, 2024. doi: 10.48550/ARXIV.2310.09656. URL <https://openreview.net/forum?id=4Ay23yeuz0>.

- Jun Zhang, Graham Cormode, Cecilia M. Procopiuc, Divesh Srivastava, and Xiaokui Xiao. PrivBayes: Private data release via bayesian networks. In *Proceedings of the 2014 ACM SIGMOD International Conference on Management of Data*, pp. 1423–1434, 2014. doi: 10.1145/2588555.2588573.
- Kevin Alex Zhang, Kalyan Veeramachaneni, and Carles Sala. SDMetrics: Synthetic data metrics. Open-source library, MIT Data To AI Lab, 2020. URL <https://docs.sdv.dev/sdmetrics/>.
- Zilong Zhao, Aditya Kunar, Robert Birke, and Lydia Y. Chen. CTAB-GAN: Effective table data synthesizing. In *Proceedings of The 13th Asian Conference on Machine Learning*, volume 157 of *Proceedings of Machine Learning Research*, pp. 97–112. PMLR, 2021. URL <https://proceedings.mlr.press/v157/zhao21a.html>.
- Zilong Zhao, Aditya Kunar, Robert Birke, Hiek Van der Scheer, and Lydia Y. Chen. CTAB-GAN+: Enhancing tabular data synthesis. *Frontiers in Big Data*, 6, 2024. doi: 10.3389/fdata.2023.1296508.

Table 31: **Fidelity and detection by configuration and size.** Median, hierarchical interval, and lineage count in each measured cell. Common-numeric cohorts, with separate availability by metric. Stored logistic AUC and grouped CatBoost AUC have distinct columns. AUC alone establishes neither privacy nor release safety. No Holm family.

Method	Configuration	Size	Marginal KS	Pair fidelity	Logistic AUC	CatBoost AUC
ARF	author default	1n	0.1128	0.9595	0.5321	0.6159
			[0.0895, 0.1433]	[0.9471, 0.9644]	[0.5156, 0.5442]	[0.5751, 0.6928]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$
ARF	author default	4n	0.1035	0.9618	0.5268	0.6120
			[0.0846, 0.1340]	[0.9487, 0.9653]	[0.5069, 0.5388]	[0.5696, 0.6877]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$
ARF	native selected	1n	0.1122	0.9599	0.5255	0.6100
			[0.0869, 0.1445]	[0.9418, 0.9626]	[0.5072, 0.5449]	[0.5789, 0.7089]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$
ARF	native selected	4n	0.1046	0.9620	0.5296	0.6091
			[0.0845, 0.1382]	[0.9464, 0.9643]	[0.5094, 0.5375]	[0.5766, 0.7075]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$
CTGAN	native selected	1n	0.2319	0.9091	0.7839	0.9522
			[0.1887, 0.2943]	[0.8313, 0.9406]	[0.6997, 0.8819]	[0.8778, 0.9841]
			$n = 22$	$n = 22$	$n = 22$	$n = 21$
CTGAN	native selected	4n	0.2300	0.9123	0.7839	0.9535
			[0.1839, 0.2947]	[0.8367, 0.9431]	[0.6997, 0.8821]	[0.8487, 0.9777]
			$n = 22$	$n = 22$	$n = 22$	$n = 22$
Chow-Liu	native selected	1n	0.1293	0.9273	0.5423	0.6868
			[0.0892, 0.2361]	[0.8834, 0.9418]	[0.5040, 0.6292]	[0.6025, 0.8986]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$
Chow-Liu	native selected	4n	0.1266	0.9283	0.5404	0.7061
			[0.0852, 0.2292]	[0.8906, 0.9465]	[0.5185, 0.6303]	[0.5910, 0.9369]
			$n = 100$	$n = 100$	$n = 100$	$n = 99$
DOPE	DOPE profile	1n	0.1216	0.9532	0.5172	0.6344
			[0.0950, 0.1680]	[0.9317, 0.9610]	[0.5000, 0.5376]	[0.5705, 0.8368]
			$n = 98$	$n = 98$	$n = 98$	$n = 98$
DOPE	DOPE profile	4n	0.1138	0.9531	0.5172	0.6405
			[0.0917, 0.1583]	[0.9346, 0.9645]	[0.5000, 0.5376]	[0.5830, 0.8312]
			$n = 98$	$n = 98$	$n = 98$	$n = 98$
Forest-Flow	native selected	1n	0.1254	0.9548	0.5187	—
			[0.0612, 0.1992]	[0.9242, 0.9799]	[0.4139, 0.7000]	—
			$n = 6$	$n = 6$	$n = 6$	—
Forest-Flow	native selected	4n	0.1231	0.9691	0.5187	—
			[0.0557, 0.2062]	[0.9321, 0.9811]	[0.4139, 0.7000]	—
			$n = 6$	$n = 6$	$n = 6$	—
GaussianCopula	native selected	1n	0.1393	0.9614	0.5246	0.7343
			[0.1171, 0.2132]	[0.9516, 0.9652]	[0.5100, 0.5518]	[0.6789, 0.8981]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$
GaussianCopula	native selected	4n	0.1336	0.9635	0.5246	0.7346
			[0.1137, 0.2115]	[0.9488, 0.9687]	[0.5100, 0.5537]	[0.6902, 0.8689]
			$n = 100$	$n = 100$	$n = 100$	$n = 99$
TVAE	native selected	1n	0.1500	0.9240	0.5918	0.7479
			[0.1199, 0.2598]	[0.8905, 0.9418]	[0.5616, 0.6392]	[0.7221, 0.8859]
			$n = 22$	$n = 22$	$n = 22$	$n = 22$
TVAE	native selected	4n	0.1382	0.9258	0.6033	0.7309
			[0.1174, 0.2431]	[0.8913, 0.9451]	[0.5544, 0.6681]	[0.7032, 0.8893]
			$n = 22$	$n = 22$	$n = 22$	$n = 22$
TabSyn	scaled schedule	1n	0.1244	0.9543	—	0.6242
			[0.0828, 0.1416]	[0.9412, 0.9598]	—	[0.5734, 0.7373]
			$n = 99$	$n = 99$	—	$n = 99$
TabSyn	scaled schedule	4n	0.1196	0.9581	—	0.6168
			[0.0829, 0.1351]	[0.9458, 0.9631]	—	[0.5866, 0.7772]
			$n = 99$	$n = 99$	—	$n = 98$
Independent marginals	native selected	1n	0.1363	0.8994	0.5213	0.8694
			[0.0866, 0.2847]	[0.8099, 0.9289]	[0.5000, 0.5347]	[0.7602, 0.9749]
			$n = 100$	$n = 100$	$n = 100$	$n = 99$
Independent marginals	native selected	4n	0.1297	0.8986	0.5166	0.8578
			[0.0813, 0.2672]	[0.8258, 0.9329]	[0.4988, 0.5488]	[0.7553, 0.9727]
			$n = 100$	$n = 100$	$n = 100$	$n = 100$

Table 32: **Alpha-precision and beta-recall.** Fixed unembedded equations, median over complete sample triples, then hierarchical interval and lineage count. Configuration and cohort are explicit; unavailable cells carry no scalar. The retained TabSyn subset remains separate. No Holm family.

Method	Configuration	Size	Cohort	α precision	β recall
ARF	author default	1n	common numeric	0.8838	0.4269
				[0.8644, 0.9144]	[0.3237, 0.4727]
				$n = 100$	$n = 100$
ARF	author default	4n	common numeric	0.8786	0.4227
				[0.8447, 0.9103]	[0.3150, 0.4663]
				$n = 100$	$n = 100$
ARF	native selected	1n	common numeric	0.8713	0.4360
				[0.8502, 0.9013]	[0.3010, 0.4646]
				$n = 100$	$n = 100$
ARF	native selected	4n	common numeric	0.8767	0.4208
				[0.8465, 0.9069]	[0.3180, 0.4607]
				$n = 100$	$n = 100$
CTGAN	native selected	1n	common numeric	unavailable	unavailable
				$n = 0$	$n = 0$
				0.7193	0.1695
CTGAN	native selected	4n	common numeric	[0.4848, 0.8751]	[0.1140, 0.2667]
				$n = 22$	$n = 22$
				unavailable	unavailable
Chow-Liu	native selected	1n	common numeric	$n = 0$	$n = 0$
				0.8585	0.3595
				[0.7403, 0.9238]	[0.1793, 0.4230]
Chow-Liu	native selected	4n	common numeric	$n = 100$	$n = 100$
				unavailable	unavailable
				$n = 0$	$n = 0$
DOPE	DOPE profile	1n	common numeric	0.9051	0.4366
				[0.8751, 0.9254]	[0.2566, 0.4863]
				$n = 98$	$n = 98$
Forest-Flow	native selected	1n	common numeric	unavailable	unavailable
				$n = 0$	$n = 0$
				unavailable	unavailable
Forest-Flow	native selected	4n	common numeric	$n = 0$	$n = 0$
				unavailable	unavailable
				$n = 0$	$n = 0$
GaussianCopula	native selected	1n	common numeric	unavailable	unavailable
				$n = 0$	$n = 0$
				0.8125	0.4301
GaussianCopula	native selected	4n	common numeric	[0.7516, 0.8499]	[0.2173, 0.4593]
				$n = 100$	$n = 100$
				unavailable	unavailable
TVAE	native selected	1n	common numeric	$n = 0$	$n = 0$
				0.8496	0.3562
				[0.7411, 0.8966]	[0.2778, 0.3991]
TVAE	native selected	4n	common numeric	$n = 22$	$n = 22$
				0.8714	0.4807
				[0.8033, 0.9182]	[0.3965, 0.5177]
TabSyn	scaled schedule	1n	common numeric	$n = 99$	$n = 99$
				0.8418	0.4591
				[0.4700, 0.9326]	[0.3520, 0.5267]
TabSyn	scaled schedule	1n	retained subset	$n = 8$	$n = 8$
				0.8920	0.4633
				[0.7919, 0.9280]	[0.4100, 0.5242]
TabSyn	scaled schedule	4n	common numeric	$n = 99$	$n = 99$
				0.8262	0.4466
				[0.4467, 0.9352]	[0.3420, 0.5686]
TabSyn	scaled schedule	4n	retained subset	$n = 8$	$n = 8$
				unavailable	unavailable
				$n = 0$	$n = 0$
Independent marginals	native selected	1n	common numeric	0.8849	0.2837
				[0.8141, 0.9181]	[0.0949, 0.3963]
				$n = 100$	$n = 100$
Independent marginals	native selected	4n	common numeric	$n = 100$	$n = 100$

Table 33: **Baseline coverage from committed publications.** The denominator is the committed publication named in the row. No Holm family.

Method	Published measured coverage	Remaining campaign
TabSyn	49 scaled-default lineages, 143 measured cells at fit11: eight complete n/4n lineages (48 cells), plus 95 n cells on 41 lineages (19 complete/22 partial groups). Native tuning/full population pending.	Pending
TabDDPM	12 retained lineages; 126 physical/132 logical cells; 21 models, fit seed 11. Default/native cohorts differ; native search incomplete.	Pending
Forest-Flow	13 default lineages at five fit seeds; 65 models/390 cells. Separate legacy six-lineage default/native panel; full population pending.	Pending
Forest-Diffusion	Pilot only; separate 100-lineage matched diffusion panel pending.	Pending
ARF	100 default/native lineages at five fit seeds, n/4n, three sample seeds: 6,000 logical cells / 5,892 physical receipts (108 aliases). Between-fit SD available; official-test certification pending.	Pending
GaussianCopula	100 default/native lineages at fit seed 11; 1200 logical cells. Five-fit common coverage pending.	Pending
Chow-Liu	100 default/native lineages at fit seed 11; 1200 logical cells. Five-fit common coverage pending.	Pending

Table 34: **Retained classical validation at 4n.** Median [95% lineage-bootstrap interval] (informative lineages). No Holm family in this table.

Method	Config.	CB retention	KS/TV	C2ST CB AUC	Distance MIA AUC
ARF	Default	0.820 [0.787, 0.872] (99)	0.103 [0.097, 0.114] (100)	0.604 [0.582, 0.626] (100)	0.553 [0.538, 0.564] (100)
ARF	Native	0.806 [0.736, 0.855] (99)	0.105 [0.098, 0.114] (100)	0.610 [0.591, 0.627] (100)	0.542 [0.536, 0.559] (100)
Chow-Liu	Default	0.643 [0.552, 0.672] (99)	0.128 [0.106, 0.154] (100)	0.721 [0.652, 0.773] (100)	0.513 [0.503, 0.525] (100)
Chow-Liu	Native	0.566 [0.460, 0.631] (99)	0.126 [0.105, 0.157] (100)	0.702 [0.641, 0.761] (99)	0.513 [0.504, 0.523] (100)
Copula	Default	0.594 [0.427, 0.762] (99)	0.126 [0.094, 0.138] (100)	0.615 [0.591, 0.671] (100)	0.514 [0.503, 0.524] (100)
Copula	Native	0.641 [0.431, 0.803] (99)	0.134 [0.125, 0.158] (100)	0.730 [0.696, 0.779] (99)	0.507 [0.494, 0.521] (100)

Table 35: **Earlier eight-lineage TabSyn subset at 4n.** Same inputs and real-data controls for every method; median [95% interval] (lineages). A separate cohort from the full TabSyn panel. No Holm family in this table.

Method	Config.	CB retention	KS/TV	C2ST CB AUC	Distance MIA AUC
ARF	Default	0.839 [0.668, 0.942] (8)	0.129 [0.084, 0.165] (8)	0.623 [0.547, 0.969] (8)	0.555 [0.506, 0.613] (8)
ARF	Native	0.752 [0.534, 0.946] (8)	0.132 [0.084, 0.163] (8)	0.606 [0.583, 0.812] (8)	0.527 [0.482, 0.558] (8)
Chow-Liu	Default	0.746 [0.666, 0.859] (8)	0.186 [0.068, 0.255] (8)	0.756 [0.525, 0.999] (8)	0.522 [0.465, 0.542] (8)
Chow-Liu	Native	0.655 [0.561, 0.778] (8)	0.188 [0.087, 0.248] (8)	0.707 [0.525, 0.999] (8)	0.499 [0.470, 0.533] (8)
Copula	Default	0.643 [0.270, 0.935] (8)	0.146 [0.067, 0.242] (8)	0.687 [0.554, 0.987] (8)	0.510 [0.466, 0.541] (8)
Copula	Native	0.680 [0.260, 0.858] (8)	0.157 [0.097, 0.242] (8)	0.698 [0.642, 0.987] (8)	0.493 [0.463, 0.525] (8)
TabSyn	Scaled default	0.911 [0.685, 0.971] (8)	0.136 [0.120, 0.157] (8)	0.670 [0.597, 0.869] (8)	0.571 [0.547, 0.619] (8)

Table 36: **Exploratory eight-lineage subset differences, CatBoost.** The denominator is the shared lineages in the cell. Holm family size is four metrics times six configurations. The corrected simultaneous intervals are unbounded.

TabSyn minus	n	Median CB difference	Simultaneous median CI	Holm p
ARF/default	8	0.010	Unbounded	1.000
ARF/native	8	0.159	Unbounded	0.469
Chow-Liu/default	8	0.157	Unbounded	1.000
Chow-Liu/native	8	0.149	Unbounded	0.188
Copula/default	8	0.242	Unbounded	1.000
Copula/native	8	0.245	Unbounded	1.000

Table 37: **DOPE-ARF legacy matched receipt at 4n: exploratory paired differences with conservative simultaneous median intervals.** The denominator is the shared-lineage intersection for that auditor. Holm family size is six. Median-of-seed aggregation.

ARF config.	Auditor	n	DOPE minus ARF	Simultaneous median CI	Holm p
Default	catboost	97	0.077	[0.019, 0.154]	1.31×10^{-5}
Native	catboost	97	0.095	[0.028, 0.175]	5.47×10^{-7}
Default	linear	92	0.027	[0.0056, 0.090]	8.53×10^{-5}
Native	linear	92	0.037	[0.00582, 0.113]	1.62×10^{-5}
Default	mlp	75	0.115	[0.043, 0.236]	4.56×10^{-8}
Native	mlp	75	0.143	[0.040, 0.253]	8.01×10^{-8}

Table 38: **First-batch diagnostics, one row per physical job.** Individual evaluations without a population interval. No Holm family.

Job prefix	Method	Pearson diff.	Alpha	Beta	Distance MIA AUC
0002af084a92	independent_marginals	0.805	0.602	0	0.474
000377f02a18	DOPE	0.071	0.836	0.558	0.504
000f66aa8642	DOPE	0.266	0.876	0.187	0.595
001897658278	Chow-Liu	0.194	0.804	0.450	0.632
0021656cfc57	Forest-Flow	0.046	0.780	0.555	0.614
00514b7a27e2	DOPE	0.143	0.875	0.239	0.517
0056ef8c748b	independent_marginals	0.149	0.696	0.361	0.637
006a72169115	GaussianCopula	0.064	0.878	0.383	0.452
006cafa72e1a	DOPE	0.069	0.838	0.725	0.710
006dfe86527f	independent_marginals	0.072	0.805	0.491	0.478
007f4ac594ee	TVAE	0.134	0.944	0.523	0.615
0090be69a85e	GaussianCopula	0.098	0.867	0.480	0.393

Table 39: **BeyondArena families that prepared.** The denominator is the 12-family fit check, distinct from the 142-family inventory. No retention auditor and no Holm family.

Family	Fit rows	2,048 B	8,192 B	Val 2,048	Val 8,192
airfoil	802	804	818	0.010	0.00448
garments	939	2,317	2,372	0.028	0.025
musk	3,738	9,273	9,350	0.146	0.090
telemonitoring	257	1,852	1,890	0.033	0.019
video	35,918	5,054	5,071	0.00247	0.00195