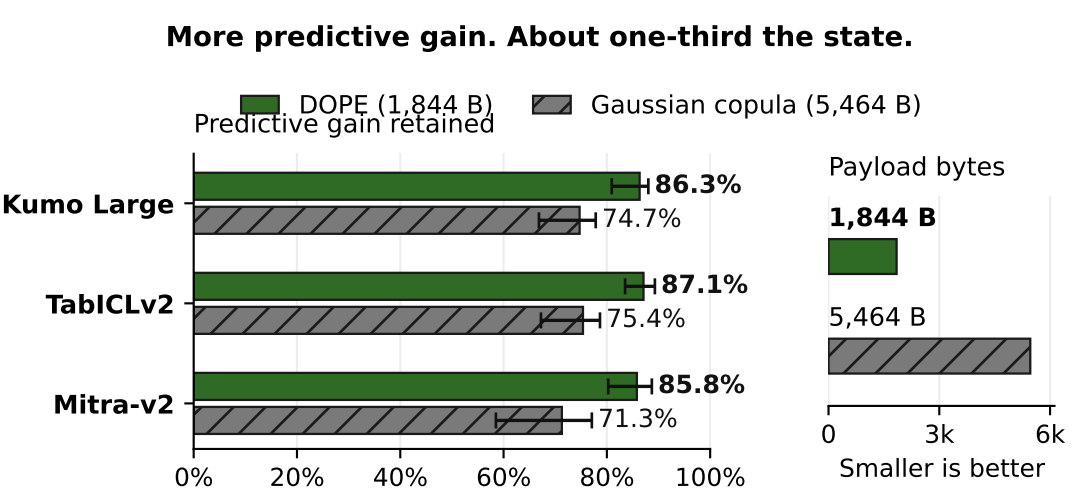
More predictive gain. About one-third the state.
DOPE (1,844 B)
Gaussian copula (5,464 B)
Predictive gain retained
Kumo Large
TabICLv2
Mitra-v2
86.3%
74.7%
87.1%
75.4%
85.8%
71.3%
0%
20%
40%
60%
80%
100%
Transfer retention (%)
Payload bytes
1,844 B
5,464 B
0
3k
6k
Smaller is better
97 matched lineages; 4n draws; whiskers: historical 95% interval
Selection-exposed validation; cold sampling closure: unmeasured
Cross-model feature-driver preservation: not yet measured