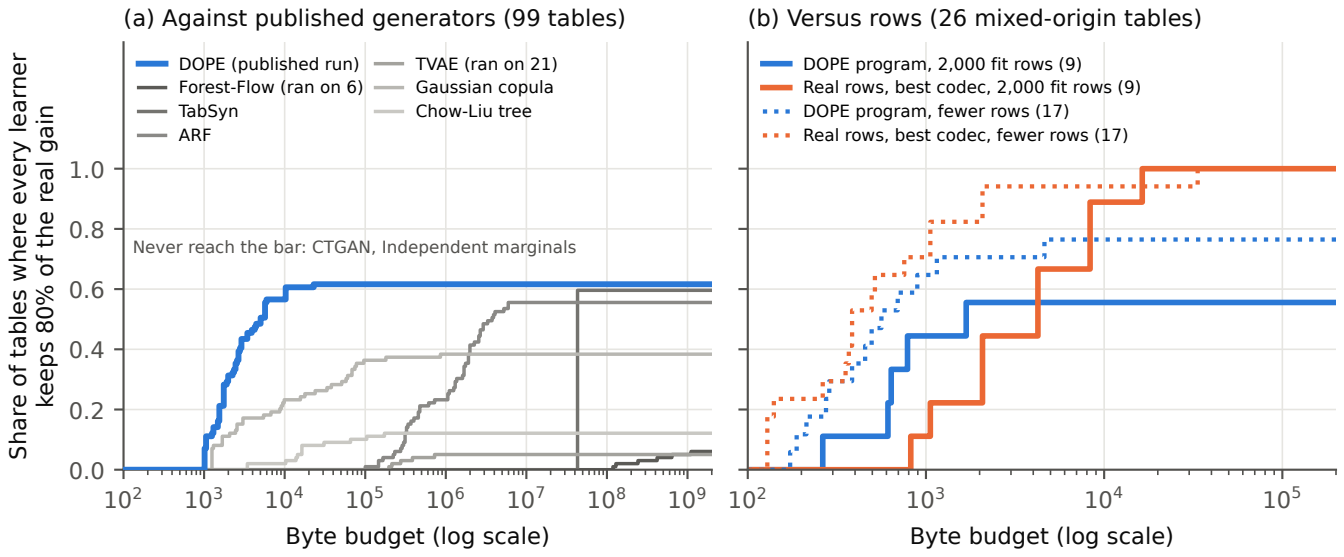
(a) Against published generators (99 tables)
DOPE (published run)
Forest-Flow (ran on 6)
TabSyn
ARF
TVAE (ran on 21)
Gaussian copula
Chow-Liu tree
Never reach the bar: CTGAN, Independent marginals
Share of tables where every learner keeps 80% of the real gain
1.0
0.8
0.6
0.4
0.2
0.0
10^2
10^3
10^4
10^5
10^6
10^7
10^8
10^9
Byte budget (log scale)
(b) Versus rows (26 mixed-origin tables)
DOPE program, 2,000 fit rows (9)
Real rows, best codec, 2,000 fit rows (9)
DOPE program, fewer rows (17)
Real rows, best codec, fewer rows (17)
10^2
10^3
10^4
10^5
Byte budget (log scale)